
Introduction to Process

Operating Systems

School of Data & Computer Science
Sun Yat-sen University

Lecture Notes: os_sysu@163.com
Instructor: Guoyang Cai
email: isscg@mail.sysu.edu.cn



中山大學
SUN YAT-SEN UNIVERSITY



■ Contents

- Basic Concepts
- Process Table and Process Control Block
- Process States and Transitions
- Operations on Process
 - Process Creation
 - Process Termination
- Unix and Linux Examples
- Process Scheduling
- Process Switching



■ Basic Concepts

- Process is a program in execution; forms the basis of all computation. Process execution must progress in sequential fashion.
 - A program is a **passive** entity containing a list of instructions stored on disk as an executable file.
 - A process is an **active** entity with some specifications of the corresponding program and a set of associated resources.
 - A program becomes a process when the executable object of this program is **loaded** into memory, given to the process scheduler.
 - A process is an **instance** of a running program; it can be assigned to, and executed on, a processor.
- Execution of program can be started via CLI entry of its name, GUI mouse clicks, etc.
- Related terms for Process
 - Job, Step, Load Module, Task, Thread.



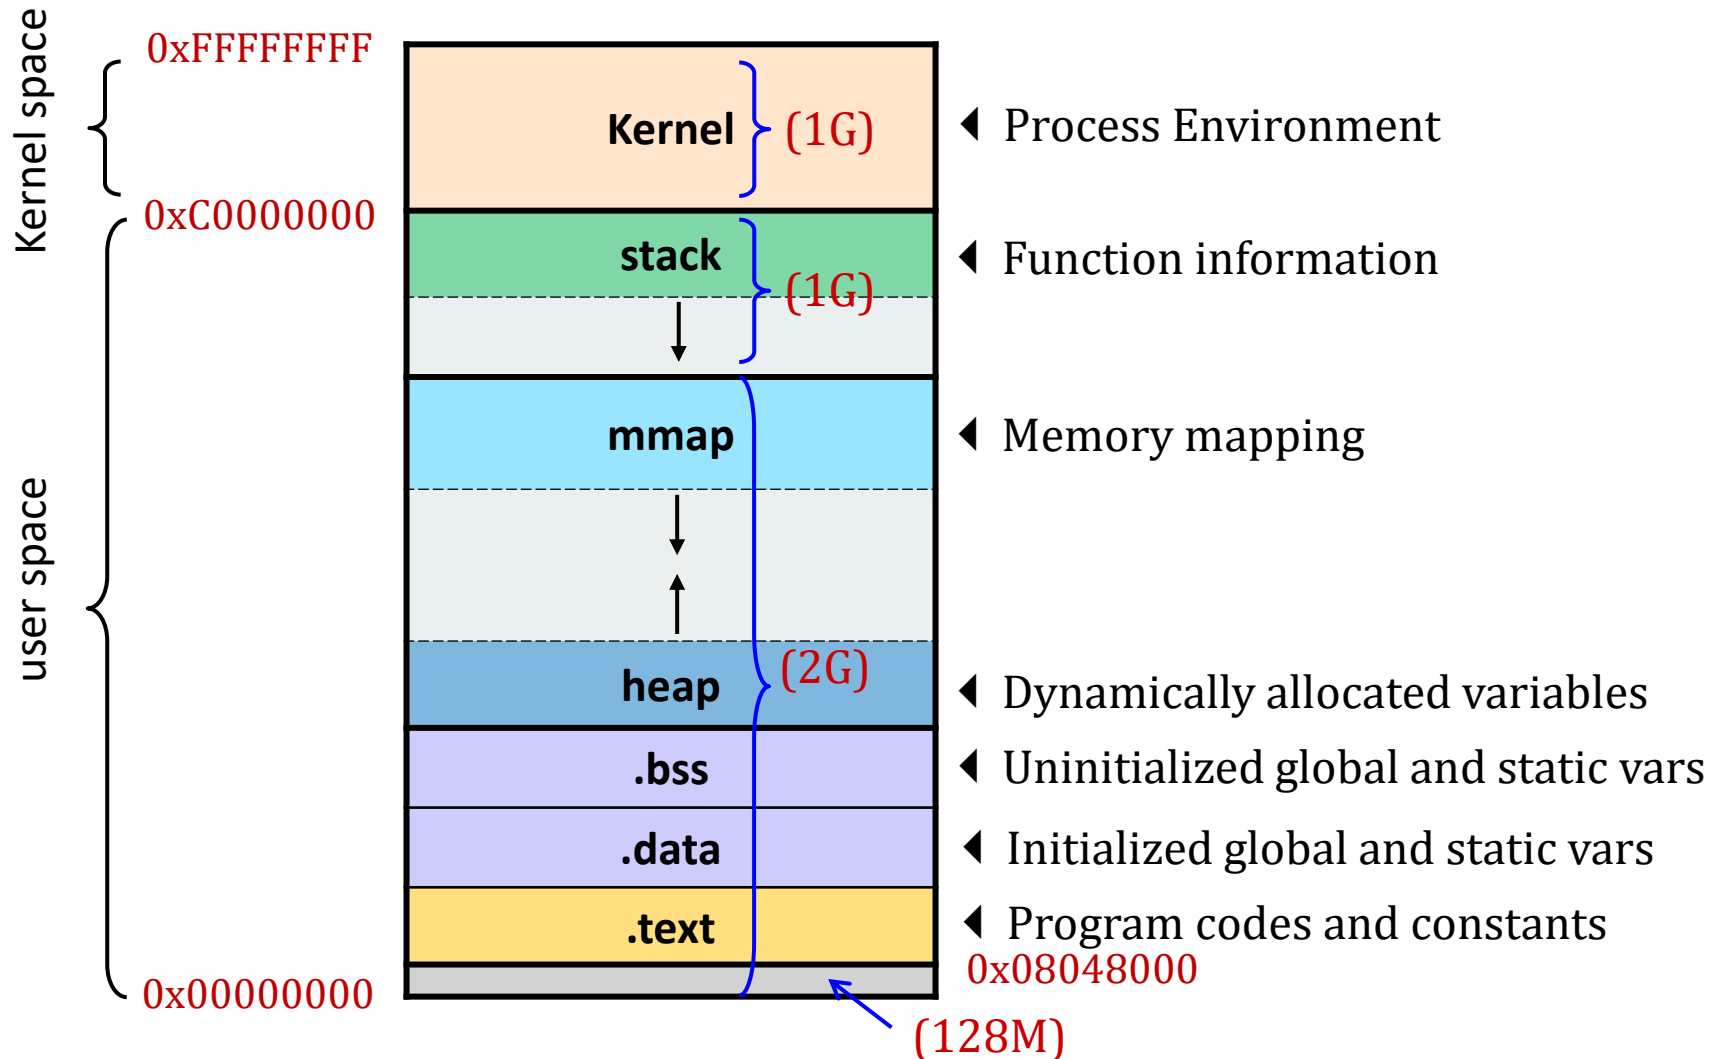
■ Basic Concepts

- A process includes some **segments**/sections:
 - Text
 - the executable (program) code
 - Data & Heap
 - Data
 - global variables
 - Heap
 - memory dynamically allocated during run time
 - Stack
 - temporary data storage
 - procedure/function parameters, return addresses, local variables
- Current activity of a program, or a process, includes its **context**.
 - program counter (PC), processor registers, etc.
- One program can be corresponding to several processes.
 - multiple users executing the same sequential program
 - concurrent program running several processes



Basic Concepts

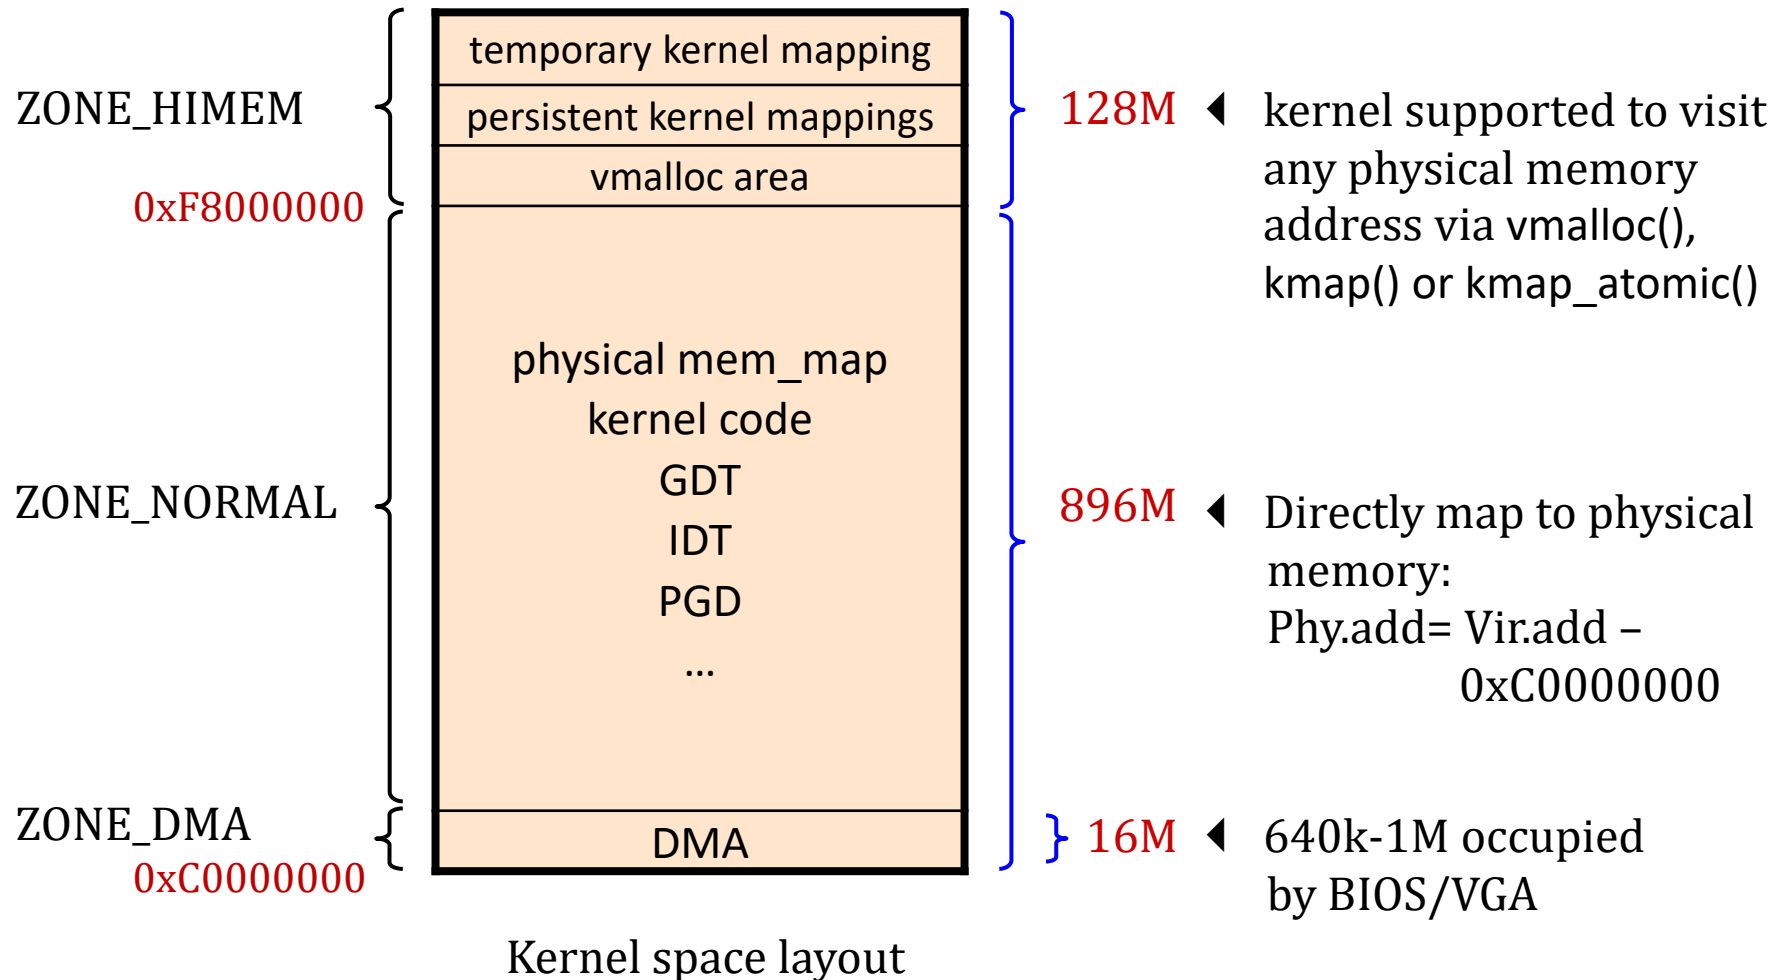
- Process Virtual Memory - Typical layout on Linux/IA-32.





■ Basic Concepts

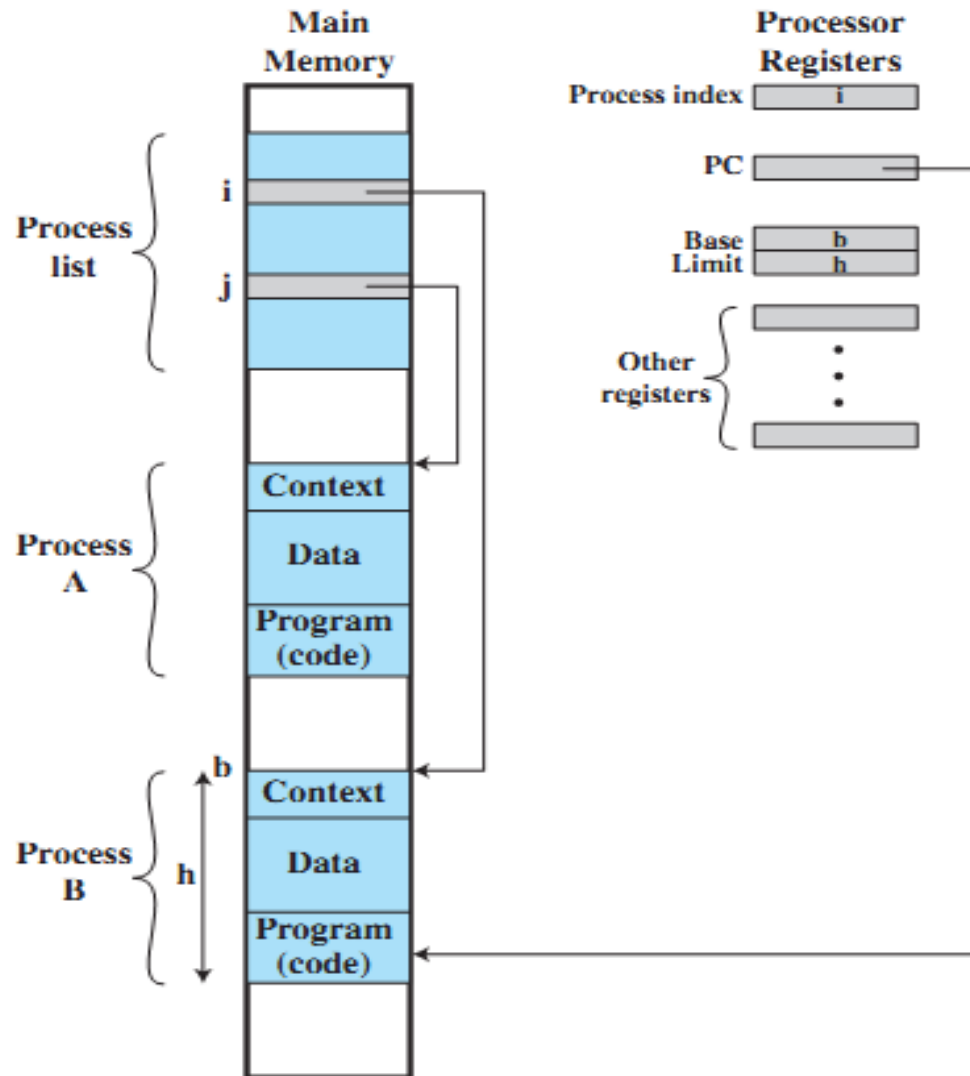
- Process Virtual Memory - Typical layout on Linux/IA-32.





Basic Concepts

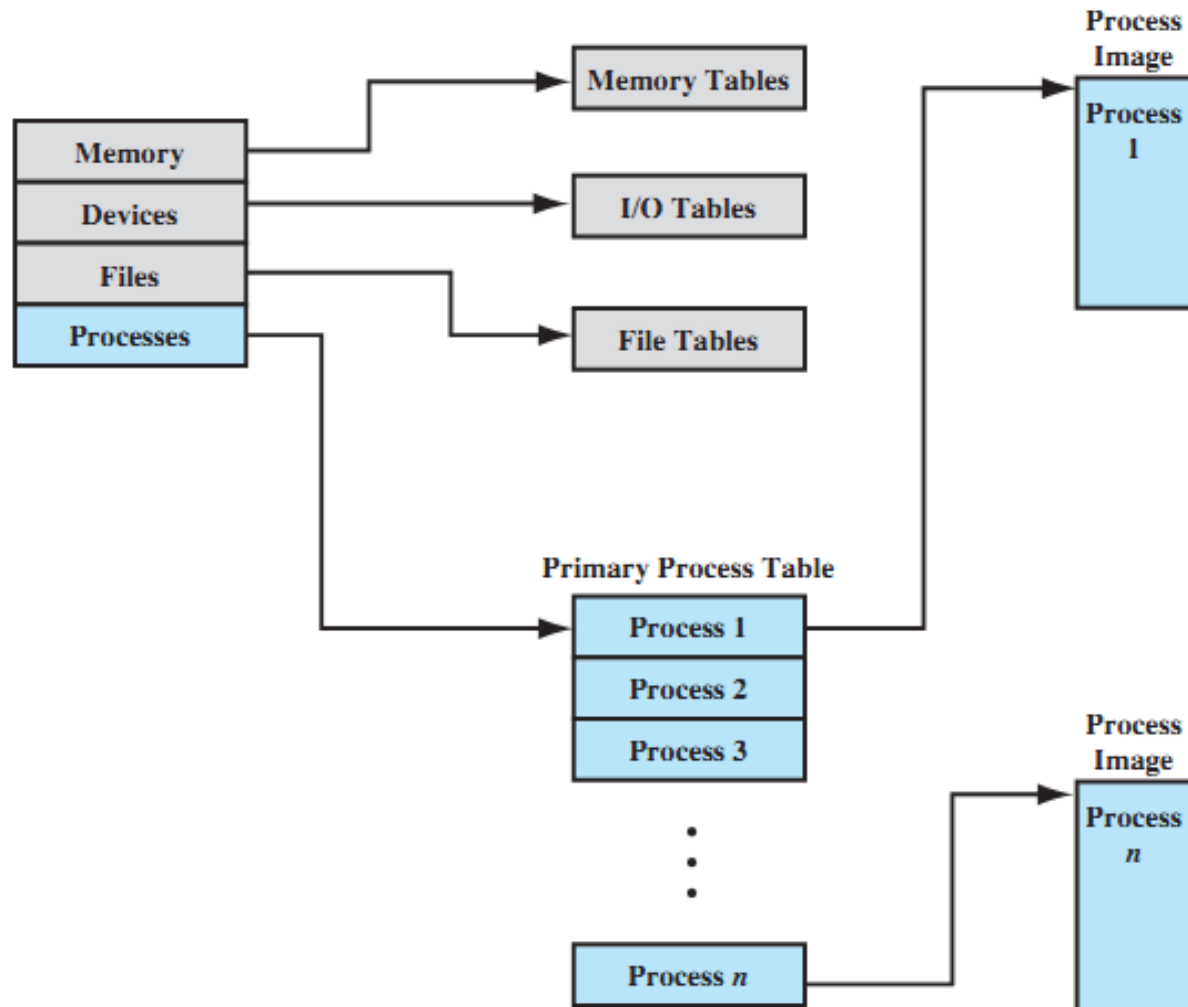
Typical Process Table Implementation.





Basic Concepts

- General Structure of OS Control Tables.





■ Basic Concepts

- Attributes of Process
 - Process ID
 - Parent process ID
 - User ID
 - Process state
 - Process priority
 - Program counter
 - CPU registers
 - Memory management information
 - I/O status information
 - Access control
 - Accounting information



■ Process Table

- *Process table* is a kernel data structure containing fields that must always be available to the kernel.
 - state field (that identifies the state of the process)
 - fields that allow kernel to locate the process in memory
 - UIDs for determining various process privileges
 - PIDs to specify relationships b/w processes (e.g. fork)
 - event descriptor (when the process in sleep state)
 - scheduling parameters to determine the order in which process moves to the states "kernel running" and "user running"
 - signal field for signals send to the process but not yet handled
 - timers that give process execution time in kernel mode and user mode
 - field that gives process size (so that kernel knows how much space to allocate for the process).



■ Process Table

■ Fields of a Typical Process Table.

Process management	Memory management	File management
Registers Program Counter Program Status Word Stack Pointer Process state Priority Scheduling parameters Process ID Parent process Process group Signals Time when process started CPU time used Children's CPU time Time of next alarm	Pointer to text segment info Pointer to data segment info Pointer to stack segment info	Root directory Working directory File descriptors User ID Group ID



■ Process Control Block (PCB)

- Each process is represented in OS by a *process control block* (PCB)
 - PCB also called a *task control block*, IBM name for information associated with a specific process.
 - It saves context of the Process.
- PCB is the data (Process Attributes) needed by OS to control process:
 - Process location information
 - Process identification information
 - Processor state information
 - Process control information.



Process control block (PCB)



■ Process Control Block (PCB)

■ Process Location Information

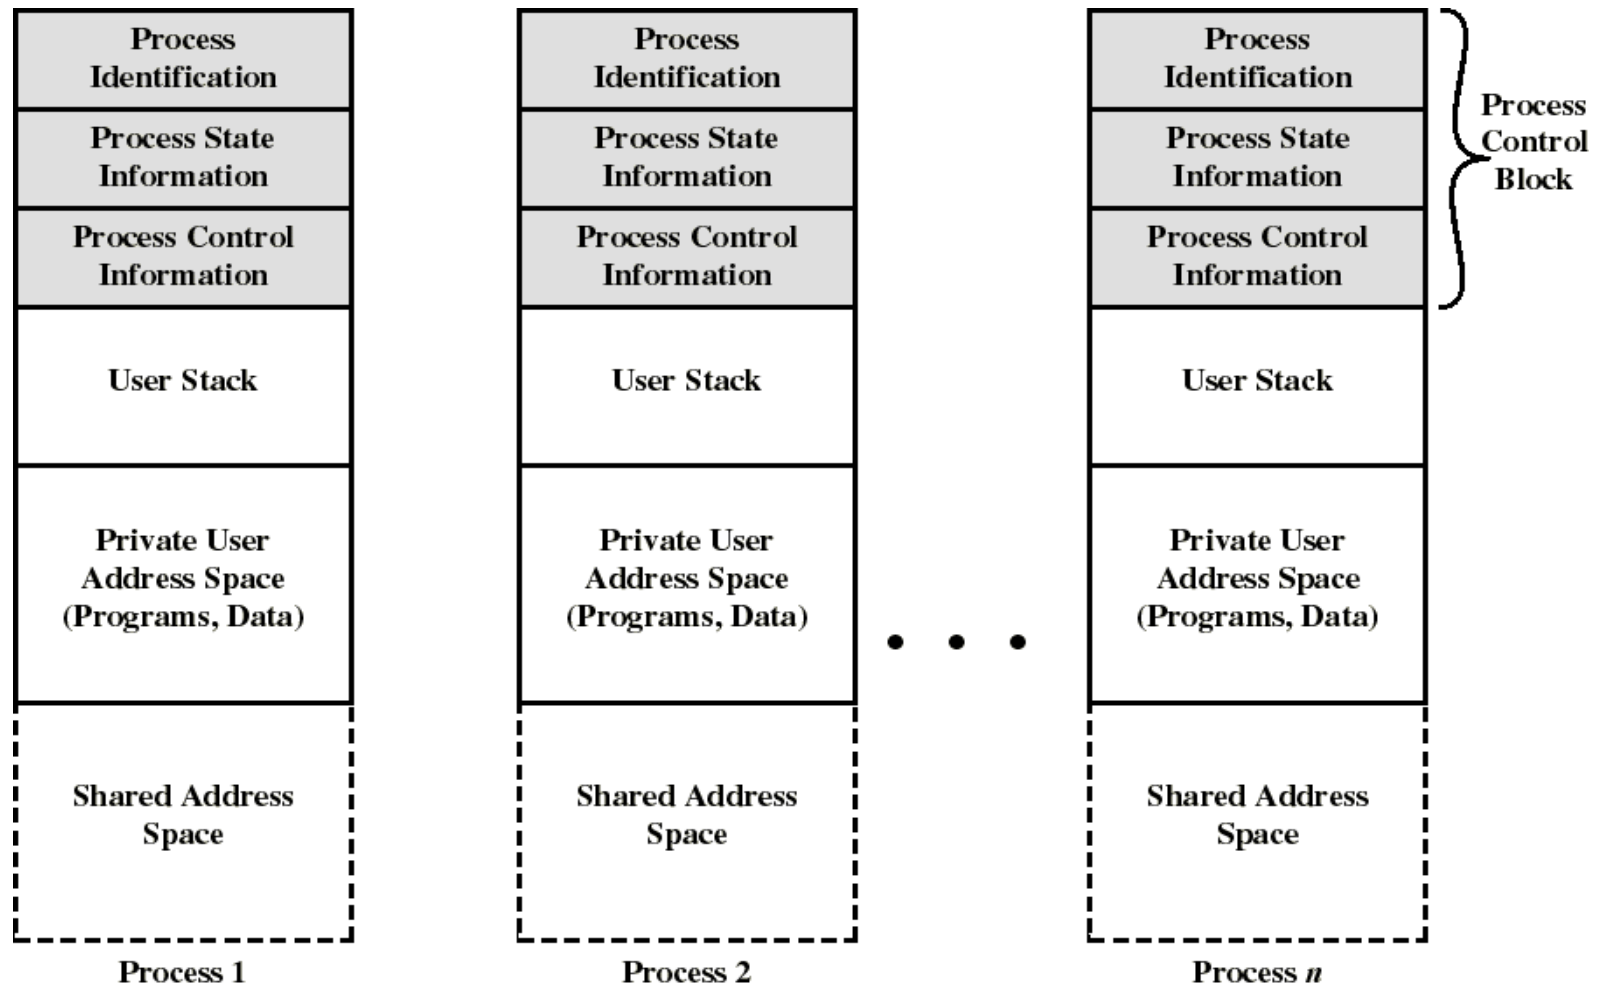
■ Process image

- Each process has an image in memory.
 - It may not occupy a contiguous range of addresses.
 - depends on memory management scheme used
 - Both a private and shared memory address space can be used.
- Each process image is pointed to by an entry in the process table.
 - For the OS to manage the process, at least part of its image must be brought into main memory.



■ Process Control Block (PCB)

■ Process Location Information.



Process Images in Memory



■ Process Control Block (PCB)

■ Process Identification Information

■ A few numeric identifiers may be used:

- Unique process identifier (PID)
 - indexes (directly or indirectly) into the process table
- User identifier (UID)
 - the user who is responsible for the job
- Identifier of the process that created this process (PPID, Parent process ID)

■ Maybe symbolic names that are related to numeric identifiers



■ Process Control Block (PCB)

■ Processor State Information

- contents of processor registers
 - User-visible registers
 - Control and status registers
 - Stack pointers.

■ Program Status Word (PSW)

- contains status information
- E.g.
 - the EFLAGS register on Pentium machines



■ Process Control Block (PCB)

■ Process Control Information

■ Scheduling and state information

- *Process state* (e.g., running, ready, blocked...)
- Priority of the process
- Event for which the process is waiting (if blocked)

■ Data structuring information

- may hold pointers to other PCBs for process queues, parent-child relationships and other structures

■ InterProcess Communication (IPC)

- may hold flags and signals for IPC

■ Resource ownership and utilization

- resource in use: open files, I/O devices...
- history of usage (of CPU time, I/O...)

■ Process privileges (Access Control)

- access to certain memory locations, to resources, etc.

■ Memory management

- pointers to segment/page tables assigned to this process



■ Process Control Block (PCB)

■ Process Control Information (cont.)

■ Program counter

- indicates the address of the next instruction to be executed for this process

■ Accounting information

- includes the amount of CPU and real time used, time limits, account numbers, job or process numbers, and so on

■ I/O status information

- includes the list of I/O devices allocated to the process, a list of open files, and so on



■ Process States and Transitions

■ Three-state Process Model

■ Running State

- the process that gets executed; its instructions are being executed

■ Ready State

- any process that is ready to be executed; the process is waiting to be assigned to a processor

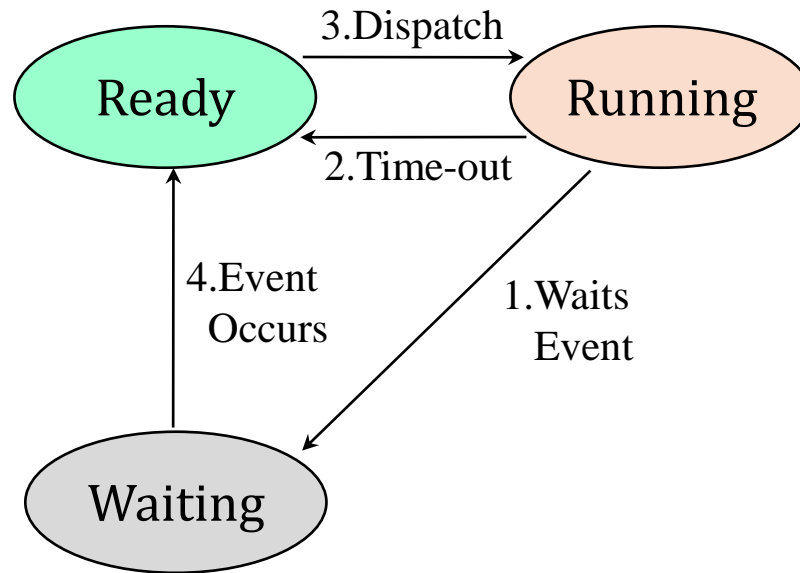
■ Waiting/Blocked State

- any process that cannot execute until its I/O completes or some other event occurs



■ Process States and Transitions

■ Three-state Process Model.



1. Process blocks for input.
2. Scheduler picks another process.
3. Scheduler picks this process.
4. Input becomes available.



■ Process States and Transitions

■ Three-state Process Model


■ Process State Transitions

- Ready → Running
 - When it is time, the dispatcher selects a new process to run.
- Running → Ready
 - The running process has expired his time slot.
 - The running process gets interrupted because a higher priority process is in the ready state.
- Running → Waiting
 - When a process requests something for which it must wait:
 - a service that the OS is not ready to perform
 - an access to a resource not yet available
 - Initiating I/O and waiting for the result
 - waiting for a process to provide input
- Waiting → Ready
 - When the event for which the process was waiting occurs



■ Process States and Transitions

- Three-state Process Model
 - Process List Structures.



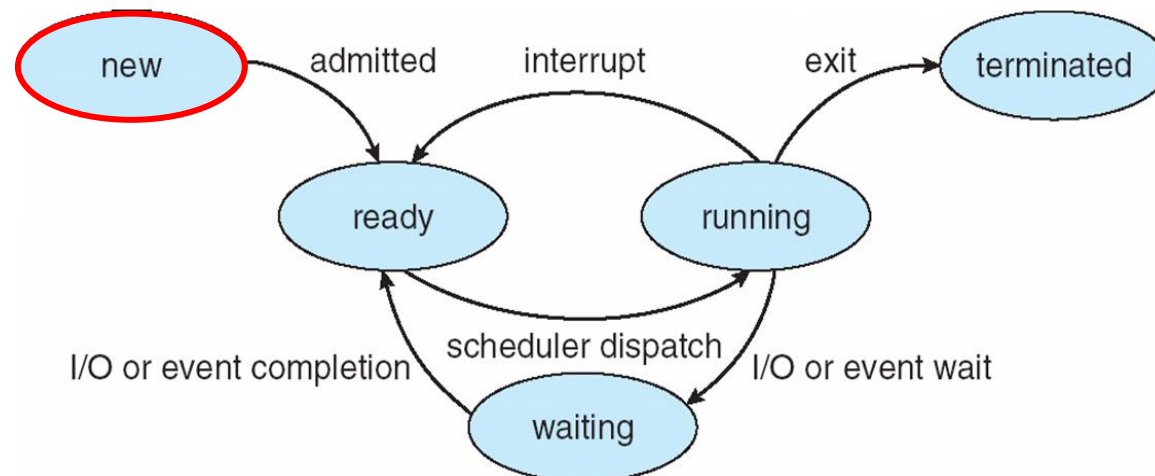


■ Process States and Transitions

■ Five-state Process Model

■ New state

- OS has performed the necessary actions to create the process:
 - has created a process identifier
 - has created tables needed to manage the process
- but has **not yet committed** to execute the process (not yet admitted):
 - because resources are limited



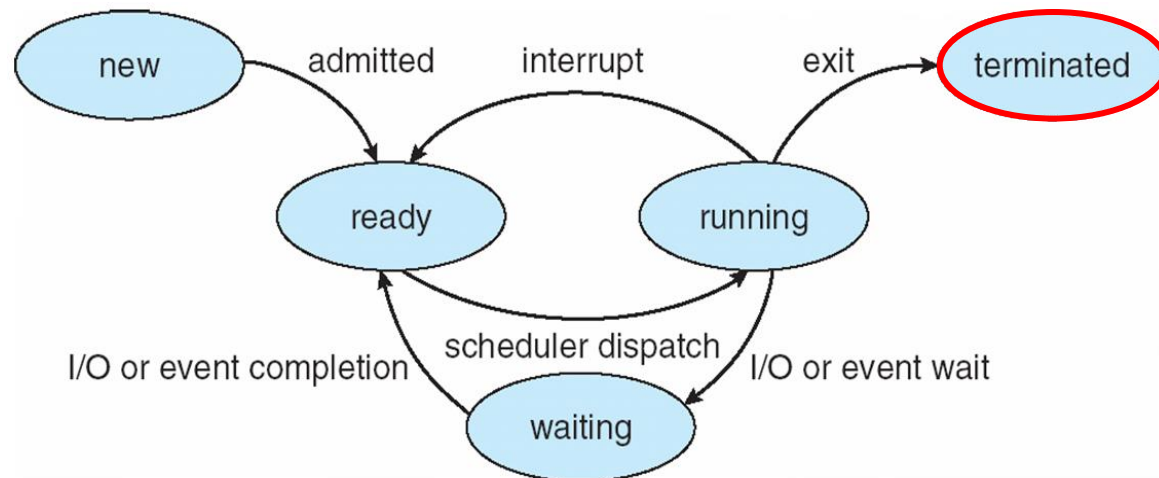


■ Process States and Transitions

■ Five-state Process Model

■ Terminated state

- Process termination moves the process to terminate state.
 - It is no longer eligible for execution.
- Tables and other information are temporarily preserved for auxiliary program.
 - E.g., accounting program that cumulates resource usage for billing the users
- The process (and its tables) gets deleted when the data is no more needed.





■ Process Creation

- When to Create a Process
 - System initialization
 - Submission of a batch job
 - User logs on
 - Created by OS to provide a service to a user
 - e.g., printing a file
 - A user request to create a new process
 - Spawned (繁衍) by an existing process
 - A program can dictate (to require or determine necessarily) the creation of a number of processes.
 - The creating process is the *parent* process and the new processes created are called the *children* of that process.



■ Process Creation

■ Details in Process Creating

- Parent process create children processes, which, in turn create other processes, forming a *tree* of processes.
- Possible resource sharing:
 - Parent process and children processes **share all** resources.
 - Children processes **share subset** of parent's resources.
 - Parent process and child process **share no** resources.
- Possible execution:
 - Parent process and children processes execute **concurrently**.
 - Parent process **waits** until children processes terminate.



■ Process Creation

- Details in Process Creating (cont.)
 - Assign a unique **process identifier** (PID).
 - typically an integer number
 - Allocate space for the **process image**.
 - Initialize **Process Control Block** (PCB).
 - many default values
 - E.g., state is New, no I/O devices or files,
 - Set up appropriate **linkages**.
 - E.g., add new process to linked list used for the scheduling queue
 - **Address space**
 - child is a duplicate of parent, or
 - child has a program loaded into it.

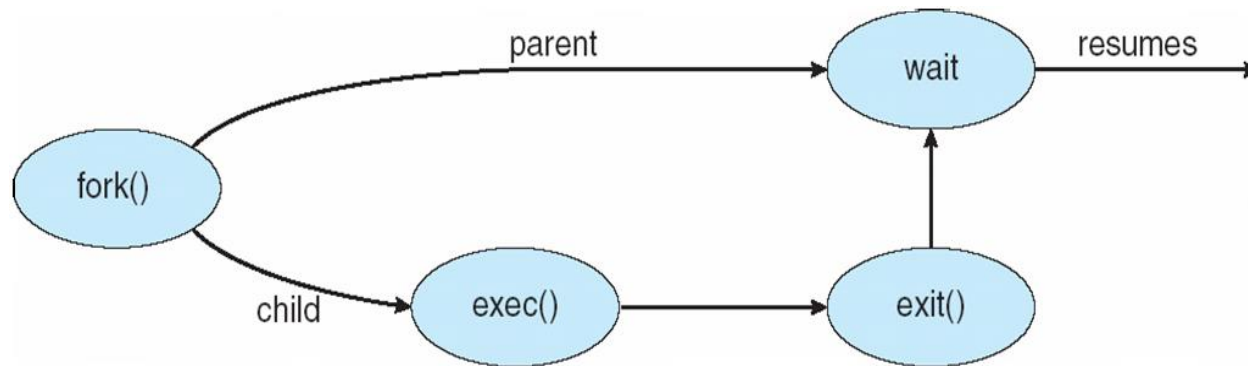


■ Process Creation

■ Details in Process Creating (cont.)

■ UNIX examples

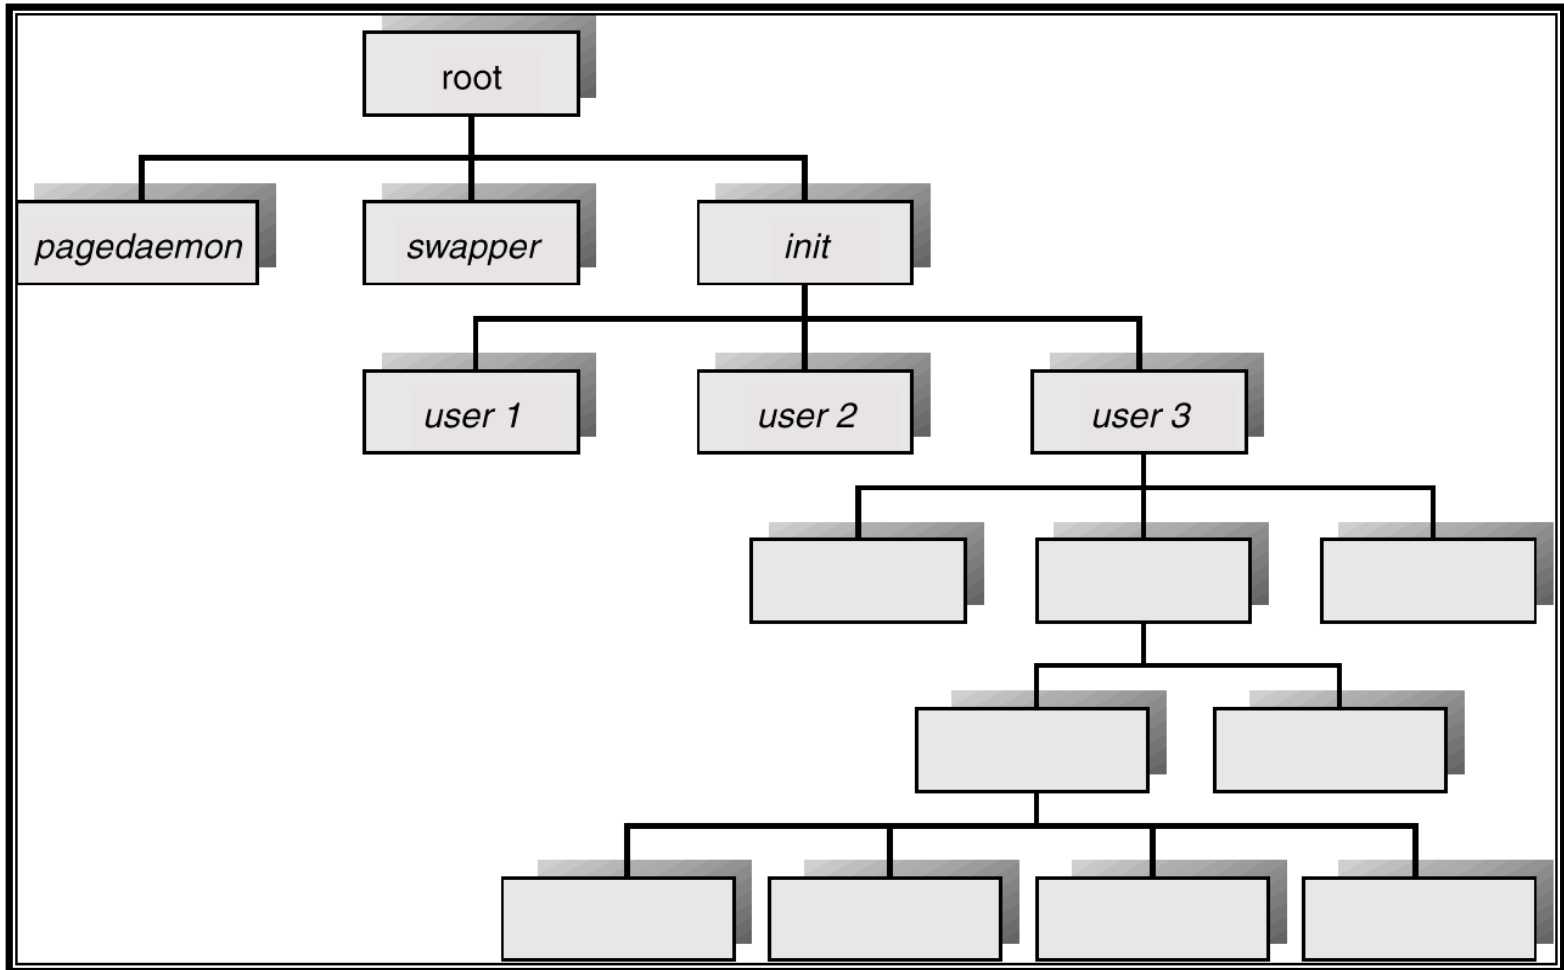
- `fork()` system call creates new process.
- `exec()` system call used after a `fork()` to replace the memory space of the process with a new program.





■ Process Creation

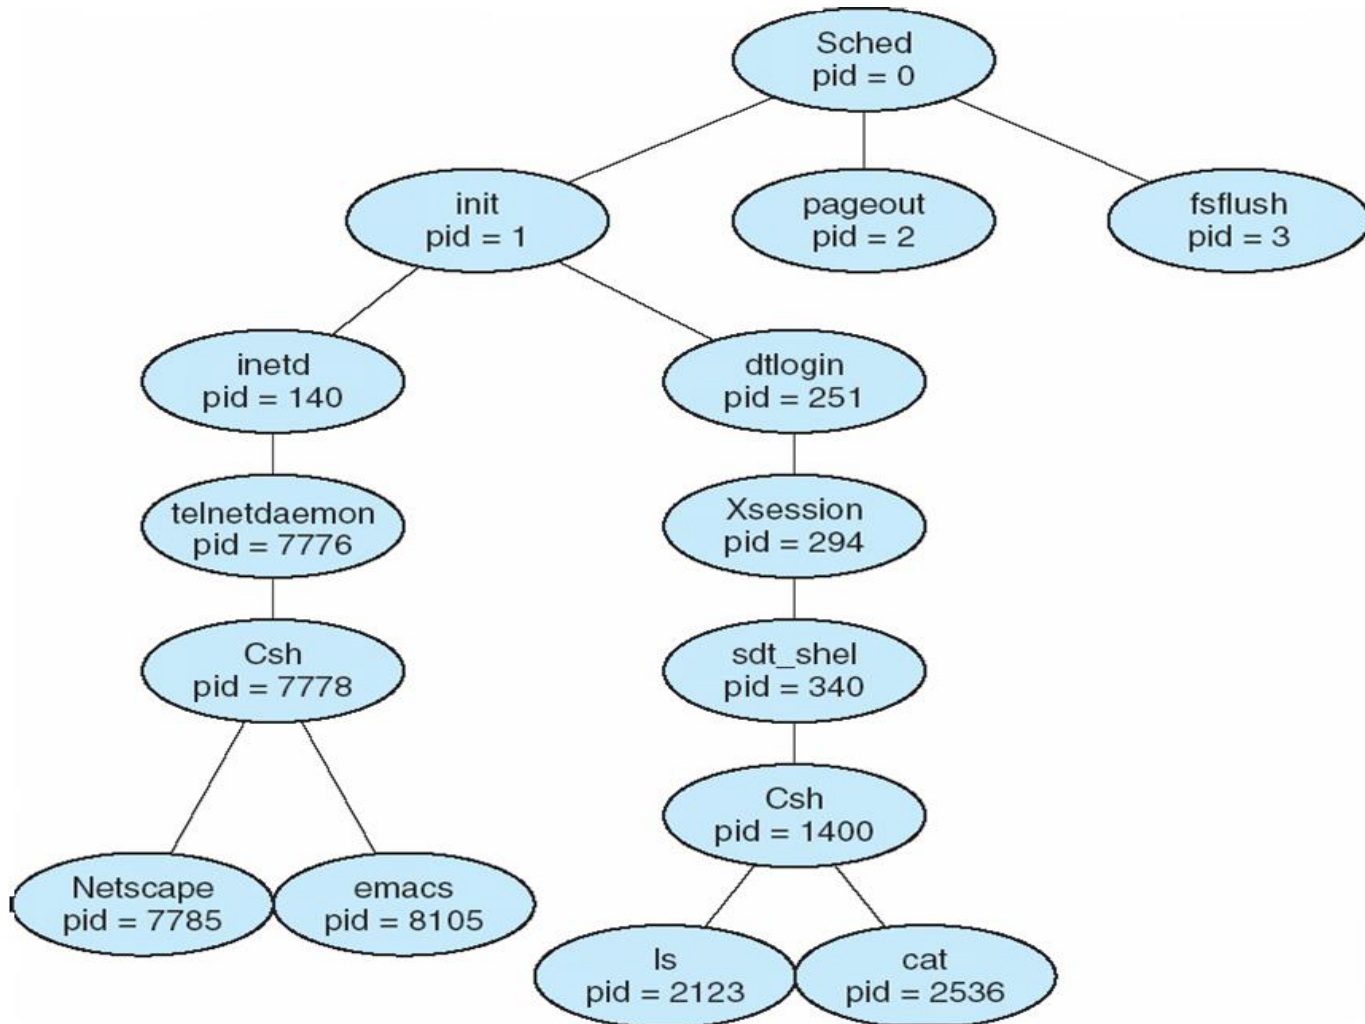
- A Tree of Processes on UNIX.





■ Process Creation

- A Tree of Processes on Typical Solaris.





■ Process Creation

■ Algorithm 6-1: fork-demo.c (forking a separate process).

```
int main(void)
{
    int count = 1;
    pid_t childpid;

    childpid = fork(); /* child duplicates parent's address space */
    if (childpid < 0) {
        perror("fork()");
        return EXIT_FAILURE;
    }
    else /* fork() returns 2 values: 0 for child pro and childpid for parent pro */
        if (childpid == 0) { /* This is child pro */
            count++;
            printf("Child pro pid = %d, count = %d (addr = %p)\n", getpid(), count,
&count);
        }
        else { /* This is parent pro */
            printf("parent pro pid = %d, child pid = %d, count = %d (addr = %p)\n",
getpid(), childpid, count, &count);
            sleep(5);
            wait(0); /* waiting for all children terminated */
        }
        printf("Testing point by %d\n", getpid()); /* child executed this statement and
became defunct before parent wait()*/
        return EXIT_SUCCESS;
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
#include <sys/wait.h>
```



Process Creation

- Algorithm 6-1: fork-demo.c (forking a separate process).

```
int main(void)
{
    int count = 1;
```

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
#include <sys/wait.h>
```

```
iisscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-1-fork-demo.c
iisscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out
Parent pro pid = 6452, child pid = 6453, count = 1 (addr = 0x7fffd86d2660)
Child pro pid = 6453, count = 2 (addr = 0x7fffd86d2660)
Testing point by 6453
Testing point by 6452
iisscgy@ubuntu:/mnt/hgfs/os-2020$
```

```
        printf("Child pro pid = %d, count = %d (addr = %p)\n", getpid(), count,
&count);
    }
    else { /* This is parent pro */
        printf("parent pro pid = %d, child pid = %d, count = %d (addr = %p)\n",
getpid(), childpid, count, &count);
        sleep(5);
        wait(0); /* waiting for all children terminated */
    }
    printf("Testing point by %d\n", getpid()); /* child executed this statement and
became defunct before parent wait()*/
    return EXIT_SUCCESS;
}
```




■ Process Creation

- Algorithm 6-1: fork-demo.c (forking a separate process).

```
int main(void)
{
    int count = 1;
```

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
#include <sys/wait.h>
```

```
isscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-1-fork-demo.c
isscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out
Parent pro pid = 6452, child pid = 6453, count = 1 (addr = 0x7fffd86d2660)
Child pro pid = 6453, count = 2 (addr = 0x7fffd86d2660)
Testing point by 6453
Testing point by 6452
isscgy@ubuntu:/mnt/hgfs/os-2020$
```

The variable `count` in child process has the same **virtual address** with that in parent process.

```
    printf("Child pro pid = %d, count = %d (addr = %p)\n",
    &count);
}
else { /* This is parent pro */
    printf("parent pro pid = %d, child pid = %d, count = %d (addr = %p)\n",
    getpid(), childpid, count, &count);
    sleep(5);
    wait(0); /* waiting for all children terminated */
}
printf("Testing point by %d\n", getpid()); /* child executed this statement and
became defunct before parent wait()*/
return EXIT_SUCCESS;
}
```



■ Process Creation

■ Algorithm 6-1: fork-demo.c (forking a separate process).

```
int main(void)
{
    int count = 1;
```

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
#include <sys/wait.h>
```

```
isscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-1-fork-demo.c
isscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out
Parent pro pid = 6452, child pid = 6453, count = 1 (addr = 0x7fffd86d2660)
Child pro pid = 6453, count = 2 (addr = 0x7fffd86d2660)
Testing point by 6453
Testing point by 6452
isscgy@ubuntu:/mnt/hgfs/os-2020$
```

```
printf("Child\n",
&count);
}
else { /* This is parent pro */
    printf("parent pro pid = %d, child pid = %d, count = %d (addr = %p)\n",
getpid(), childpid, count, &count);
    sleep(5);
    wait(0); /* waiting for all children terminated */
}
printf("Testing point by %d\n", getpid()); /* child executed this statement and
became defunct before parent wait()*/
return EXIT_SUCCESS;
}
```

The value of `count` in child process is different from that in parent process. They are mapped to different **physical addresses** in different process images.



Process Creation

- Algorithm 6-1: fork-demo.c (forking a separate process).

```
int main(void)
{
    int count = 1;
```

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
#include <sys/wait.h>
```

```
isscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-1-fork-demo.c
isscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out
Parent pro pid = 6452, child pid = 6453, count = 1 (addr = 0x7fffd86d2660)
Child pro pid = 6453, count = 2 (addr = 0x7fffd86d2660)
Testing point by 6453
Testing point by 6452
isscgy@ubuntu:/mnt/hgfs/os-2020$
```

Testing point executed both by child pro
and parent pro

```
%d, count = %d (addr = %p)\n", getpid(), count,
```

```
    }
    else { /* This is parent pro */
        printf("parent pro pid = %d, child pid = %d, count = %d (addr = %p)\n",
getpid(), childpid, count, &count);
        sleep(5);
        wait(0); /* waiting for all children terminated */
    }
    printf("Testing point by %d\n", getpid()); /* child executed this statement and
became defunct before parent wait()*/
    return EXIT_SUCCESS;
}
```



■ Process Creation

■ Algorithm 6-2: vfork-demo.c (vforking a sharing-space process).

```
int main(void)
{
    int count = 1;
    pid_t childpid;

    childpid = vfork(); /* child shares parent's address space */
    if (childpid < 0) {
        perror("fork()");
        return EXIT_FAILURE;
    }
    else /* vfork() returns 2 values: 0 for child pro and childpid for parent pro */
        if (childpid == 0) { /* This is child pro, parent hung up until child exit */
            count++;
            printf("Child pro pid = %d, count = %d (addr = %p)\n", getpid(), count,
&count);

            printf("Child taking a nap ...\n");
            sleep(10); printf("Child waking up!\n");
            _exit(0); /* or exec(0); "return" will cause stack smashing */
        }
        else { /* This is parent pro, start when the vforked child terminated */
            printf("parent pro pid = %d, child pid = %d, count = %d (addr = %p)\n",
getpid(), childpid, count, &count);
            wait(0); /* not waitting this vforked child terminated */
        }
        printf("Testing point by %d\n", getpid()); /* executed by parent pro only */
        return EXIT_SUCCESS;
}
```



Process Creation

- Algorithm 6-2: vfork-demo.c (vforking a sharing-space process).

```
int main(void)
{
    int count = 1;
    pid_t childpid;
```

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
#include <sys/wait.h>
```

```
iisscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-2-vfork-demo.c
iisscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out
Child pro pid = 15418, count = 2 (addr = 0x7ffd866a34b0)
Child taking a nap ...
Child waking up!
Parent pro pid = 15417, child pid = 15418, count = 2 (addr = 0x7ffd866a34b0)
Testing point by 15417
iisscgy@ubuntu:/mnt/hgfs/os-2020$
```

```
    &count);
    printf("Child taking a nap ...\n");
    sleep(10); printf("Child waking up!\n");
    _exit(0); /* or exec(0); "return" will cause stack smashing */
}
else { /* This is parent pro, start when the vforked child terminated */
    printf("parent pro pid = %d, child pid = %d, count = %d (addr = %p)\n",
getpid(), childpid, count, &count);
    wait(0); /* not waitting this vforked child terminated */
}
printf("Testing point by %d\n", getpid()); /* executed by parent pro only */
return EXIT_SUCCESS;
}
```



Process Creation

Algorithm 6-2: vfork-demo.c (vforking a sharing-space process).

```
int main(void)
{
    int count = 1;
    pid_t childpid;
```

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
#include <sys/wait.h>
```

```
iisscg@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-2-vfork-demo.c
iisscg@ubuntu:/mnt/hgfs/os-2020$ ./a.out
Child pro pid = 15418, count = 2 (addr = 0x7ffd866a34b0)
Child taking a nap ...
Child waking up!
Parent pro pid = 15417, child pid = 15418, count = 2 (addr = 0x7ffd866a34b0)
Testing point by 15417
iisscg@ubuntu:/mnt/hgfs/os-2020$
```

```
&count);
```

```
printf("Child taking a nap
sleep(10); printf("Child wa
_exit(0); /* or exec(0); "r
```

```
}
```

```
else { /* This is parent pro, start when the vforked child terminated */
```

```
printf("parent pro pid = %d, child pid = %d, count = %d (addr = %p)\n",
```

```
getpid(), childpid, count, &count);
```

```
wait(0); /* not waitting this vforked child terminated */
```

```
}
```

```
printf("Testing point by %d\n", getpid()); /* executed by parent pro only */
```

```
return EXIT_SUCCESS;
```

```
}
```

The variable **count** in child process has the same **virtual address** with that in parent process.



Process Creation

- Algorithm 6-2: vfork-demo.c (vforking a sharing-space process).

```
int main(void)
{
    int count = 1;
    pid_t childpid;
```

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
#include <sys/wait.h>
```

```
iisscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-2-vfork-demo.c
iisscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out
Child pro pid = 15418, count = 2 (addr = 0x7ffd866a34b0)
Child taking a nap ...
Child waking up!
Parent pro pid = 15417, child pid = 15418, count = 2 (addr = 0x7ffd866a34b0)
Testing point by 15417
iisscgy@ubuntu:/mnt/hgfs/os-2020$
```

```
&count);
```

```
printf("Child taking a nap\n");
sleep(10); printf("Child waking up\n");
_exit(0); /* or exec(0);
```

```
}
```

```
else { /* This is parent pro, start when the vforked child terminated */
```

```
printf("parent pro pid = %d, child pid = %d, count = %d (addr = %p)\n",
getpid(), childpid, count, &count);
```

```
wait(0); /* not waiting this vforked child terminated */
```

```
}
```

```
printf("Testing point by %d\n", getpid()); /* executed by parent pro only */
```

```
return EXIT_SUCCESS;
```

```
}
```

The value of `count` in child process is the same as that in parent process. They are mapped to the same **physical address** in the same process images.



Process Creation

- Algorithm 6-2: vfork-demo.c (vforking a sharing-space process).

```
int main(void)
{
    int count = 1;
    pid_t childpid;
```

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
#include <sys/wait.h>
```

```
iisscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-2-vfork-demo.c
iisscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out
Child pro pid = 15418, count = 2 (addr = 0x7ffd866a34b0)
Child taking a nap ...
Child waking up!
Parent pro pid = 15417, child pid = 15418, count = 2 (addr = 0x7ffd866a34b0)
Testing point by 15417
iisscgy@ubuntu:/mnt/hgfs/os-2020$
```

parent pro is hung up until the vforked child terminated.

```
    sleep(10); printf("Child waking up!\n");
    _exit(0); /* or exec(0); "return" will cause stack smashing */
}
else { /* This is parent pro, start when the vforked child terminated */
    printf("parent pro pid = %d, child pid = %d, count = %d (addr = %p)\n",
        getpid(), childpid, count, &count);
    wait(0); /* not waitting this vforked child terminated */
}
printf("Testing point by %d\n", getpid()); /* executed by parent pro only */
return EXIT_SUCCESS;
}
```




Process Creation

- Algorithm 6-2: vfork-demo.c (vforking a sharing-space process).

```
int main(void)
{
    int count = 1;
    pid_t childpid;
```

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <unistd.h>
#include <sys/wait.h>
```

```
iisscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-2-vfork-demo.c
iisscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out
Child pro pid = 15418, count = 2 (addr = 0x7ffd866a34b0)
Child taking a nap ...
Child waking up!
Parent pro pid = 15417, child pid = 15418, count = 2 (addr = 0x7ffd866a34b0)
Testing point by 15417
iisscgy@ubuntu:/mnt/hgfs/os-2020$
```

Child exited before the testing point and it was executed by parent only

```
        _exit(0); /* or exec(0); "return" will cause stack smashing */
    }
    else { /* This is parent pro, start when the vforked child terminated */
        printf("parent pro pid = %d, child pid = %d, count = %d (addr = %p)\n",
            getpid(), childpid, count, &count);
        wait(0); /* not waitting this vforked child terminated */
    }
    printf("Testing point by %d\n", getpid()); /* executed by parent pro only */
    return EXIT_SUCCESS;
}
```



■ Process Termination

- A Process terminates when one of the following events happened
 - Batch job issues Halt instruction.
 - User logs off.
 - Process executes a service request to terminate.
 - Parent kills child process.
 - Error and fault conditions.



■ Process Termination

- Reasons for process termination
 - Normal/Error/Fatal exit
 - Time limit exceeded
 - Time overrun
 - process waited longer than a specified maximum for an event
 - Memory unavailable
 - Memory bounds violation
 - Protection error
 - e.g., write to read-only file
 - Arithmetic error
 - I/O failure
 - Invalid instruction
 - happens when trying to execute data
 - Privileged instruction
 - Operating system intervention (OS介入)
 - such as when deadlock occurs.
 - Parent request to terminate one child
 - Parent terminates so child processes terminate.



■ Process Termination

- Procedure of process termination
 - A process may execute last statement and ask the operating system to terminate it by `exit()` system call.
 - Its entry in the process table remains there until her parent, if exists, calls `wait()`.
 - Its resources are deallocated by operating system.
 - Parent may terminate execution of child processes:
 - Child has exceeded allocated resources.
 - Mission assigned to child is no longer required.
 - If parent process is exiting:
 - Some OSES do not allow child to continue if its parent terminates.
 - Cascading termination (级联终止) – all children terminated.



■ Process Termination

■ Procedure of process termination

■ Prototype of `wait()`

```
#include <sys/wait.h>
/* pid_t wait(int *status); */
```

```
pid_t pid;
int status;
pid = wait(&status);
```

- When a process terminates, its resources are deallocated by the operating system. However, its entry in the process table must remain there until the parent calls `wait()`, because the process table contains the process's exit status.



■ Process Termination

■ Zombies and Orphans

- A process that has terminated, but whose parent has not yet called `wait()`, is defunct and known as a **zombie process** (僵尸进程).
 - All processes transition to this state when they terminate, but generally they exist as zombies only briefly. Once the parent calls `wait()`, the process identifier of the zombie process and its entry in the process table are released.
- Now consider what would happen if a parent did not invoke `wait()` and instead terminated, thereby leaving its child processes as **orphans** (孤儿进程).
- Linux and UNIX address this scenario by assigning the **init process** as the new parent to orphan processes (**adoption** of orphans).
 - The init process is the root of the process hierarchy in UNIX and Linux systems.
 - The init process periodically invokes `wait()`, thereby allowing the exit status of any orphaned process to be collected and releasing the orphan's process identifier and process-table entry.



■ Process Termination

■ Algorithm 6-3: fork-demo-nowait.c (fork without waiting).

```
int main(void)
{
    int count = 1;
    pid_t childpid;

    childpid = fork(); /* child duplicates parent's address space */
    if (childpid < 0) {
        perror("fork()");
        return EXIT_FAILURE;
    }
    else
        if (childpid == 0) { /* This is child pro */
            count++;
            printf("child pro pid = %d, count = %d (addr = %p)\n", getpid(), count,
&count);

            printf("child sleeping ...\n");
            sleep(10); /* parent exits during this period, child became an orphan */
            printf("\nchild waking up!\n");
        }
        else { /* This is parent pro */
            printf("parent pro pid = %d, child pid = %d, count = %d (addr = %p)\n",
getpid(), childpid, count, &count);
        }
        printf("\nTesting point by %d\n", getpid()); /* executed by parent and child */
        return EXIT_SUCCESS;
}
```



■ Process Termination

■ Algorithm 6-3: fork-demo-nowait.c (fork without waiting).

```
iisscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-3-fork-demo-nowait.c
iisscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out
Parent pro pid = 15431, child pid = 15432, count = 1 (addr = 0x7ffe1a5ce6d0)

Testing point by 15431
child pro pid = 15432, count = 2 (addr = 0x7ffe1a5ce6d0)
child sleeping ...
iisscgy@ubuntu:/mnt/hgfs/os-2020$ ps -l
F S  UID      PID  PPID  C PRI  NI ADDR SZ WCHAN  TTY          TIME CMD
0 S  1000     1954   1944  0  80   0 -  6150 wait  pts/0        00:00:03 bash
1 S  1000    15432   1484  0  80   0 -  1128 hrtime pts/0        00:00:00 a.out
0 R  1000    15433   1954  0  80   0 -  7667 -    pts/0        00:00:00 ps
iisscgy@ubuntu:/mnt/hgfs/os-2020$
child waking up!

Testing point by 15432
```

```
        printf("parent pro pid = %d, child pid = %d, count = %d (addr = %p)\n",
getpid(), childpid, count, &count);
    }
    printf("\nTesting point by %d\n", getpid()); /* executed by parent and child */
    return EXIT_SUCCESS;
}
```




Process Termination

Algorithm 6-3: fork-demo-nowait.c (fork without waiting).

```

iisscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-3-fork-demo-nowait.c
iisscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out
Parent pro pid = 15431, child pid = 15432, count = 1 (addr = 0x7ffe1a5ce6d0)

Testing point by 15431
child pro pid = 15432, count = 2 (addr = 0x7ffe1a5ce6d0)
child sleeping ...
iisscgy@ubuntu:/mnt/hgfs/os-2020$ ps -l

```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	1000	1954	1944	0	80	0	-	6150	wait	pts/0	00:00:03	bash
1	S	1000	15432	1484	0	80	0	-	1128	hrttime	pts/0	00:00:00	a.out
0	R	1000	15433	1954	0	80	0	-	7667	-	pts/0	00:00:00	ps

```

iisscgy@ubuntu:/mnt
child waking up!

Testing point by 15432

```

The parent process terminated with an orphan of pid = 15432 left.

```

        printf("parent pro pid = %d, child pid = %d, count = %d (addr = %p)\n",
getpid(), childpid, count, &count);
    }
    printf("\nTesting point by %d\n", getpid()); /* executed by parent and child */
    return EXIT_SUCCESS;
}

```



Process Termination

Algorithm 6-3: fork-demo-nowait.c (fork without waiting).

```
iisscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-3-fork-demo-nowait.c
iisscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out
Parent pro pid = 15431, child pid = 15432, count = 1 (addr = 0x7ffe1a5ce6d0)

Testing point by 15431
child pro pid = 15432, count = 2 (addr = 0x7ffe1a5ce6d0)
child sleeping ...
iisscgy@ubuntu:/mnt/hgfs/os-2020$ ps -l
F S  UID      PID  PPID  C PRI  NI ADDR SZ WCHAN  TTY          TIME CMD
0 S  1000    1954   1944  0  80   0 -  6150 wait  pts/0      00:00:03 bash
1 S  1000    15432  1484  0  80   0 -  1128 hrtime pts/0      00:00:00 a.out
0 R  1000    15433  1954  0  80   0 -  7667 -    pts/0      00:00:00 ps
iisscgy@ubuntu:/mnt/hgfs/os-2020$
child waking up!

Testing point by 15432
```

What happens here? The terminal (bash) and the forked child are working asynchronously.

```
ld pid = %d, count = %d (addr = %p)\n",
```

```
    }
    printf("\nTesting point by %d\n", getpid()); /* executed by parent and child */
    return EXIT_SUCCESS;
}
```



■ Process Termination

■ Algorithm 6-4: fork-demo-wait.c (fork and wait).

```
int main(void)
{
    int count = 1;
    pid_t childpid, terminatedid;

    childpid = fork(); /* child duplicates parent's address space */
    if (childpid < 0) {
        perror("fork()");
        return EXIT_FAILURE;
    }
    else
        if (childpid == 0) { /* This is child pro */
            count++;
            printf("child pro pid = %d, count = %d (addr = %p)\n", getpid(), count,
&count);

            printf("child sleeping ...\n");
            sleep(5); /* parent wait() during this period */
            printf("\nchild waking up!\n");
        }
        else { /* This is parent pro */
            terminatedid = wait(0);
            printf("parent pro pid = %d, terminated pid = %d, count = %d (addr =
%p)\n", getpid(), terminatedid, count, &count);
        }
        printf("\nTesting point by %d\n", getpid()); /* executed by child and parent */
        return EXIT_SUCCESS;
}
```



■ Process Termination

■ Algorithm 6-4: fork-demo-wait.c (fork and wait).

```
iSScgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-4-fork-demo-wait.c
iSScgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out
child pro pid = 15444, count = 2 (addr = 0x7ffd425373bc)
child sleeping ...

child waking up!

Testing point by 15444
Parent pro pid = 15443, terminated pid = 15444, count = 1 (addr = 0x7ffd425373bc)

Testing point by 15443
iSScgy@ubuntu:/mnt/hgfs/os-2020$ ps
  PID TTY          TIME CMD
  1954 pts/0        00:00:03 bash
  15445 pts/0        00:00:00 ps
iSScgy@ubuntu:/mnt/hgfs/os-2020$
```

```
    else { /* This is parent pro */
        terminatedid = wait(0);
        printf("parent pro pid = %d, terminated pid = %d, count = %d (addr = %p)\n", getpid(), terminatedid, count, &count);
    }
    printf("\nTesting point by %d\n", getpid()); /* executed by child and parent */
    return EXIT_SUCCESS;
}
```



■ Process Termination

■ Algorithm 6-4: fork-demo-wait.c (fork and wait).

```
iisscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-4-fork-demo-wait.c
iisscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out
child pro pid = 15444, count = 2 (addr = 0x7ffd425373bc)
child sleeping ...

child waking up!

Testing point by 15444
Parent pro pid = 15443, terminated pid = 15444, count = 1 (addr = 0x7ffd425373bc)

Testing point by 15444
iisscgy@ubuntu:/mnt/hgfs/os-2020$ ps
  PID TTY          PPID  PTY  TIME COMMAND
 1954 pts/0        0    00:00:03 bash
 15445 pts/0        0    00:00:00 ps
iisscgy@ubuntu:/mnt/hgfs/os-2020$
```

The parent process is waiting until child process terminated.

```
    else { /* This is parent pro */
        terminatedid = wait(0);
        printf("parent pro pid = %d, terminated pid = %d, count = %d (addr = %p)\n", getpid(), terminatedid, count, &count);
    }
    printf("\nTesting point by %d\n", getpid()); /* executed by child and parent */
    return EXIT_SUCCESS;
}
```




■ Process Termination

- Algorithm 6-4: fork-demo-wait.c (fork and wait).

```
isscg@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-4-fork-demo-wait.c
isscg@ubuntu:/mnt/hgfs/os-2020$ ./a.out
child pro pid = 15444, count = 2 (addr = 0x7ffd425373bc)
child sleeping ...

child waking up!

Testing point by 15444
Parent pro pid = 15443, terminated pid = 15444, count = 1 (addr = 0x7ffd425373bc)

Testing point by 15443
isscg@ubuntu:/mnt/hgfs/os-2020$ ps
  PID TTY          PPID  PID  TIME
 1954 pts/0        1954    0   00:00
 15445 pts/0        15445   0   00:00
```

The testing point achieved first by child and then by parent

```
    else { /* This is parent pro */
        terminatedid = wait(0);
        printf("parent pro pid = %d, terminated pid = %d, count = %d (addr = %p)\n", getpid(), terminatedid, count, &count);
    }
    printf("\nTesting point by %d\n", getpid()); /* executed by child and parent */
    return EXIT_SUCCESS;
}
```



■ Process Termination

- **Algorithm 6-5-0:** sleeper.c (a demo process sleeping for 5 seconds).

```
/* gcc -o alg.6-5-0-sleeper.o alg.6-5-0-sleeper.c */
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char* argv[])
{
    int secnd = 5;

    if (argc > 1) {
        secnd = atoi(argv[1]);
        if ( secnd <= 0 || secnd > 10)
            secnd = 5;
    }

    printf("\nsleeper pid = %d, ppid = %d\nsleeper is taking a nap for %d
seconds\n", getpid(), getppid(), secnd); /* ppid - its parent pro id */

    sleep(secnd);
    printf("\nsleeper wakes up and returns\n");

    return 0;
}
```



Process Termination

Algorithm 6-5: vfork-execv-wait.c (vfork, execv and wait) (1)

```
int main(void)
{
    pid_t childpid;

    childpid = vfork();
    /* child shares parent's address space */
    if (childpid < 0) {
        perror("fork()");
        return EXIT_FAILURE;
    }
    else
        if(childpid == 0) { /* This is child pro */
            printf("This is child, pid = %d, taking a nap for 2 sencods \n", getpid());
            sleep(2); /* parent hung up and do nothing */

            char filename[80];
            struct stat buf;
            strcpy(filename, "./alg.6-5-0-sleeper.o");
            if(stat(filename, &buf) == -1) {
                perror("\nsleeper stat()");
                _exit(0);
            }
            char *argv1[] = {filename, argv[1], NULL};
            printf("child waking up and again execv() a sleeper: %s %s\n\n", argv1[0],
argv1[1]);
            execv(filename, argv); /* parent resume at the point 'execv' called */
        }
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/types.h>
#include <unistd.h>
#include <sys/stat.h>
#include <wait.h>
```




■ Process Termination

■ Algorithm 6-5: vfork-execv-wait.c (vfork, execv and wait) (2)

```
        else { /* This is parent pro, start when the vforked child terminated */
            printf("This is parent, pid = %d, childpid = %d \n", getpid(), childpid);
/* parent executed this statement during the EXECV time */
            int retpid = wait(0); /* without wait(), the spawned EXECV may became an
orphan */
            printf("\nwait() returns childpid = %d\n", retpid);
        }

    return EXIT_SUCCESS;
}
```



■ Process Termination

- Algorithm 6-5: vfork-execv-wait.c (vfork, execv and wait) (2)

```
iisscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-5-vfork-execv-wait.c
iisscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out 6
This is child, pid = 15477, taking a nap for 2 seconds ...
child waking up and again execv() a sleeper: ./alg.6-5-0-sleeper.o 6

This is parent, pid = 15476, childpid = 15477

sleeper pid = 15477, ppid = 15476
sleeper is taking a nap for 6 seconds

sleeper wakes up and returns

wait() returns childpid = 15477
iisscgy@ubuntu:/mnt/hgfs/os-2020$
```



■ Process Termination

■ Algorithm 6-5: vfork-execv-wait.c (vfork, execv and wait) (2)

```
iisscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-5-vfork-execv-wait.c
iisscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out 6
This is child, pid = 15477, taking a nap for 2 seconds ...
child waking up and again execv() a sleeper: ./alg.6-5-0-sleeper.o 6

This is parent, pid = 15476, childpid = 15477

sleeper pid = 15477, ppid = 15476
sleeper is taking a nap for 6 seconds

sleeper wakes up and returns

wait() returns childpid = 15477
iisscgy@ubuntu:/mnt/hgfs/os-2020$
```

The sleeper inherits the pid (15477) of the vforked child



■ Process Termination

- Algorithm 6-5: vfork-execv-wait.c (vfork, execv and wait) (2)

```
i SSCgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-5-vfork-execv-wait.c
i SSCgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out 6
This is child, pid = 15477, taking a nap for 2 seconds ...
child waking up and again execv() a sleeper: ./alg.6-5-0-sleeper.o 6

This is parent, pid = 15476, childpid = 15477

sleeper pid = 15477, ppid = 15476
sleeper is taking a nap for 6 sec

sleeper wakes up and returns

wait() returns childpid = 15477
i SSCgy@ubuntu:/mnt/hgfs/os-2020$
```

parent pro resumed at the point 'execv' called where vforked pro terminated and sleeper spawned as child in the same childpid but with duplicated address space and returned to parent without any stack smashing. parent and child executed asynchronously.



■ Process Termination

- Algorithm 6-5: vfork-execv-wait.c (vfork, execv and wait) (2)

```
iisscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-5-vfork-execv-wait.c
iisscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out 6
This is child, pid = 15477, taking a nap for 2 seconds ...
child waking up and again execv() a sleeper: ./alg.6-5-0-sleeper.o 6

This is parent, pid = 15476, childpid = 15477

sleeper pid = 15477, ppid = 15476
sleeper is taking a nap for 6 seconds

sleeper wakes up and returns

wait() returns childpid = 15477
iisscgy@ubuntu:/mnt/hgfs/os-2020$
```

Any way parent needs to wait()
his children, or the spawned
sleeper pro may become an
orphan



■ Process Termination

■ Algorithm 6-6: vfork-execv-nowait.c (vfork, execv without waiting) (1)

```
int main(int argc, char* argv[])
{
    pid_t childpid;

    childpid = vfork();
    /* child shares parent's address space */
    if (childpid < 0) {
        perror("fork()");
        return EXIT_FAILURE;
    }
    else
        if (childpid == 0) { /* This is child pro */
            printf("This is child, pid = %d, taking a nap for 2 seconds ... \n",
getpid());
            sleep(2); /* parent hung up and do nothing */

            char filename[80];
            struct stat buf;
            strcpy(filename, "./alg.6-5-0-sleeper.o");
            if (stat(filename, &buf) == -1) {
                perror("\nsleeper stat()");
                _exit(0);
            }
            char *argv1[] = {filename, argv[1], NULL};
            printf("child waking up and again execv() a sleeper: %s %s\n\n", argv1[0],
argv1[1]);
            execv(filename, argv); /* parent resume at the point 'execv' called */
        }
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/types.h>
#include <unistd.h>
#include <sys/stat.h>
#include <wait.h>
```



■ Process Termination

■ Algorithm 6-6: vfork-execv-nowait.c (vfork, execv without waiting) (2)

```
else { /* This is parent pro, start when the vforked child terminated */
    printf("This is parent, pid = %d, childpid = %d \n",getpid(), childpid);
    /* parent executed this statement during the EXECV time */
    printf("parent calling shell ps\n");
    system("ps -l");
    sleep(1);
    return EXIT_SUCCESS;
    /* parent exits without wait() and child may become an orphan */
}
}
```



Process Termination

Algorithm 6-6: vfork-execv-nowait.c (vfork, execv without waiting) (2)

```

isscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-6-vfork-execv-nowait.c
isscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out 6
This is child, pid = 17159, taking a nap for 2 seconds ...
child waking up and again execv() a sleeper: ./alg.6-5-0-sleeper.o 6

This is parent, pid = 17158, childpid = 17159
parent calling shell ps

sleeper pid = 17159, ppid = 17158
sleeper is taking a nap for 6 seconds
F S    UID      PID      PPID    C PRI  NI ADDR SZ WCHAN  TTY          TIME CMD
0 S    1000     1954     1944    0  80   0  -  6150 wait  pts/0        00:00:03 bash
0 S    1000    17158     1954    0  80   0  -  1128 wait  pts/0        00:00:00 a.out
0 S    1000    17159    17158    0  80   0  -  1128 hrtime pts/0        00:00:00 alg.6-5-0-sleep
0 S    1000    17169    17158    0  80   0  -  1158 wait  pts/0        00:00:00 sh
0 R    1000    17170    17169    0  80   0  -  7667 -      pts/0        00:00:00 ps

isscgy@ubuntu:/mnt/hgfs/os-2020$ ps -l
F S    UID      PID      PPID    C PRI  NI ADDR SZ WCHAN  TTY          TIME CMD
0 S    1000     1954     1944    0  80   0  -  6150 wait  pts/0        00:00:03 bash
0 S    1000    17159    1484    0  80   0  -  1128 hrtime pts/0        00:00:00 alg.6-5-0-sleep
0 R    1000    17171     1954    0  80   0  -  7667 -      pts/0        00:00:00 ps

isscgy@ubuntu:/mnt/hgfs/os-2020$
sleeper wakes up and returns
ps -q 1484
PID TTY          TIME CMD
1484 ?           00:00:00 systemd
isscgy@ubuntu:/mnt/hgfs/os-2020$

```




Process Termination

Algorithm 6-6: vfork-execv-nowait.c (vfork, execv without waiting) (2)

```
isscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-6-vfork-execv-nowait.c
isscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out 6
This is child, pid = 17159, taking a nap for 2 seconds ...
child waking up and again execv() a sleeper: ./alg.6-5-0-sleeper.o 6
```

```
This is parent, pid = 17158, childpid = 17159
parent calling shell ps
```

bash is the parent pro of
start main()

```
sleeper pid = 17159, ppid = 17158
sleeper is taking a nap for 6 seconds
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	1000	1954	1944	0	80	0	-	6150	wait	pts/0	00:00:03	bash
0	S	1000	17158	1954	0	80	0	-	1128	wait	pts/0	00:00:00	a.out
0	S	1000	17159	17158	0	80	0	-	1128	hrttime	pts/0	00:00:00	alg.6-5-0-sleep
0	S	1000	17169	17158	0	80	0	-	1158	wait	pts/0	00:00:00	sh
0	R	1000	17170	17169	0	80	0	-	7667	-	pts/0	00:00:00	ps

```
isscgy@ubuntu:/mnt/hgfs/os-2020$ ps -l
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	1000	1954	1944	0	80	0	-	6150	wait	pts/0	00:00:03	bash
0	S	1000	17159	1484	0	80	0	-	1128	hrttime	pts/0	00:00:00	alg.6-5-0-sleep
0	R	1000	17171	1954	0	80	0	-	7667	-	pts/0	00:00:00	ps

```
isscgy@ubuntu:/mnt/hgfs/os-2020$
```

```
sleeper wakes up and returns
```

```
ps -q 1484
```

PID	TTY	TIME	CMD
1484	?	00:00:00	systemd

```
isscgy@ubuntu:/mnt/hgfs/os-2020$
```



Process Termination

Algorithm 6-6: vfork-execv-nowait.c (vfork, execv without waiting) (2)

```
isscg@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-6-vfork-execv-nowait.c
```

```
isscg@ubuntu:/mnt/hgfs/os-2020$ ./a.out 6
```

```
This is child, pid = 17159, taking a nap for 2 seconds
```

```
child waking up and again execv() a sleeper: ./a
```

```
This is parent, pid = 17158, childpid = 17159  
parent calling shell ps
```

```
sleeper pid = 17159, ppid = 17158
```

```
sleeper is taking a nap for 6 seconds
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	1000	1954	1944	0	80	0	-	6150	wait	pts/0	00:00:03	bash
0	S	1000	17158	1954	0	80	0	-	1128	wait	pts/0	00:00:00	a.out
0	S	1000	17159	17158	0	80	0	-	1128	hrttime	pts/0	00:00:00	alg.6-5-0-sleep
0	S	1000	17169	17158	0	80	0	-	1158	wait	pts/0	00:00:00	sh
0	R	1000	17170	17169	0	80	0	-	7667	-	pts/0	00:00:00	ps

```
isscg@ubuntu:/mnt/hgfs/os-2020$ ps -l
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	1000	1954	1944	0	80	0	-	6150	wait	pts/0	00:00:03	bash
0	S	1000	17159	1484	0	80	0	-	1128	hrttime	pts/0	00:00:00	alg.6-5-0-sleep
0	R	1000	17171	1954	0	80	0	-	7667	-	pts/0	00:00:00	ps

```
isscg@ubuntu:/mnt/hgfs/os-2020$
```

```
sleeper wakes up and returns
```

```
ps -q 1484
```

PID	TTY	TIME	CMD
1484	?	00:00:00	systemd

```
isscg@ubuntu:/mnt/hgfs/os-2020$
```

The start main() is the parent pro of sleeper who inherits the vorked child's pid from execv()



Process Termination

Algorithm 6-6: vfork-execv-nowait.c (vfork, execv without waiting) (2)

```
isscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-6-vfork-execv-nowait.c
isscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out 6
This is child, pid = 17159, taking a nap for 2 seconds ...
child waking up and again execv() a sleeper: ./alg.6-5-0-sleeper.o 6
```

```
This is parent, pid = 17158, childpid = 17159
parent calling shell ps
```

The start main() is the parent pro of system()

```
sleeper pid = 17159, ppid = 17158
sleeper is taking a nap for 6 seconds
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	1000	1954	1944	0	80	0	-	6150	wait	pts/0	00:00:03	bash
0	S	1000	17158	1954	0	80	0	-	1128	wait	pts/0	00:00:00	a.out
0	S	1000	17159	17158	0	80	0	-	1128	hrttime	pts/0	00:00:00	alg.6-5-0-sleep
0	S	1000	17169	17158	0	80	0	-	1158	wait	pts/0	00:00:00	sh
0	R	1000	17170	17169	0	80	0	-	7667	-	pts/0	00:00:00	ps

```
isscgy@ubuntu:/mnt/hgfs/os-2020$ ps -l
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	1000	1954	1944	0	80	0	-	6150	wait	pts/0	00:00:03	bash
0	S	1000	17159	1484	0	80	0	-	1128	hrttime	pts/0	00:00:00	alg.6-5-0-sleep
0	R	1000	17171	1954	0	80	0	-	7667	-	pts/0	00:00:00	ps

```
isscgy@ubuntu:/mnt/hgfs/os-2020$
```

```
sleeper wakes up and returns
```

```
ps -q 1484
```

PID	TTY	TIME	CMD
1484	?	00:00:00	systemd

```
isscgy@ubuntu:/mnt/hgfs/os-2020$
```



Process Termination

Algorithm 6-6: vfork-execv-nowait.c (vfork, execv without waiting) (2)

```
isscg@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-6-vfork-execv-nowait.c
isscg@ubuntu:/mnt/hgfs/os-2020$ ./a.out 6
This is child, pid = 17159, taking a nap for 2 seconds ...
child waking up and again execv() a sleeper: ./alg.6-5-0-sleeper.o 6
```

```
This is parent, pid = 17158, childpid = 17159
parent calling shell ps
```

```
sleeper pid = 17159, ppid = 17158
sleeper is taking a nap for 6 seconds
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	1000	1954	1944	0	80	0	-	6150	wait	pts/0	00:00:03	bash
0	S	1000	17158	1954	0	80	0	-	1128	wait	pts/0	00:00:00	a.out
0	S	1000	17159	17158	0	80	0	-	1128	hrttime	pts/0	00:00:00	alg.6-5-0-sleep
0	S	1000	17169	17158	0	80	0	-	1158	wait	pts/0	00:00:00	sh
0	R	1000	17170	17169	0	80	0	-	7667	-	pts/0	00:00:00	ps

The sh is the parent pro of ps from system("ps -l")

```
isscg@ubuntu:/mnt/hgfs/os-2020$ ps -l
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	1000	1954	1944	0	80	0	-	6150	wait	pts/0	00:00:03	bash
0	S	1000	17159	1484	0	80	0	-	1128	hrttime	pts/0	00:00:00	alg.6-5-0-sleep
0	R	1000	17171	1954	0	80	0	-	7667	-	pts/0	00:00:00	ps

```
isscg@ubuntu:/mnt/hgfs/os-2020$
```

```
sleeper wakes up and returns
```

```
ps -q 1484
```

PID	TTY	TIME	CMD
1484	?	00:00:00	systemd

```
isscg@ubuntu:/mnt/hgfs/os-2020$
```




Process Termination

Algorithm 6-6: vfork-execv-nowait.c (vfork, execv without waiting) (2)

```
isscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-6-vfork-execv-nowait.c
isscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out 6
This is child, pid = 17159, taking a nap for 2 seconds ...
child waking up and again execv() a sleeper: ./alg.6-5-0-sleeper.o 6
```

```
This is parent, pid = 17158, childpid = 17159
parent calling shell ps
```

```
sleeper pid = 17159, ppid = 17158
sleeper is taking a nap for 6 seconds
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN
0	S	1000	1954	1944	0	80	0	-	6150	wait
0	S	1000	17158	1954	0	80	0	-	1128	wait
0	S	1000	17159	17158	0	80	0	-	1128	hrtim
0	S	1000	17169	17158	0	80	0	-	1158	wait
0	R	1000	17170	17169	0	80	0	-	7667	-

Start main() terminated and control back to bash, I type "ps -l" from terminal

```
isscgy@ubuntu:/mnt/hgfs/os-2020$ ps -l
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	1000	1954	1944	0	80	0	-	6150	wait	pts/0	00:00:03	bash
0	S	1000	17159	1484	0	80	0	-	1128	hrtim	pts/0	00:00:00	alg.6-5-0-sleep
0	R	1000	17171	1954	0	80	0	-	7667	-	pts/0	00:00:00	ps

```
isscgy@ubuntu:/mnt/hgfs/os-2020$
```

```
sleeper wakes up and returns
```

```
ps -q 1484
```

PID	TTY	TIME	CMD
1484	?	00:00:00	systemd

```
isscgy@ubuntu:/mnt/hgfs/os-2020$
```



Process Termination

Algorithm 6-6: vfork-execv-nowait.c (vfork, execv without waiting) (2)

```
isscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-6-vfork-execv-nowait.c
isscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out 6
This is child, pid = 17159, taking a nap for 2 seconds ...
child waking up and again execv() a sleeper: ./alg.6-5-0-sleeper.o 6
```

```
This is parent, pid = 17158, childpid = 17159
parent calling shell ps
```

```
sleeper pid = 17159, ppid = 17158
sleeper is taking a nap for 6 seconds
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	1000	1954	1944	0	80	0	-	6150	wait	pts/0	00:00:03	bash
0	S	1000	17158	1954	0	80	0	-	1128	wait	pts/0	00:00:00	a.out
0	S	1000	17159	17158	0	80	0	-	1128	hrttime	pts/0	00:00:00	alg.6-5-0-sleep
0	S	1000	17169	17158	0	80	0	-	1158	wait	pts/0	00:00:00	sh
0	R	1000	17170	17169	0	80	0	-	7667	-	pts/0	00:00:00	ps

```
isscgy@ubuntu:/mnt/hgfs/os-2020$ ps -l
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	1000	1954	1944	0	80	0	-	6150	wait	pts/0	00:00:03	bash
0	S	1000	17159	1484	0	80	0	-	1128	hrttime	pts/0	00:00:00	alg.6-5-0-sleep
0	R	1000	17171	1954	0	80	0	-	7667	-	pts/0	00:00:00	ps

```
isscgy@ubuntu:/mnt/hgfs/os-2020$
```

```
sleeper wakes up and returns
```

```
ps -q 1484
```

PID	TTY	TIME	CMD
1484	?	00:00:00	systemd

```
isscgy@ubuntu:/mnt/hgfs/os-2020$
```

It showed that sleeper missed his parent and adopted by 1484



■ Process Termination

■ Algorithm 6-6: vfork-execv-nowait.c (vfork, execv without waiting) (2)

```
isscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-6-vfork-execv-nowait.c
isscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out 6
This is child, pid = 17159, taking a nap for 2 seconds ...
child waking up and again execv() a sleeper: ./alg.6-5-0-sleeper.o 6

This is parent, pid = 17158, childpid = 17159
parent calling shell ps

sleeper pid = 17159, ppid = 17158
sleeper is taking a nap for 6 seconds
F S    UID    PID    PPID    C PRI  NI ADDR SZ WCHAN  TTY          TIME CMD
0 S    1000    1954    1944    0  80   0  -   6150 wait  pts/0        00:00:03 bash
0 S    1000    17158    1954    0  80   0  -   1128 wait  pts/0        00:00:00 a.out
0 S    1000    17159    17158    0  80   0  -   1128 hrtime pts/0        00:00:00 alg.6-5-0-sleep
0 S    1000    17169    17158    0  80   0  -   1158 wait  pts/0        00:00:00 sh
0 R    1000    17170    17169    0  80   0  -   7667 -      pts/0        00:00:00 ps

isscgy@ubuntu:/mnt/hgfs/os-2020$ ps -l
F S    UID    PID    PPID    C PRI  NI ADDR SZ WCHAN  TTY          TIME CMD
0 S    1000    1954    1944    0  80   0  -   6150 wait  pts/0        00:00:03 bash
0 S    1000    17159    1484    0  80   0  -   1128 hrtime pts/0        00:00:00 alg.6-5-0-sleep
0 R    1000    17171    1954    0  80   0  -   7667 -      pts/0        00:00:00 ps

isscgy@ubuntu:/mnt/hgfs/os-2020$
sleeper wakes up and returns
ps -q 1484
  PID TTY          TIME CMD
 1484 ?            00:00:00 systemd
isscgy@ubuntu:/mnt/hgfs/os-2020$
```

The terminal (bash) and the sleeper are working asynchronously.



Process Termination

Algorithm 6-6: vfork-execv-nowait.c (vfork, execv without waiting) (2)

```
isscgy@ubuntu:/mnt/hgfs/os-2020$ gcc alg.6-6-vfork-execv-nowait.c
isscgy@ubuntu:/mnt/hgfs/os-2020$ ./a.out 6
This is child, pid = 17159, taking a nap for 2 seconds ...
child waking up and again execv() a sleeper: ./alg.6-5-0-sleeper.o 6

This is parent, pid = 17158, childpid = 17159
parent calling shell ps

sleeper pid = 17159, ppid = 17158
sleeper is taking a nap for 6 seconds
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	1000	1954	1944	0	80	0	-	6150	wait	pts/0	00:00:03	bash
0	S	1000	17158	1954	0	80	0	-	1128	wait	pts/0	00:00:00	a.out
0	S	1000	17159	17158	0	80	0	-	1128	hrttime	pts/0	00:00:00	alg.6-5-0-sleep
0	S	1000	17169	17158	0	80	0	-	1158	wait	pts/0	00:00:00	sh
0	R	1000	17170	17169	0	80	0	-	7667	-	pts/0	00:00:00	ps

```
isscgy@ubuntu:/mnt/hgfs/os-2020$ ps -l
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	1000	1954	1944	0	80	0	-	6150	wait	pts/0	00:00:03	bash
0	S	1000	17159	1484	0	80	0	-	1128	hrttime	pts/0	00:00:00	alg.6-5-0-sleep
0	R	1000	17171	1954	0	80	0	-	7667	-	pts/0	00:00:00	ps

```
isscgy@ubuntu:/mnt/hgfs/os-2020$
sleeper wakes up and returns
ps -q 1484
```

PID	TTY	TIME	CMD
1484	?	00:00:00	systemd

```
isscgy@ubuntu:/mnt/hgfs/os-2020$
```

pid 1484 is a daemon of "systemd" (in place of "init")