
Introduction to Process

Operating Systems

School of Data & Computer Science
Sun Yat-sen University

Lecture Notes: os_sysu@163.com
Instructor: Guoyang Cai
email: isscgymail@mail.sysu.edu.cn





■ Contents

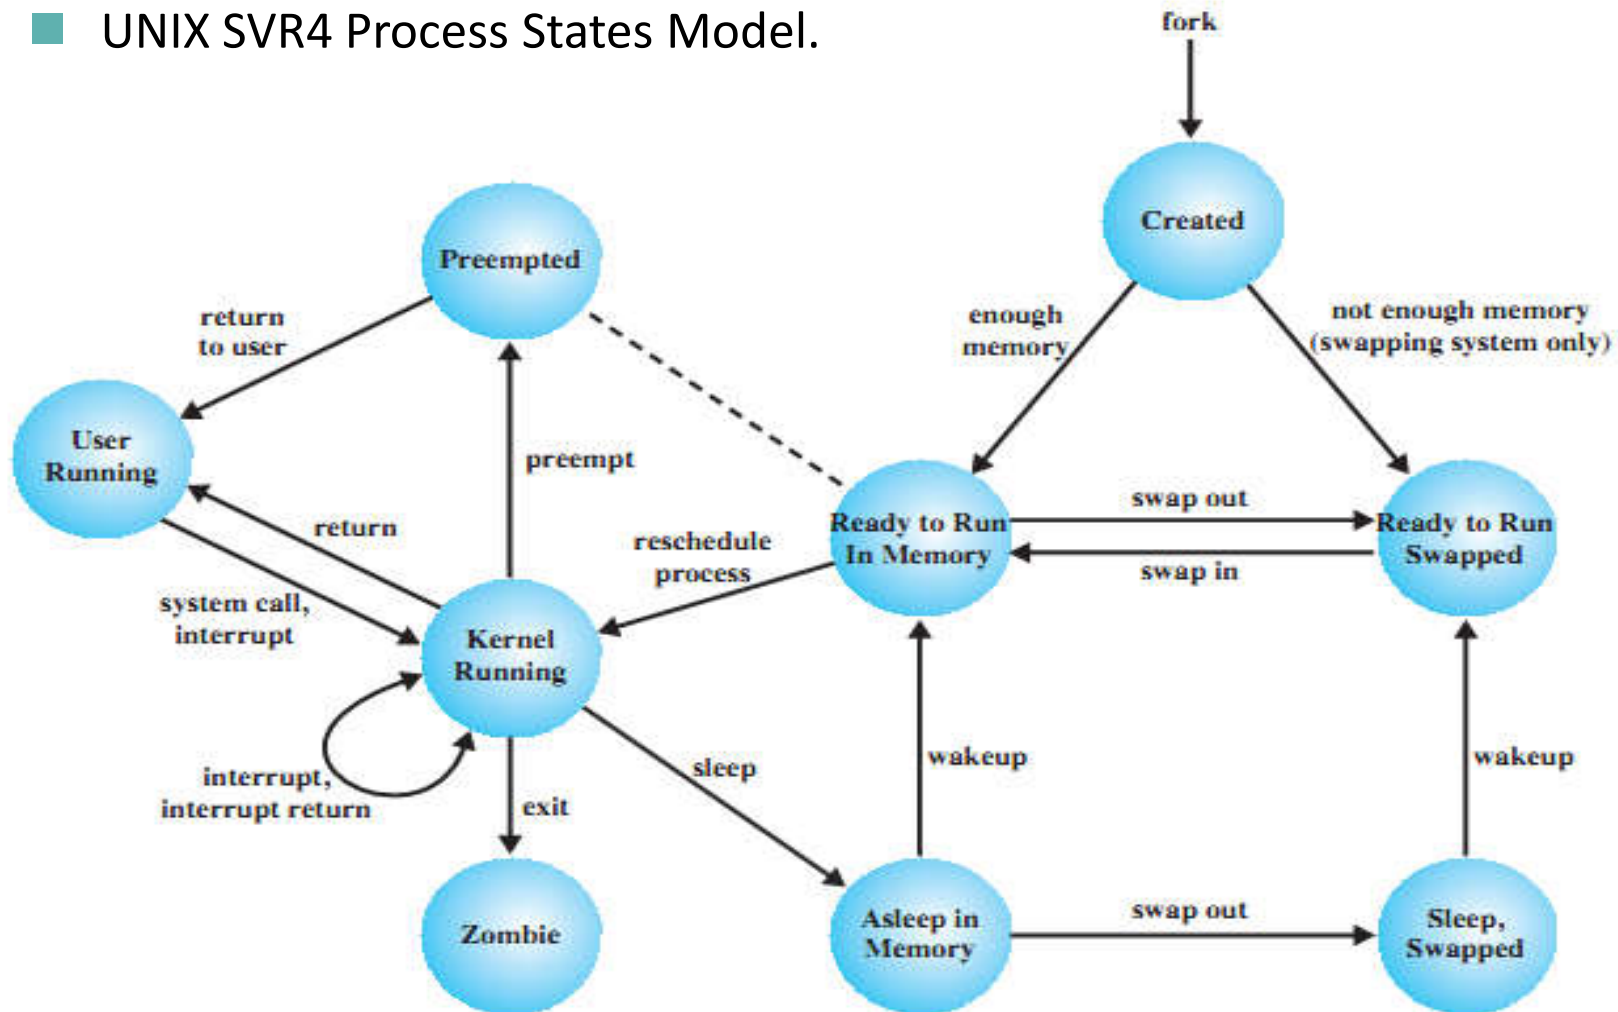
- Basic Concepts
- Process Table and Process Control Block
- Process States and Transitions
- Operations on Process
 - Process Creation
 - Process Termination
- Unix and Linux Examples
- Process Scheduling
- Process Switching

■ UNIX SVR4 Process States

- **User Running** - Executing in user mode.
- **Kernel Running** - Executing in kernel mode.
- **Ready to Run, in Memory** - Ready to run as soon as the kernel schedules it.
- **Asleep in Memory** - Unable to execute until an event occurs; process is in main memory (a blocked state).
- **Ready to Run, Swapped** - Process is ready to run, but the swapper must swap the process into main memory before the kernel can schedule it to execute.
- **Sleeping, Swapped** - The process is awaiting an event and has been swapped to secondary storage (a blocked state).
- **Preempted** - Process is returning from kernel to user mode, but the kernel preempts it and does a process switch to schedule another process.
- **Created** - Process is newly created and not yet ready to run.
- **Zombie** - Process no longer exists, but it leaves a record for its parent process to collect.

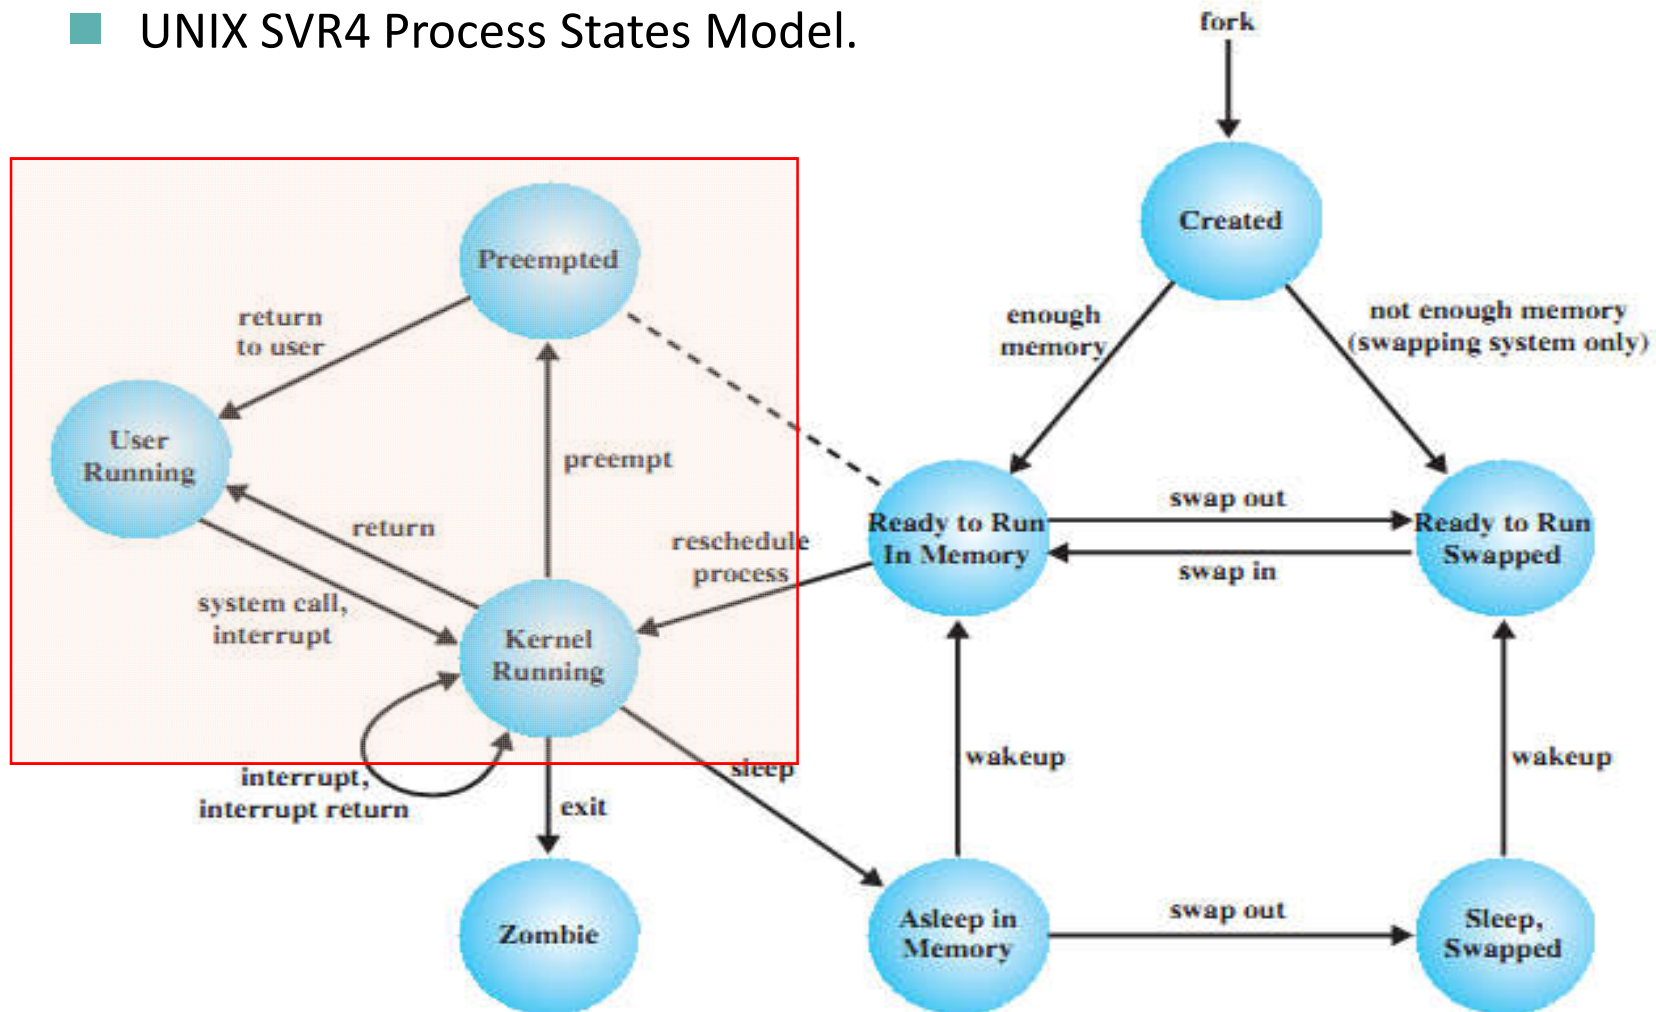
■ UNIX SVR4 Process States

- UNIX SVR4 Process States Model.



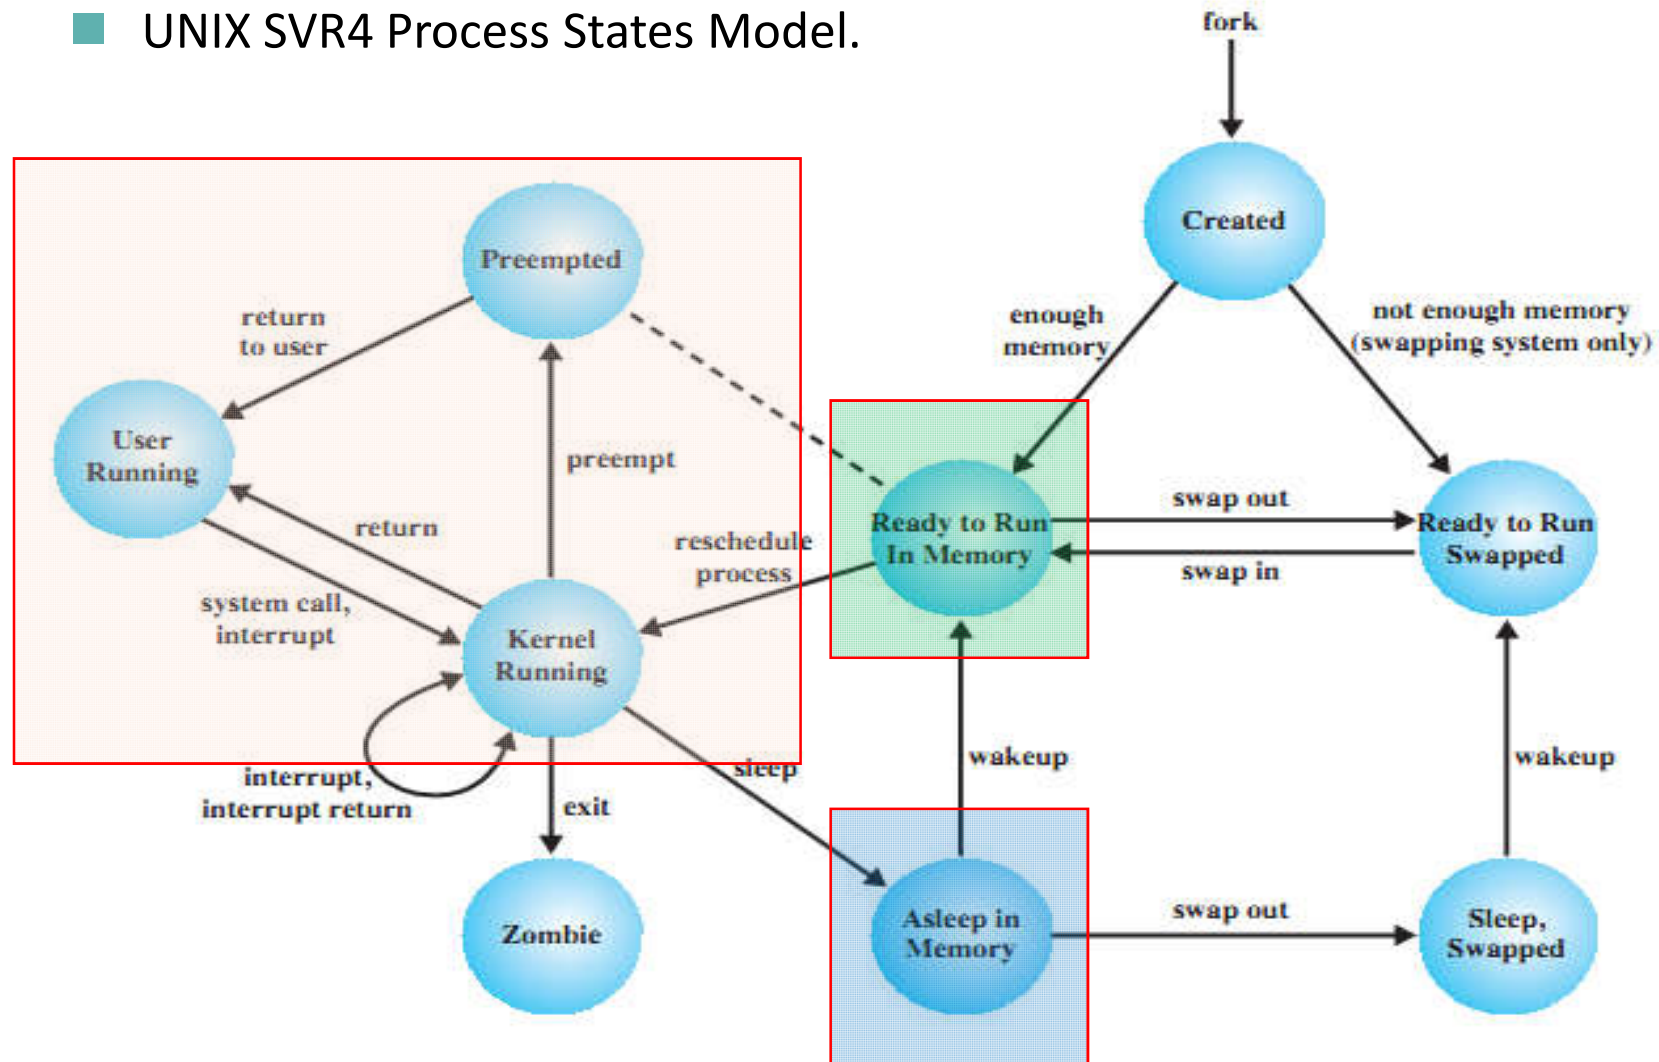
■ UNIX SVR4 Process States

- UNIX SVR4 Process States Model.



■ UNIX SVR4 Process States

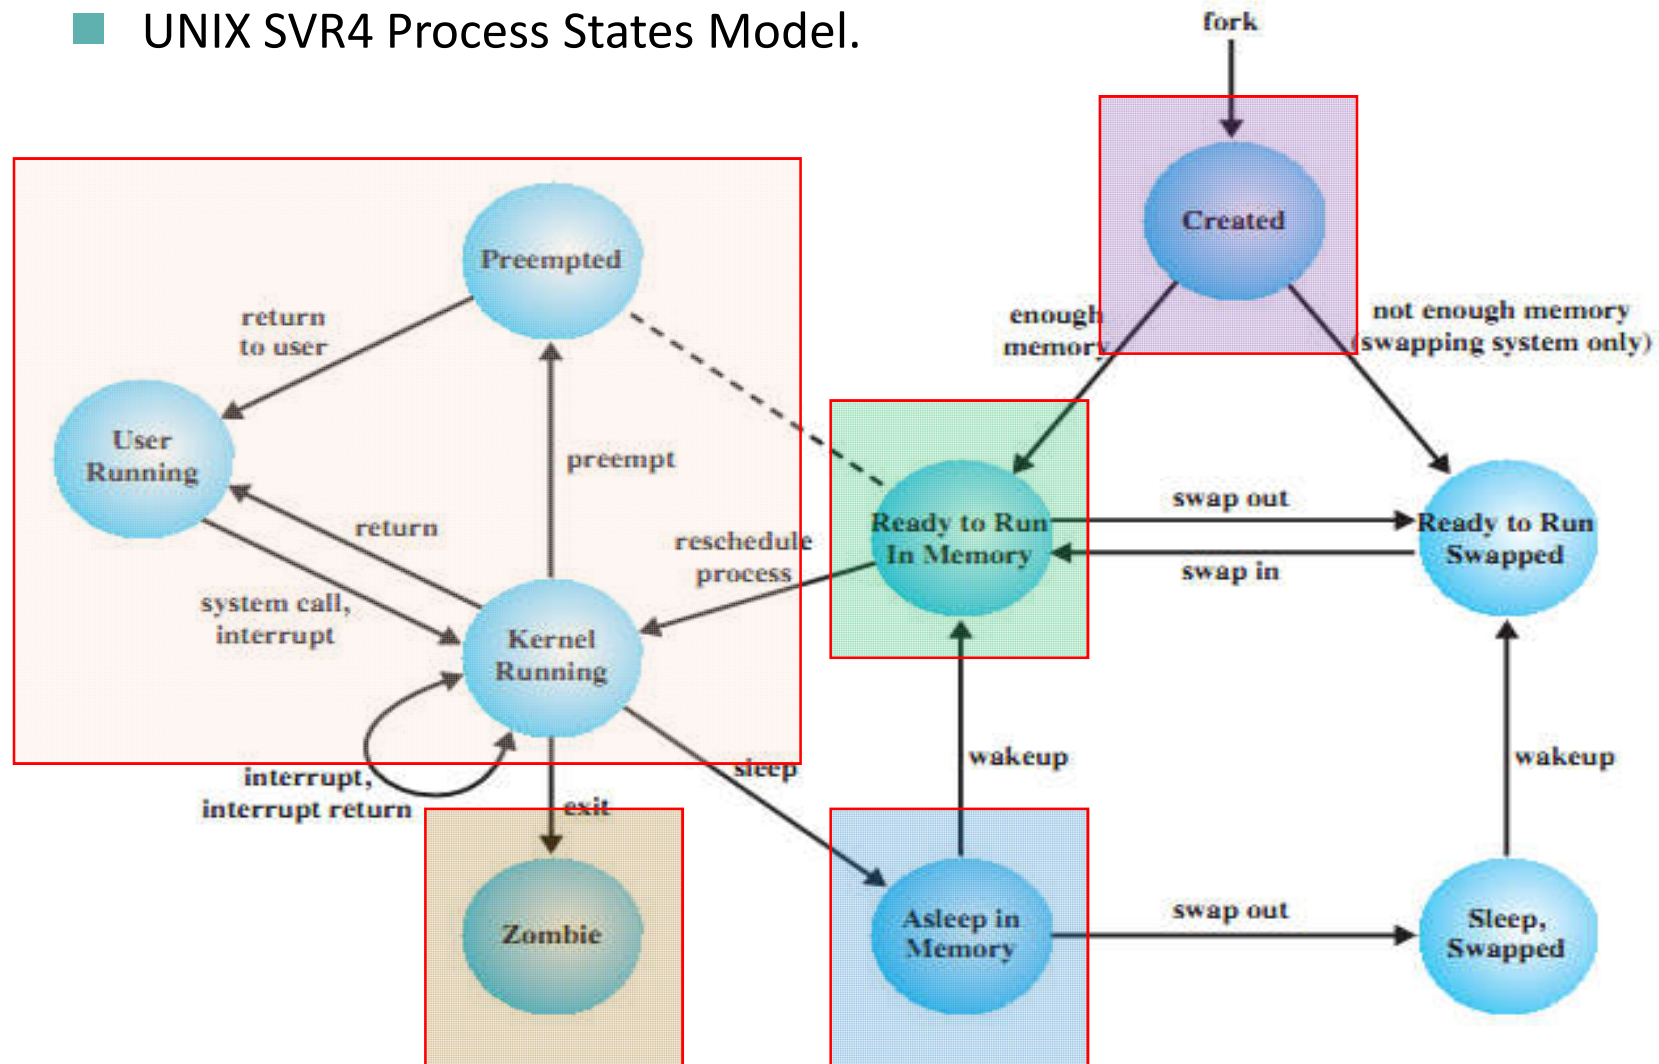
- UNIX SVR4 Process States Model.



Three-state Transition

■ UNIX SVR4 Process States

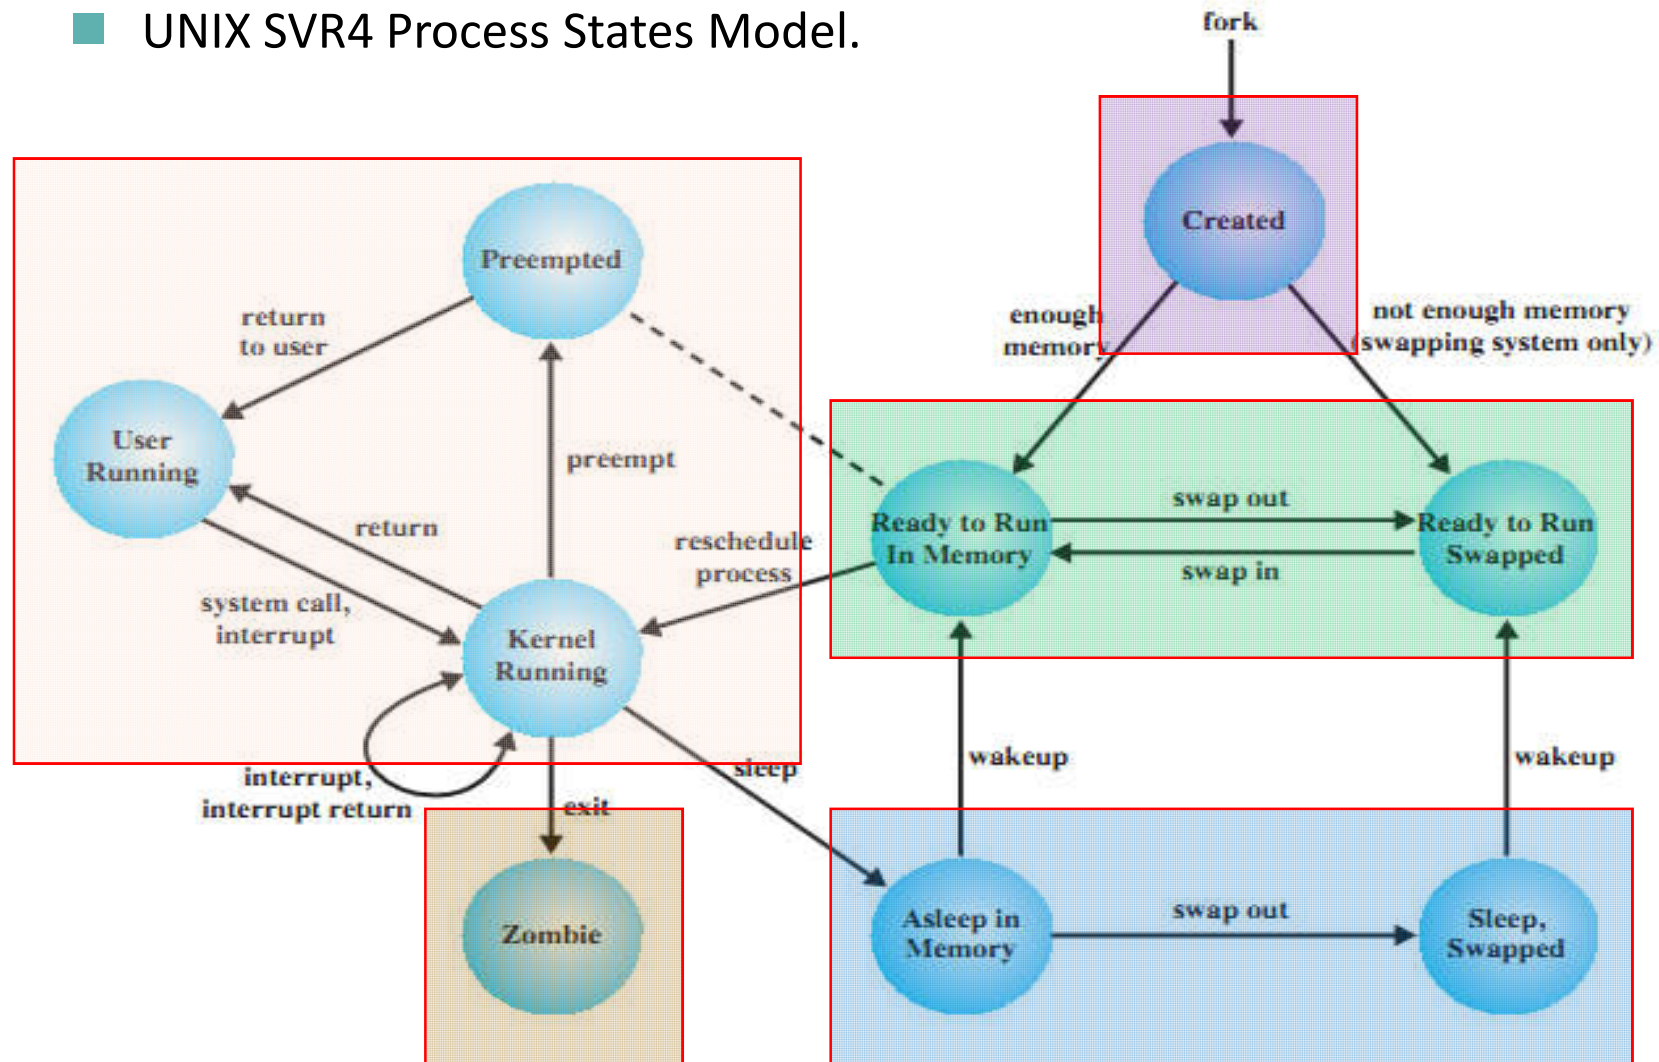
- UNIX SVR4 Process States Model.



Five-state Transition

■ UNIX SVR4 Process States

- UNIX SVR4 Process States Model.

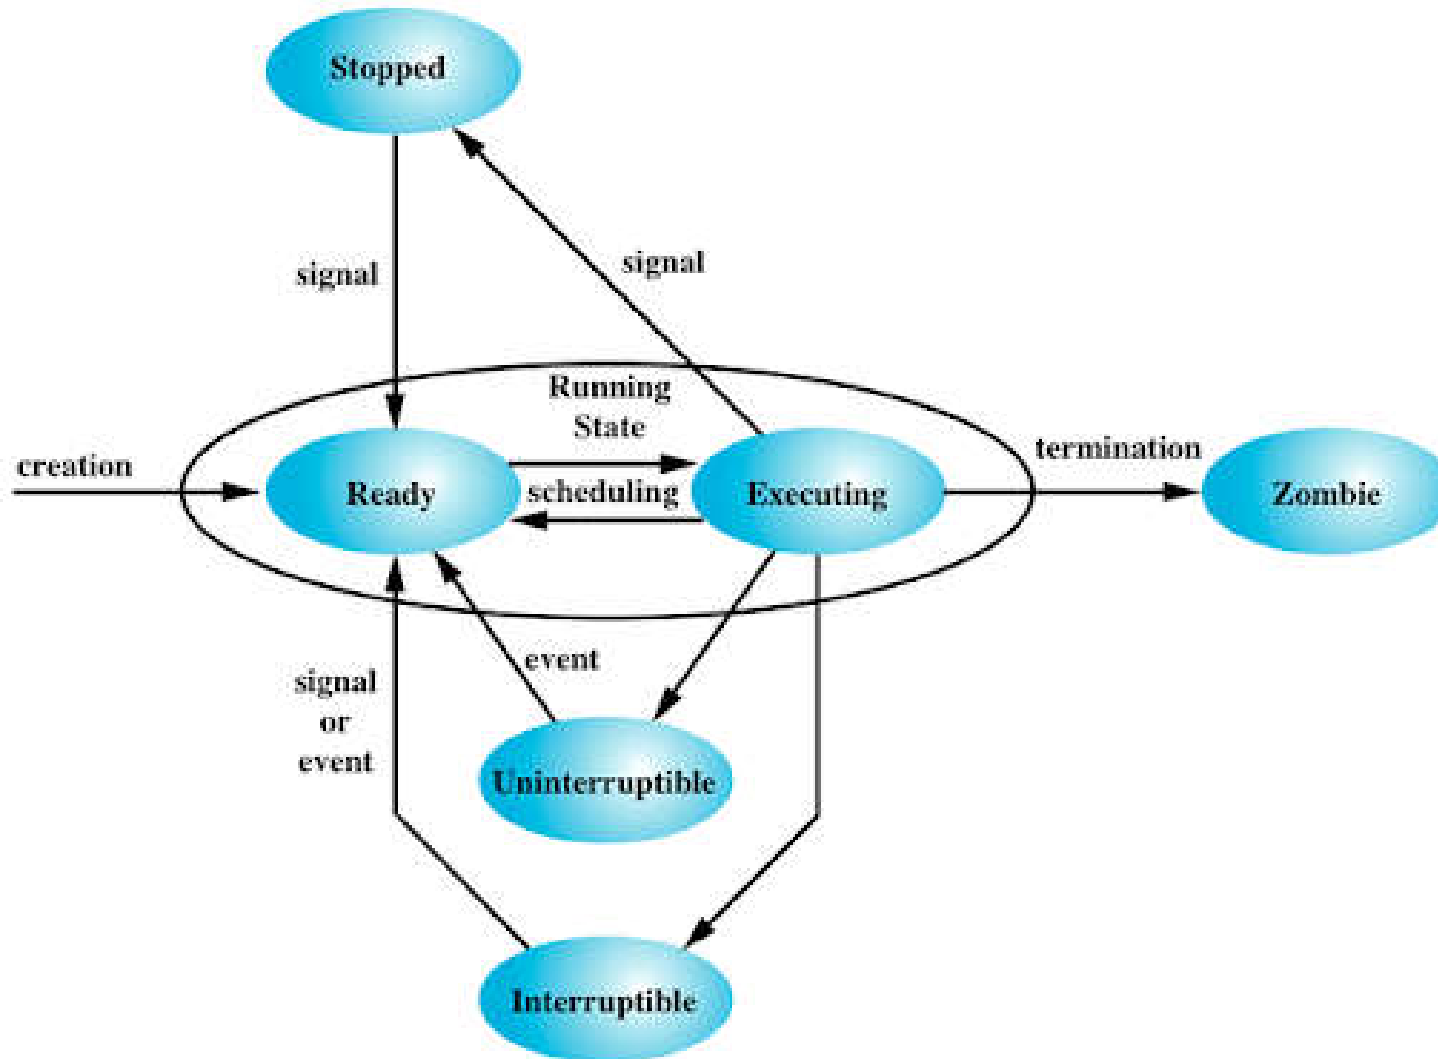


Five-state Transition with Swapping



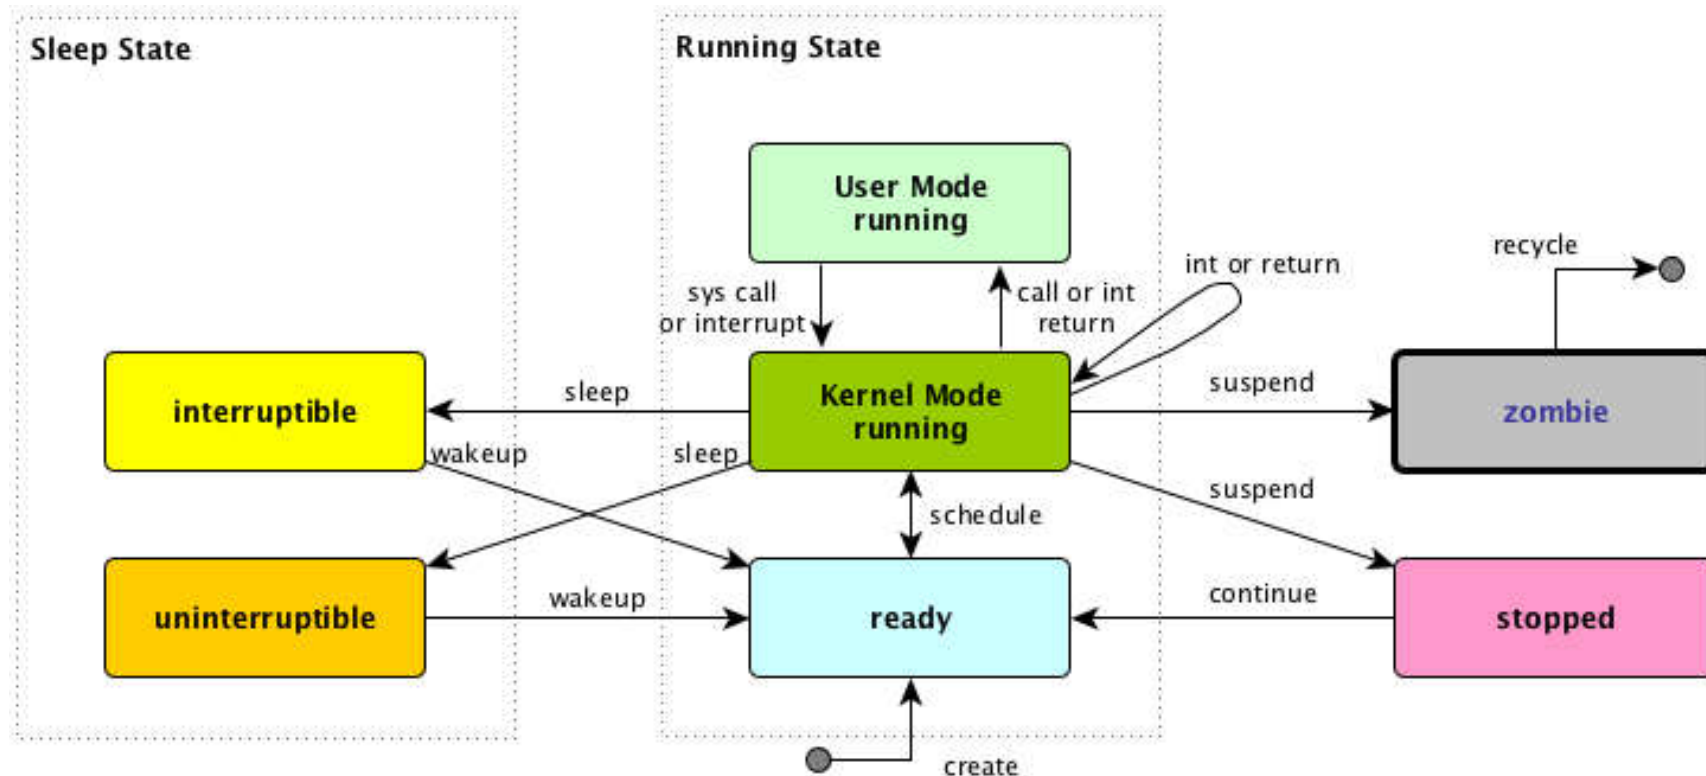
■ Process Representation in Linux

■ Linux Process State.



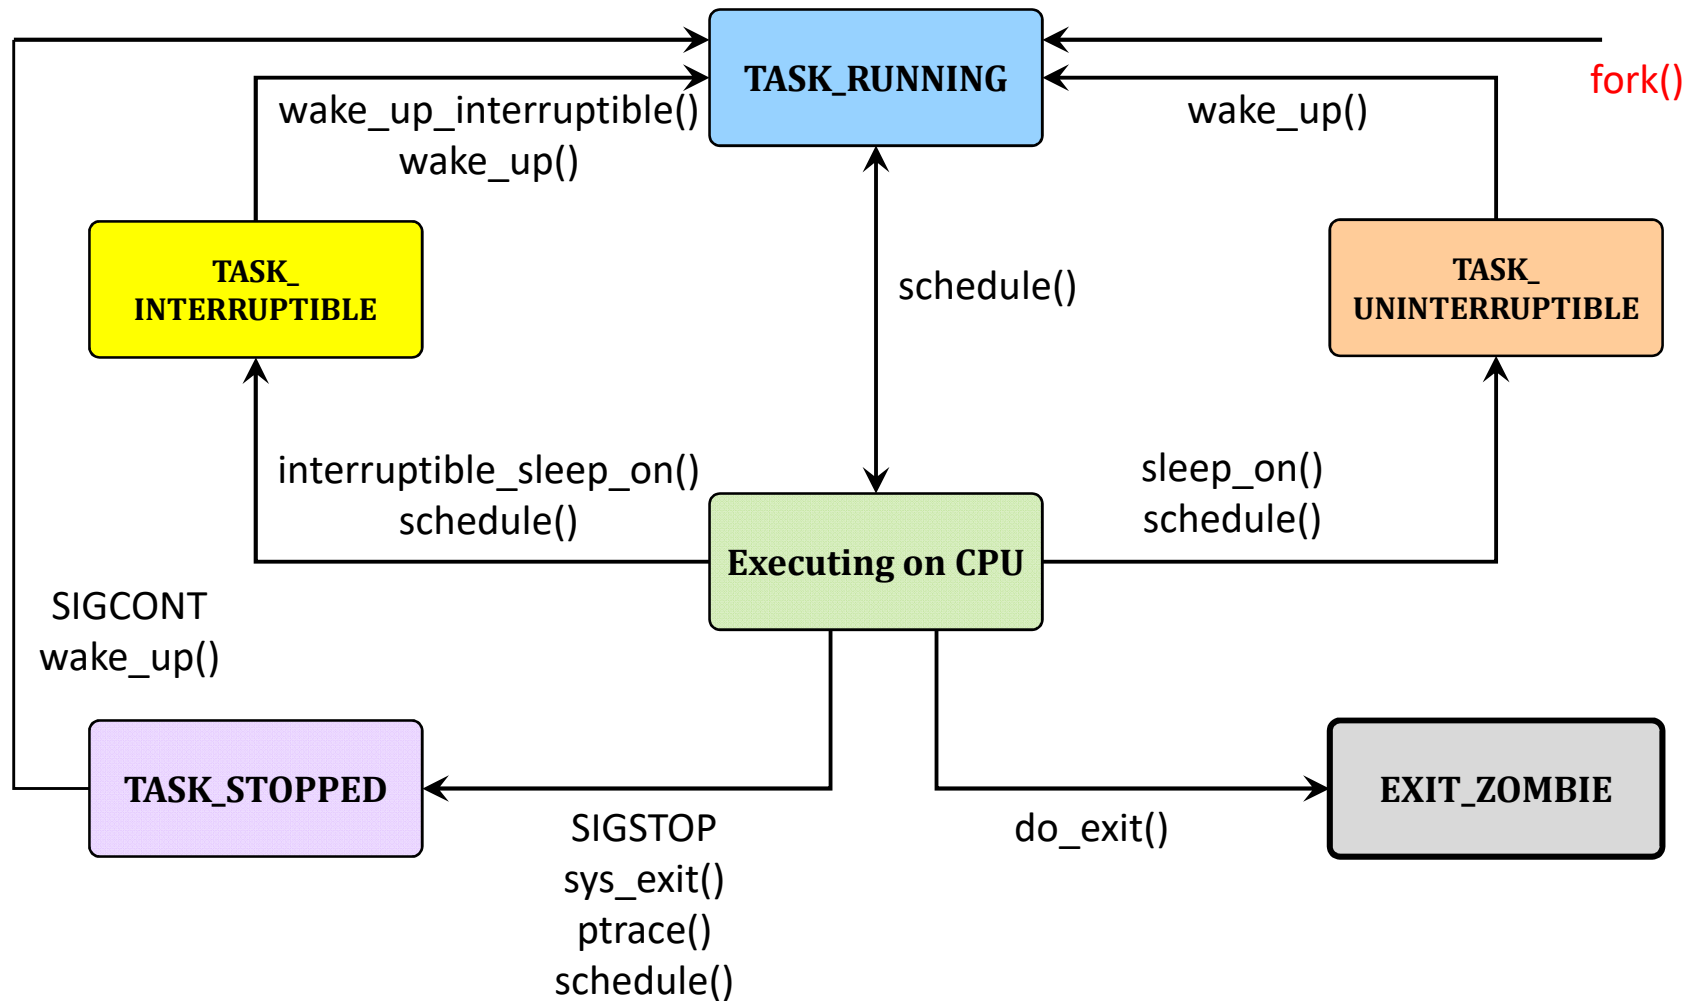
■ Process Representation in Linux

■ Linux Process State.



■ Process Representation in Linux

■ Linux Process State.



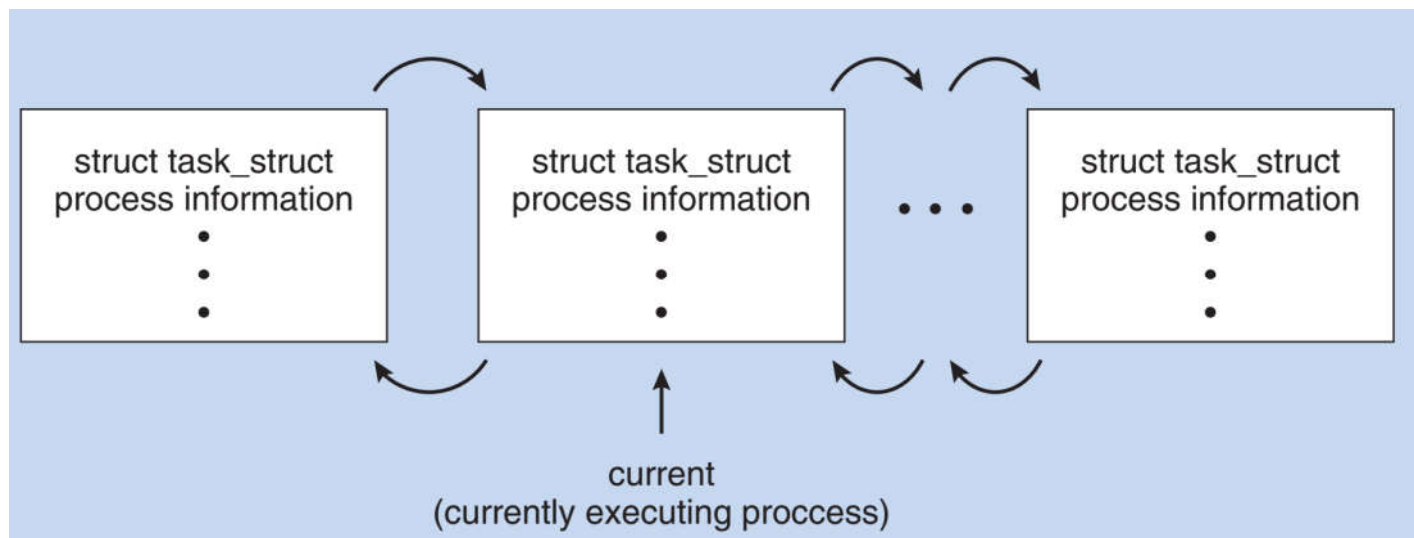
■ Process Representation in Linux

- The PCB in Linux is represented by the C structure `task_struct` included in `<linux/sched.h>`.

- `task_struct` is declared in `linux/sched.h`.

```
$ sudo apt-get install linux-source
$ tar ...
$ vim /usr/src/linux-source-X.XX.X/include/linux/sched.h
```

- The Linux kernel uses a *circular doubly-linked list* of `struct task_struct` to store these process descriptors.



■ Process Representation in Linux

- **task_struct** contains all the necessary information for representing a process, including
 - PID
 - the state of the process
 - processor registers
 - scheduling and memory-management information
 - list of open files
 - pointers to the process's parent and a list of its children and siblings.
- Some of these fields include:

```
long state; /* state of the process */
struct sched_entity se; /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```


■ Process Representation in Linux

- Here are a few fields from kernel 2.6.15-1.2054_FC5, starting at line 701: (what about these fields in your current Ubuntu version?)

```
701.struct task_struct {
702.    volatile long state; /*-1 unrunnable,0 runnable,>0 stopped */
703.    struct thread_info *thread_info;
704.    . . .
767.    /* PID/PID hash table linkage. */
768.    struct pid pids[PIDTYPE_MAX];
769.    . . .
798.    char comm[TASK_COMM_LEN]; /* executable name excluding path
```

- Type definition of volatile long.

```
#define TASK_RUNNING 0
#define TASK_INTERRUPTIBLE 1
#define TASK_UNINTERRUPTIBLE 2
#define TASK_STOPPED 4
#define TASK_TRACED 8 /* in tsk->exit_state */
#define EXIT_ZOMBIE 16
#define EXIT_DEAD 32 /* in tsk->state_again */
#define TASK_NONINTERACTIVE 64
```

■ Overview


- The objective of multiprogramming is to have some process running at all times, to maximize CPU utilization.
 - The objective of time sharing is to switch the CPU among processes so frequently that users can interact with each program while it is running.
- *Process scheduler* selects among available processes for next execution on CPU.
 - For a single-processor system, there will be only one running process and the rest will have to wait until the CPU is free and can be rescheduled.

■ Overview

- *Scheduling queues* of processes are to be maintained:
 - Job Queue
 - set of all processes in the system
 - Ready Queue
 - set of all processes residing in main memory, ready and waiting to execute.
 - Device/Waiting Queues
 - set of processes waiting for an I/O device
- Processes *migrate* among the various queues.

Overview

- The Ready and wait Queues.



■ Overview

- Most processes can be described as either I/O-bound or CPU-bound:
 - I/O-bound process
 - spends more time doing I/O than computations
 - short CPU bursts
 - CPU-bound process
 - spends more time doing computations
 - long CPU bursts

■ Types of Process Schedulers

- There are three types/levels of Process Schedulers
 - Long-term Scheduler
 - High-level Scheduler, or
 - Jobs Scheduler.
 - Medium-term Scheduler
 - Medium-level Scheduler,
 - Swapping Scheduler, or
 - Emergency Scheduler.
 - Short-term Scheduler
 - Low-level Scheduler,
 - CPU Scheduler,
 - Micro Scheduler, or
 - Process/Thread Scheduler on narrow sense (狭义上的进程/线程调度).

■ Long-term Schedulers

- Long-term process scheduler selects which programs/processes should be brought into the *ready queue*.
 - determines which **programs** are admitted to the system for processing
 - controls the degree of multiprogramming
 - strives for good process mix of I/O-bound and CPU-bound processes
 - Long-term scheduler is invoked infrequently.
 - seconds, minutes
 - may be slow
- If more processes are admitted for processing:
 - less likely that all processes will be blocked
 - bringing better CPU usage
 - each process has less fraction of the CPU time

■ Short-term Schedulers

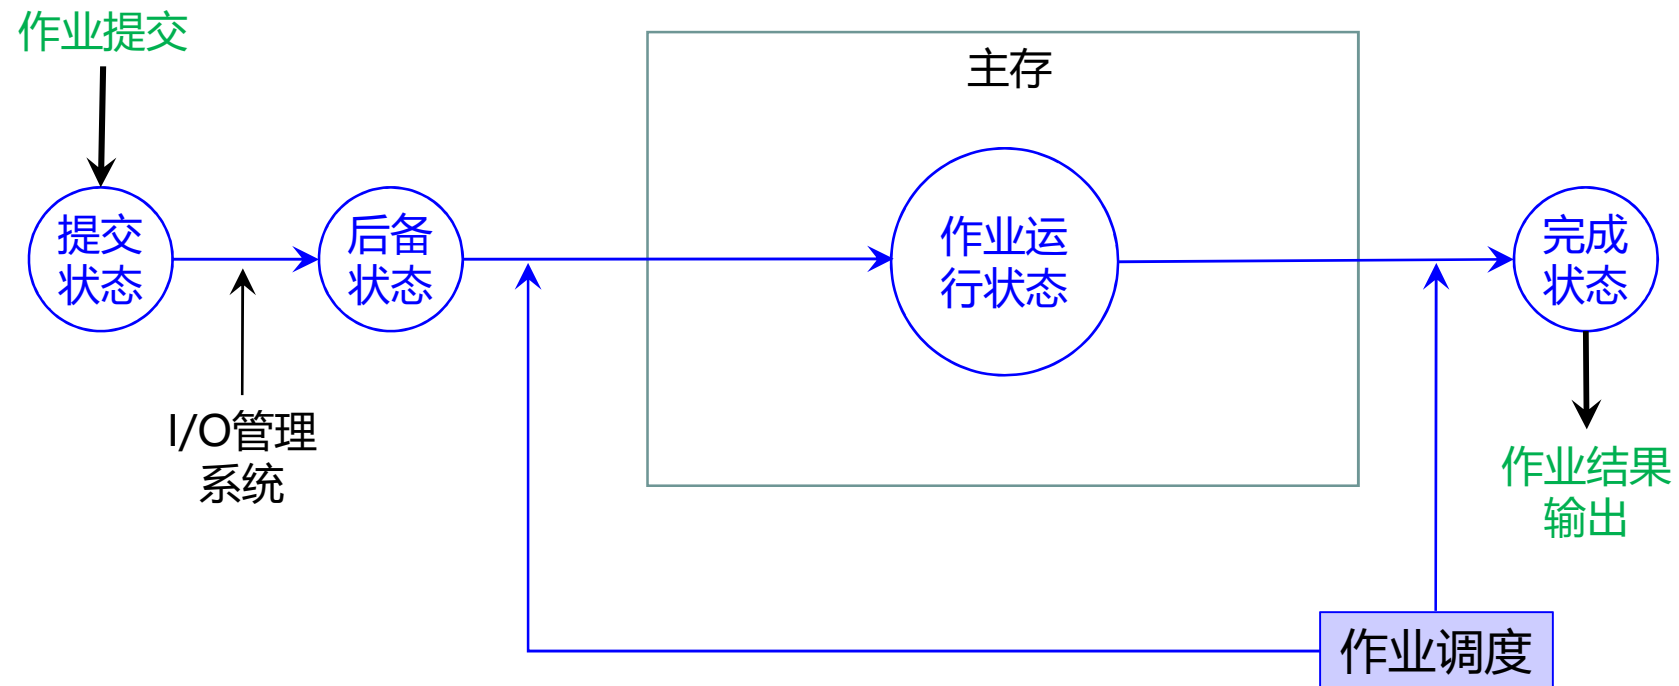
- Short-term process scheduler selects which process should be *executed* next and allocates CPU — also called CPU scheduling (处理机调度).
 - CPU scheduling determines which *process* is going to execute next according to a scheduling algorithm.
 - Short-term scheduler is also known as the *dispatcher* (which is part of it) that moves the processor from one process to another, and prevents a single process from monopolizing (独占) processor time.
 - Short-term scheduler is invoked on an event that may lead to choose another process for execution:
 - Clock interrupts
 - I/O interrupts
 - Operating system calls and traps
 - Signals.
- Short-term scheduler is invoked very frequently
 - milliseconds
 - must be fast.

■ Medium-term Schedulers

- Medium-term process scheduler selects which job/process should be *swapped* out if system is overloaded.
 - So far, all processes have to be (at least partly) in main memory.
 - Even with virtual memory, keeping too many processes in main memory will deteriorate the system's performance.
 - OS may need to swap out some processes *to disk*, and then later swap them back in.
 - Swapping decisions is based on the need of multiprogramming management.

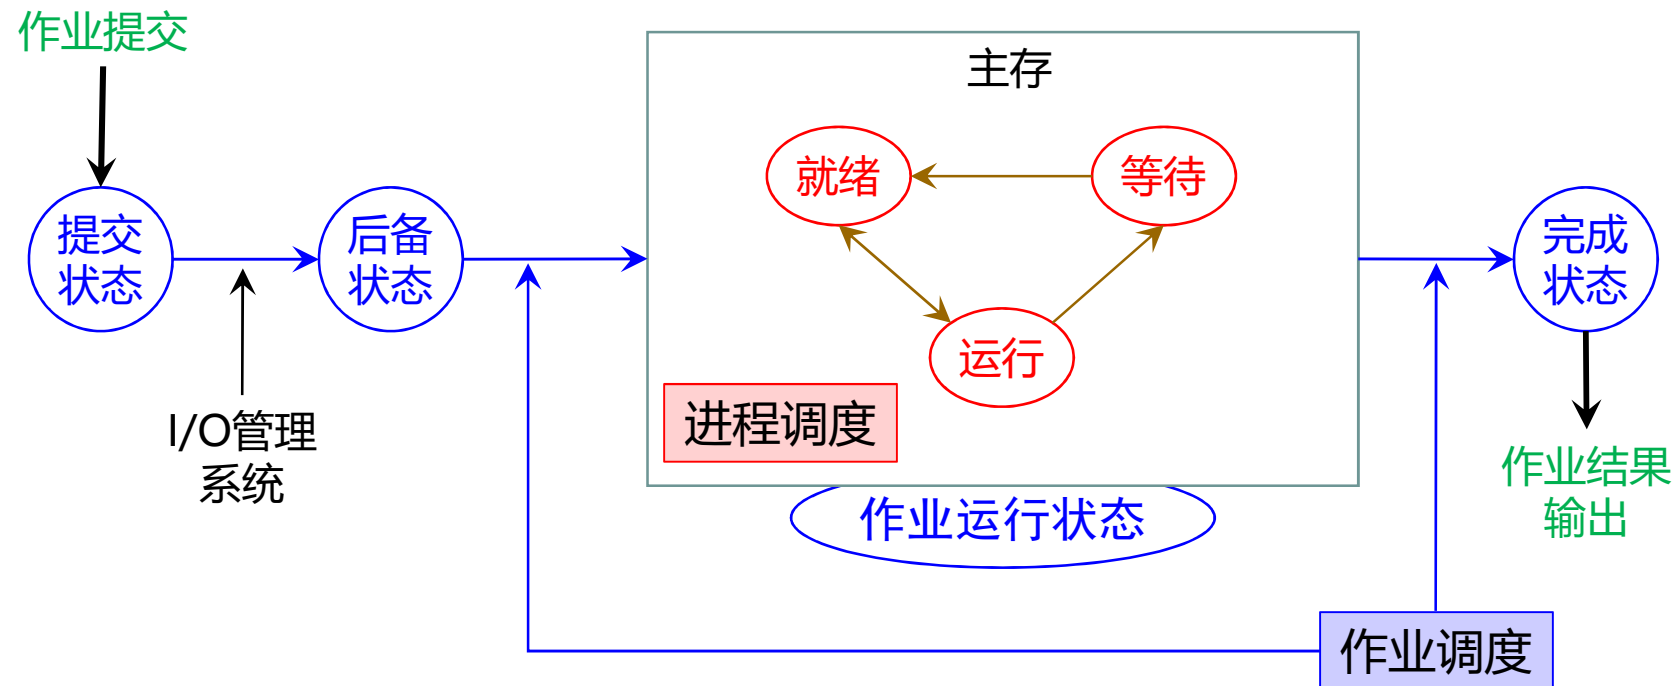
■ Schematic View of Schedulers

- Job scheduling, Swapping scheduling and Process/Thread scheduling.



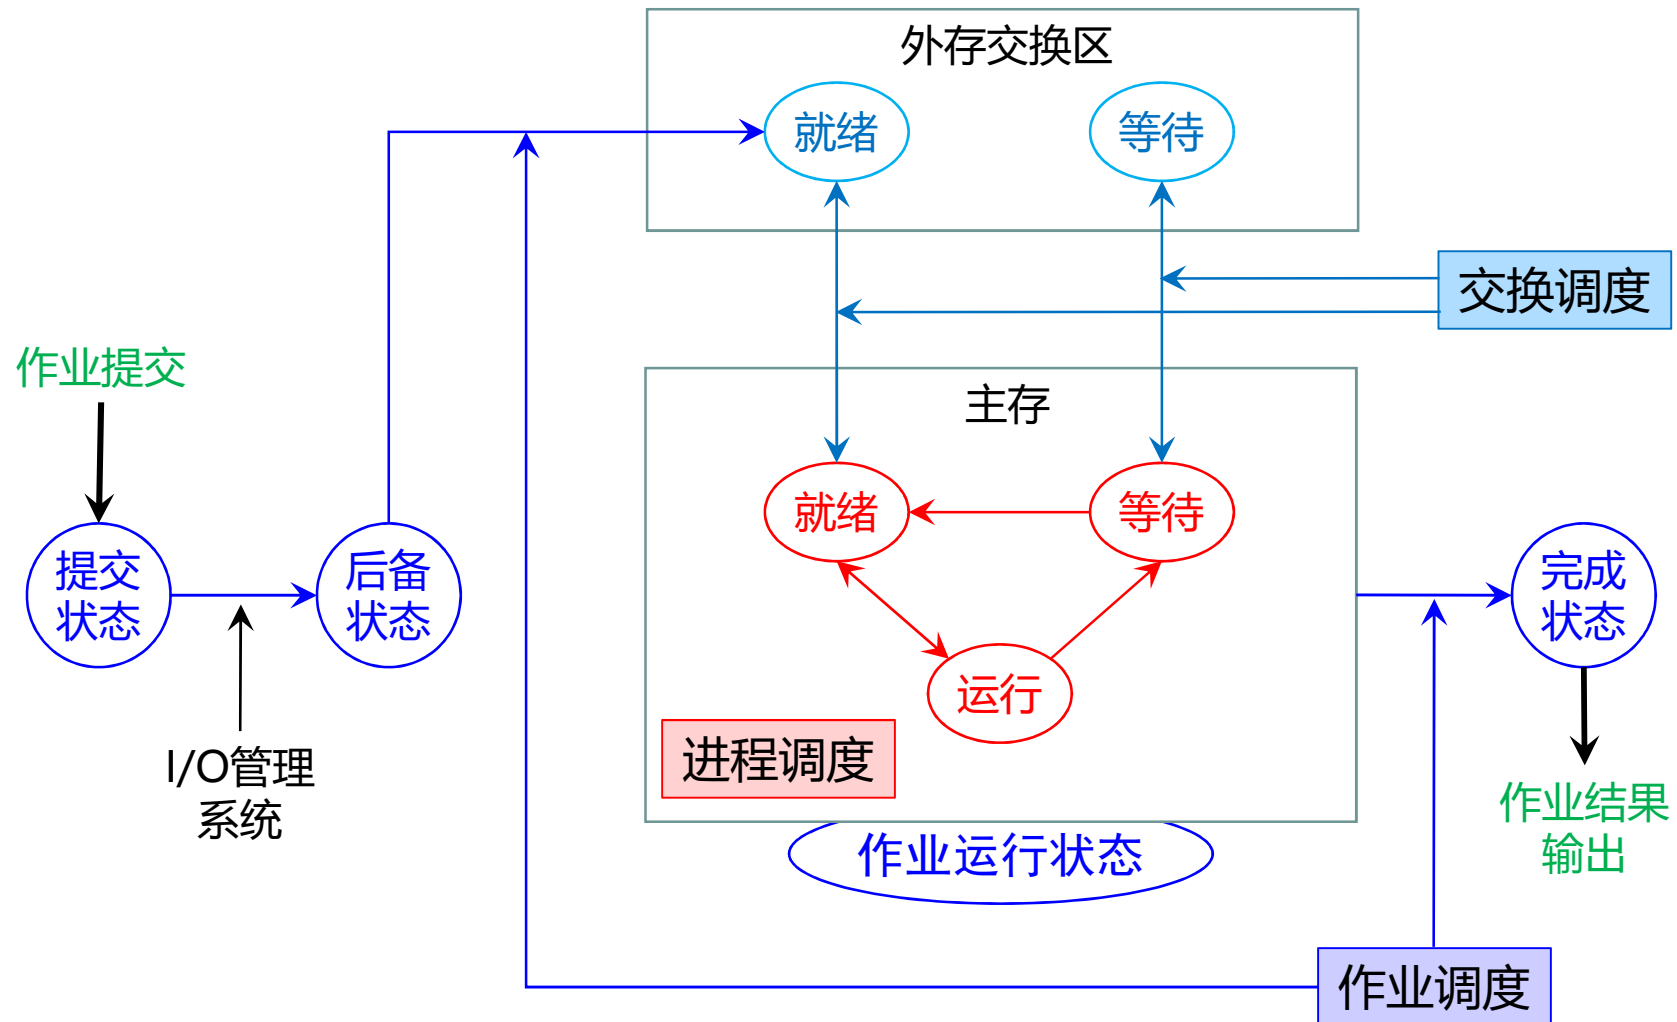
■ Schematic View of Schedulers

- Job scheduling, Swapping scheduling and Process/Thread scheduling.




■ Schematic View of Schedulers

- Job scheduling, Swapping scheduling and Process/Thread scheduling.



■ Process Swapping

■ Schematic View of Swapping.




■ Process Swapping

■ Dynamics of Swapping

- A process can be swapped temporarily out of memory to a backing storage, and then brought back into memory for continued execution.
- Backing storage
 - fast **disk** large enough to accommodate copies of all memory images for all users
 - must provide direct access to these memory images
- Roll out, roll in
 - swapping variant used for priority-based scheduling algorithms; lower-priority process is swapped out so higher-priority process can be loaded and executed
- Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped.
- Modified versions of swapping are found on many systems.
 - e.g., UNIX, Linux, and Windows.
- System maintains a ready queue of ready-to-run processes which have memory images on disk.

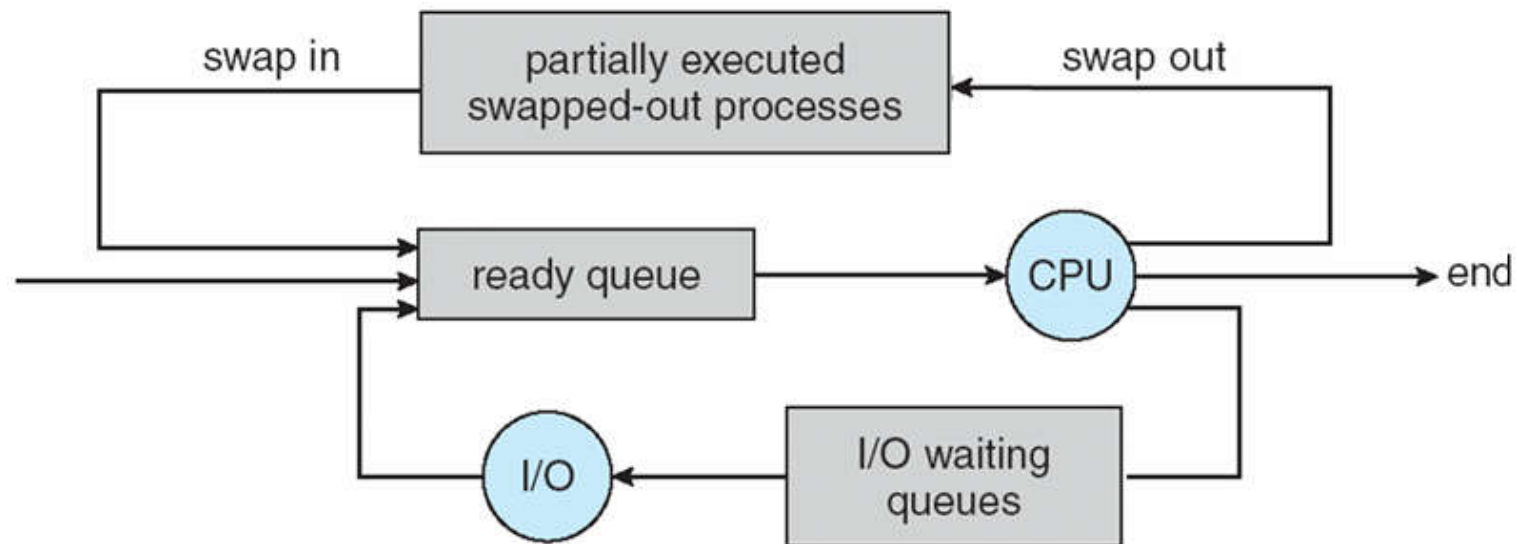
■ Process Swapping

■ Swapping Example.



■ Process Swapping

- Addition of Medium-term Scheduling.





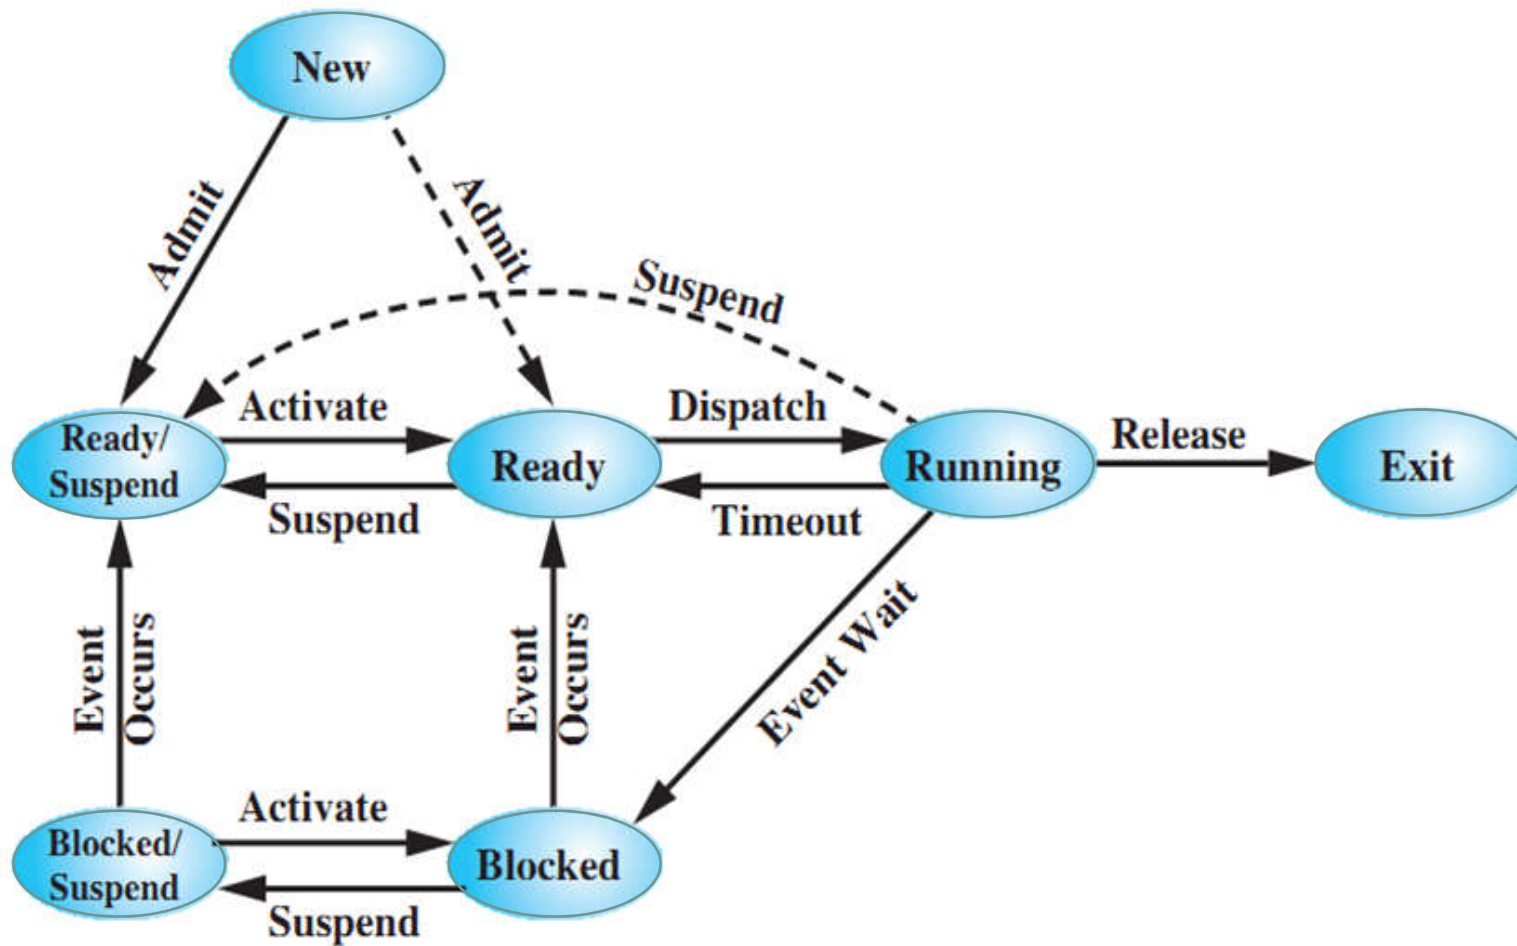
■ Process Swapping

■ Support of Swapping

- OS may need to suspend some processes, i.e., to swap them out to disk and then swap them back in.
- Two new states might need to be added:
 - Blocked Suspend: blocked processes which have been swapped out to disk
 - Ready Suspend: ready processes which have been swapped out to disk

■ Process Scheduling Activities

- A Seven-state Process Model.



■ Process Scheduling Activities

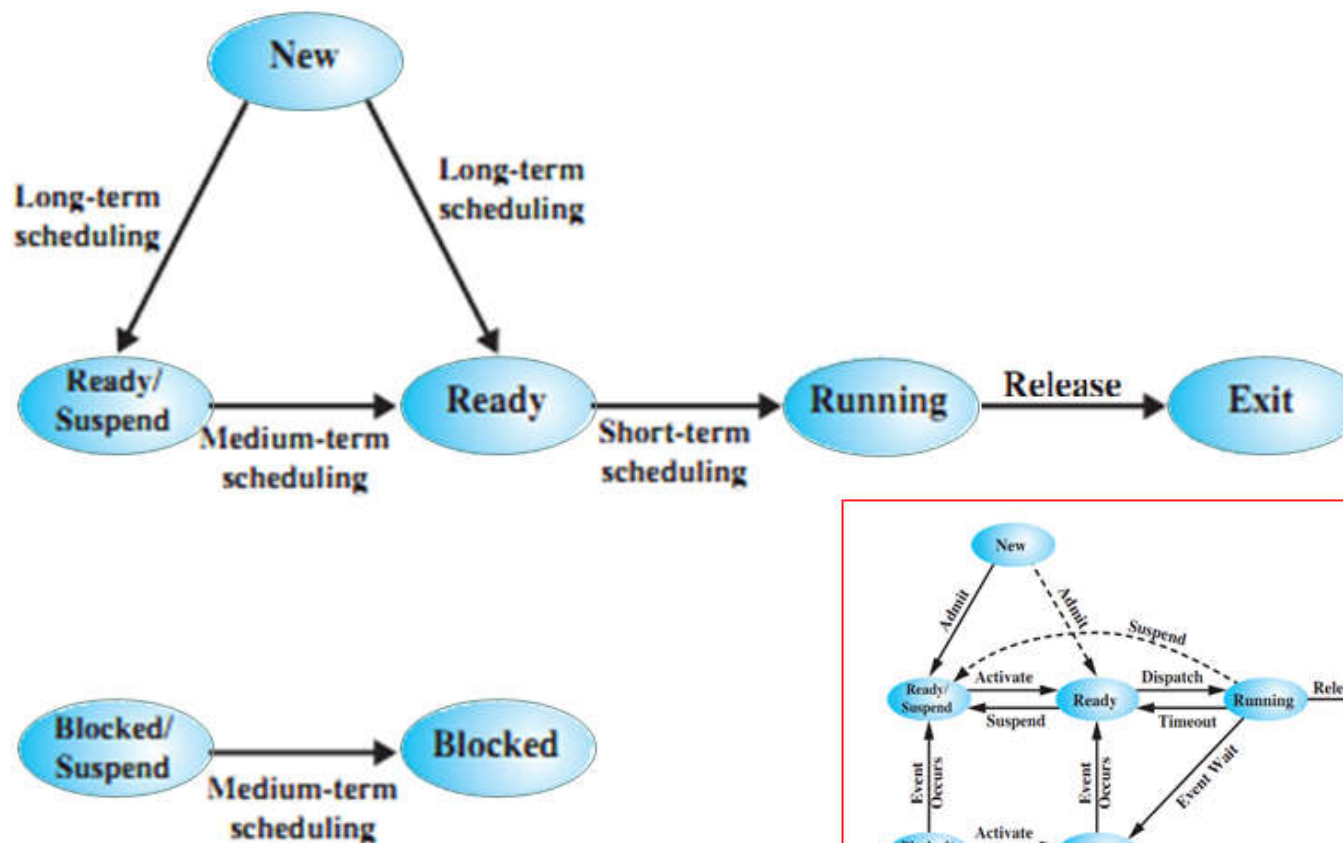
■ A Seven-state Process Model

■ Additional state transitions

- Blocked → Blocked Suspend
 - When all processes are blocked, the OS will make room to bring a ready process in memory.
- Blocked Suspend → Ready Suspend
 - when the event for which it has been waiting occurs (state info is available to OS).
- Ready Suspend → Ready
 - when no more ready processes in main memory
- Ready → Ready Suspend (unlikely)
 - when there are no blocked processes and memory must to be freed for adequate performance

■ Process Scheduling Activities

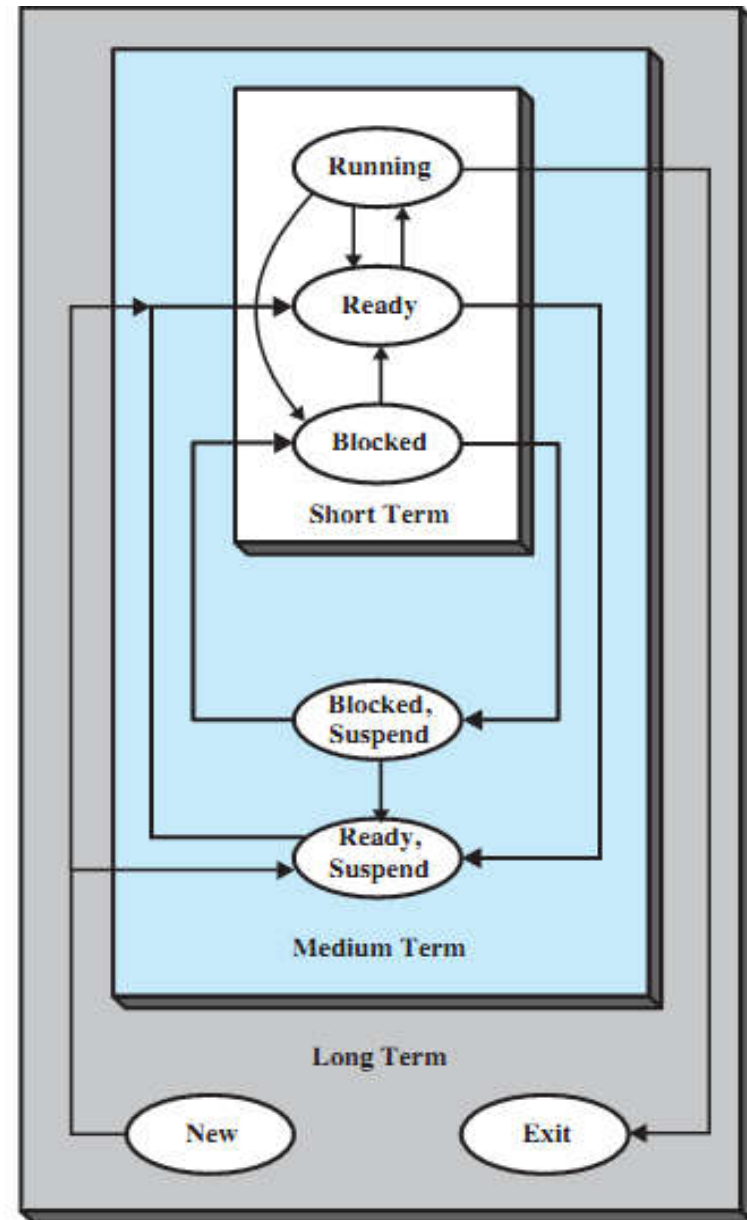
- A Seven-state Process Model
 - Scheduling Levels and Scheduling Activities.



■ Process Scheduling Activities

- A Seven-state Process Model.

- Another view of the three levels of scheduling.



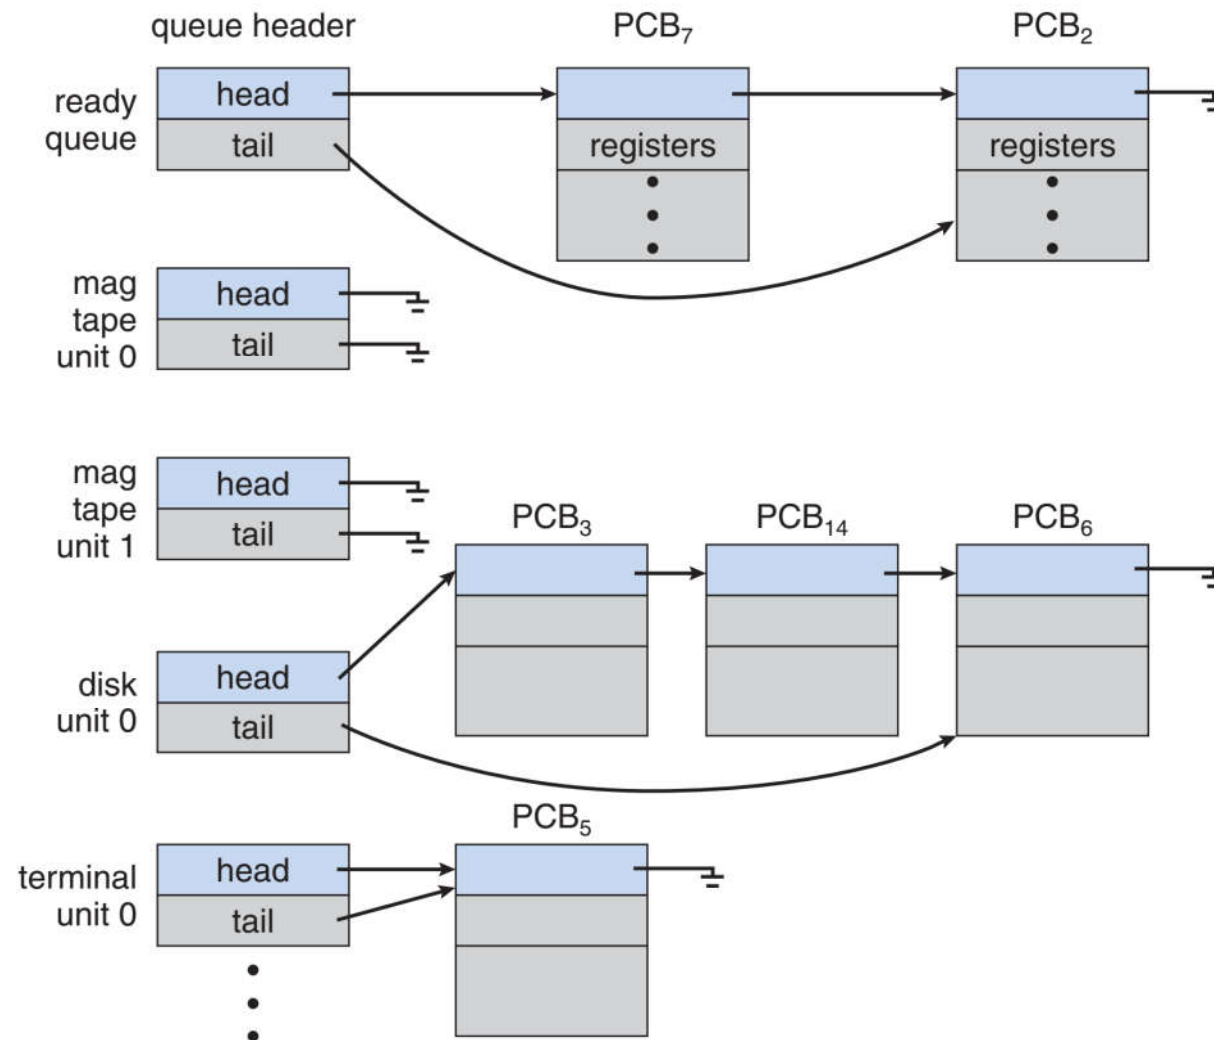
■ Process Scheduling Queues

■ Job Queue, Ready Queue and Device Queue

- As processes enter the system, they are put into a *job queue*, which consists of all processes in the system.
- The processes that are residing in main memory and are ready and waiting to execute are kept on a list called the *ready queue*.
 - The ready queue is generally stored as a linked list. The header contains pointers to the first and final PCBs in the list. Each PCB includes a pointer field that points to the next PCB in the ready queue.
- The list of processes waiting for a particular I/O device is called a *device queue*. Each device has its own device queue.
 - Suppose A process makes an I/O request to a shared device such as a disk, but the disk may be busy with the I/O request of some other process. The process therefore may have to wait for the disk.
 - A device queue is not a queue of devices, but a queue of process waiting for a particular I/O device.

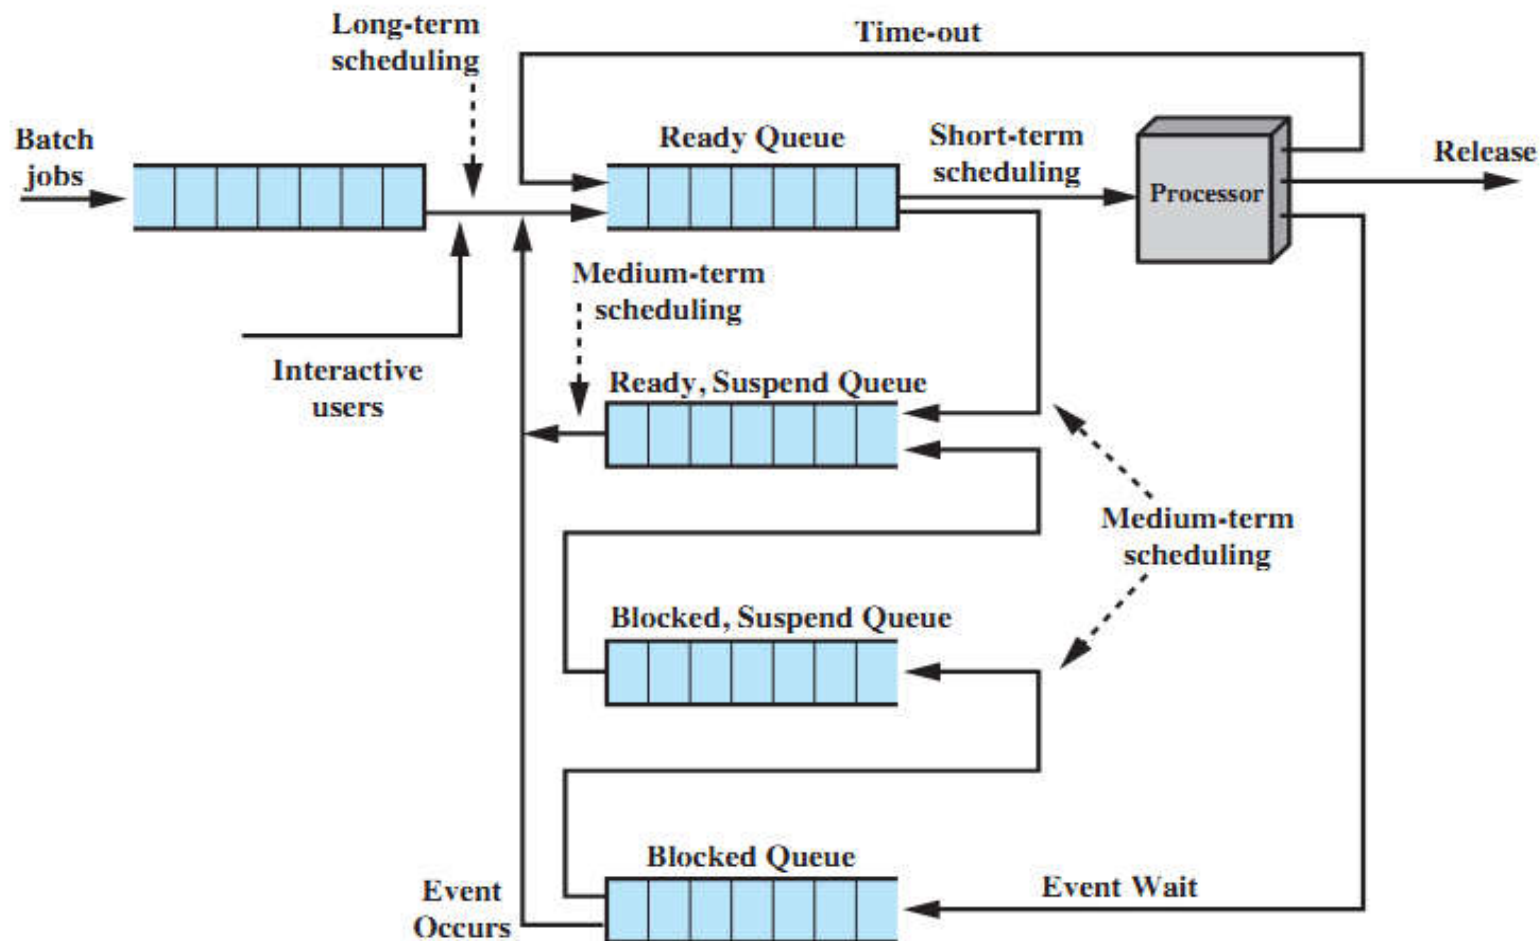
■ Process Scheduling Queues

- The ready queue and various I/O device queues.



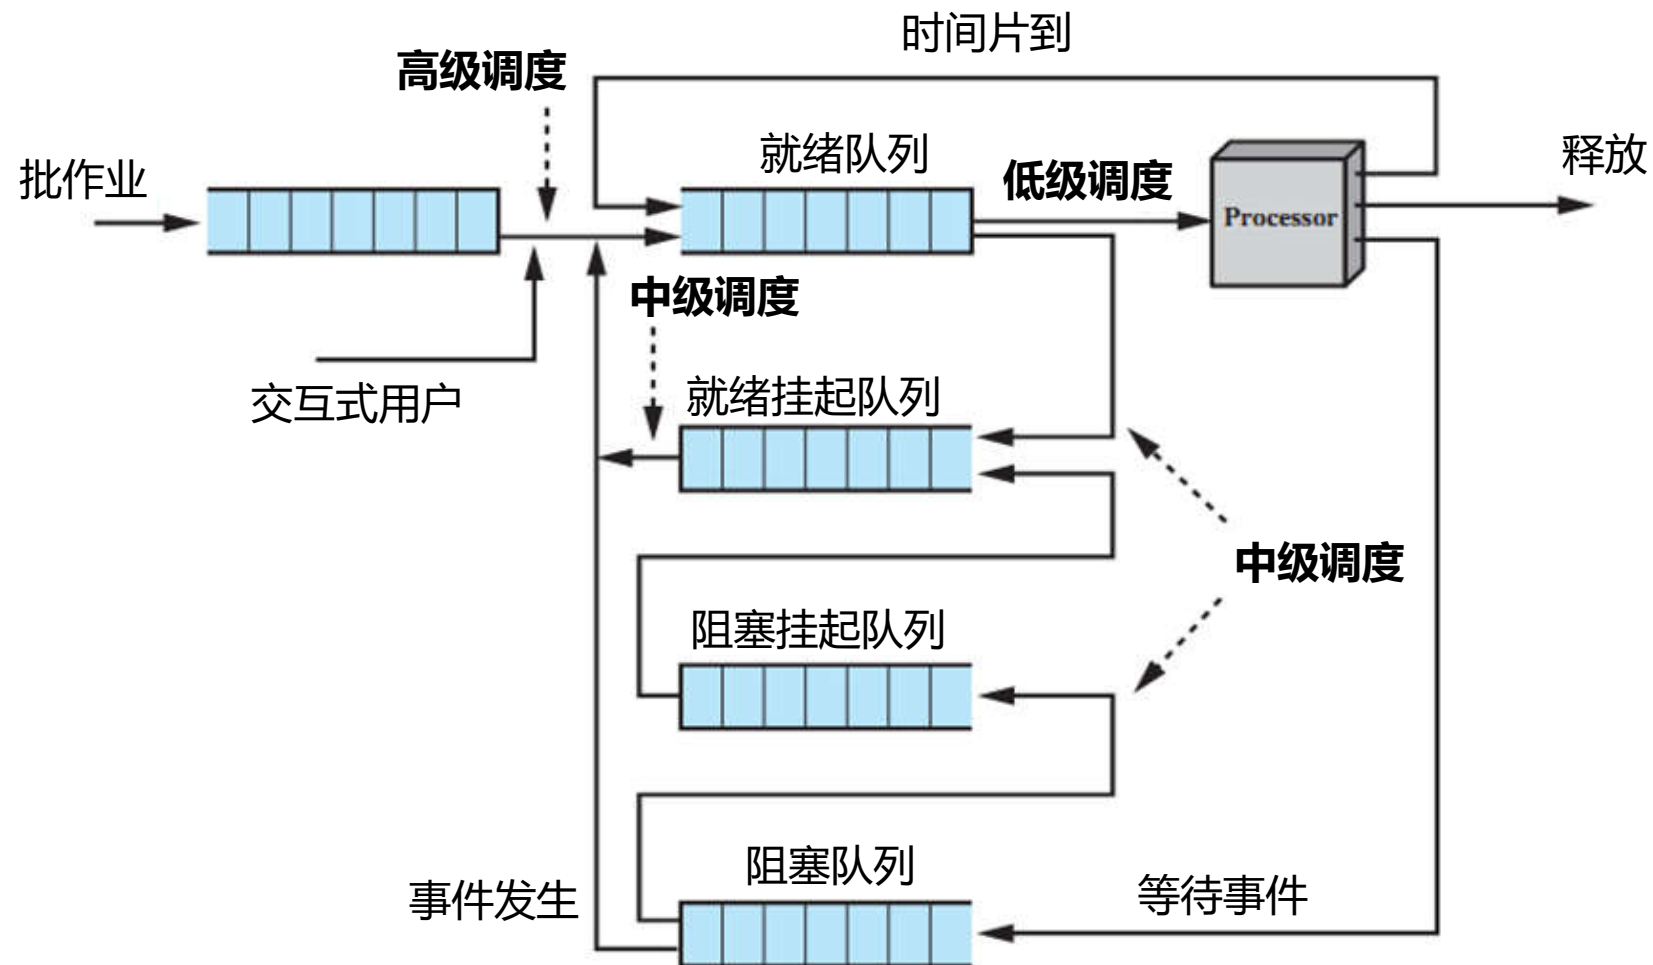
■ Process Scheduling Queues

- Queuing diagram for process scheduling
 - Processes migrate among the various queues.



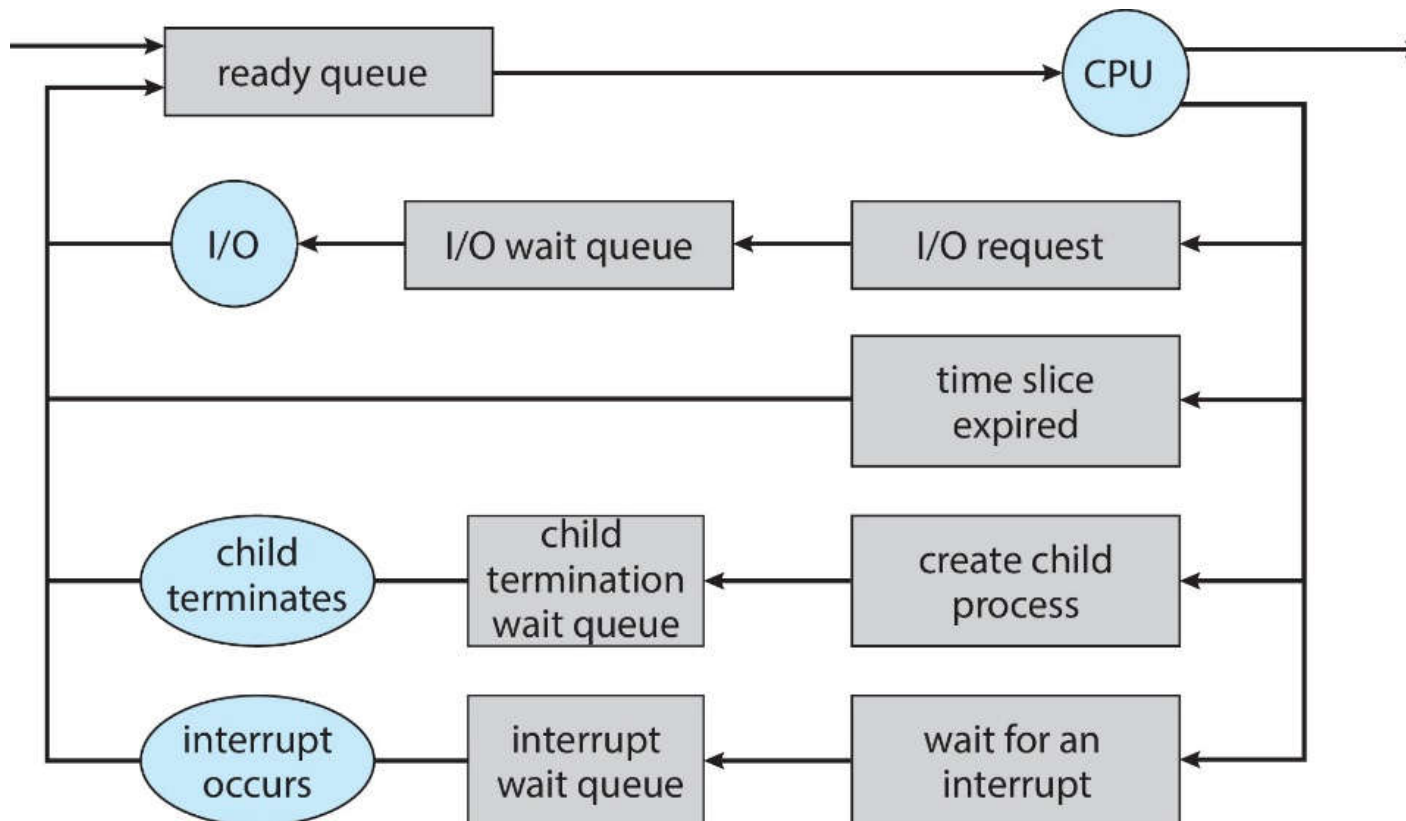
Process Scheduling Queues

- Queuing diagram for process scheduling
 - Processes migrate among the various queues.



■ Process Scheduling Queues

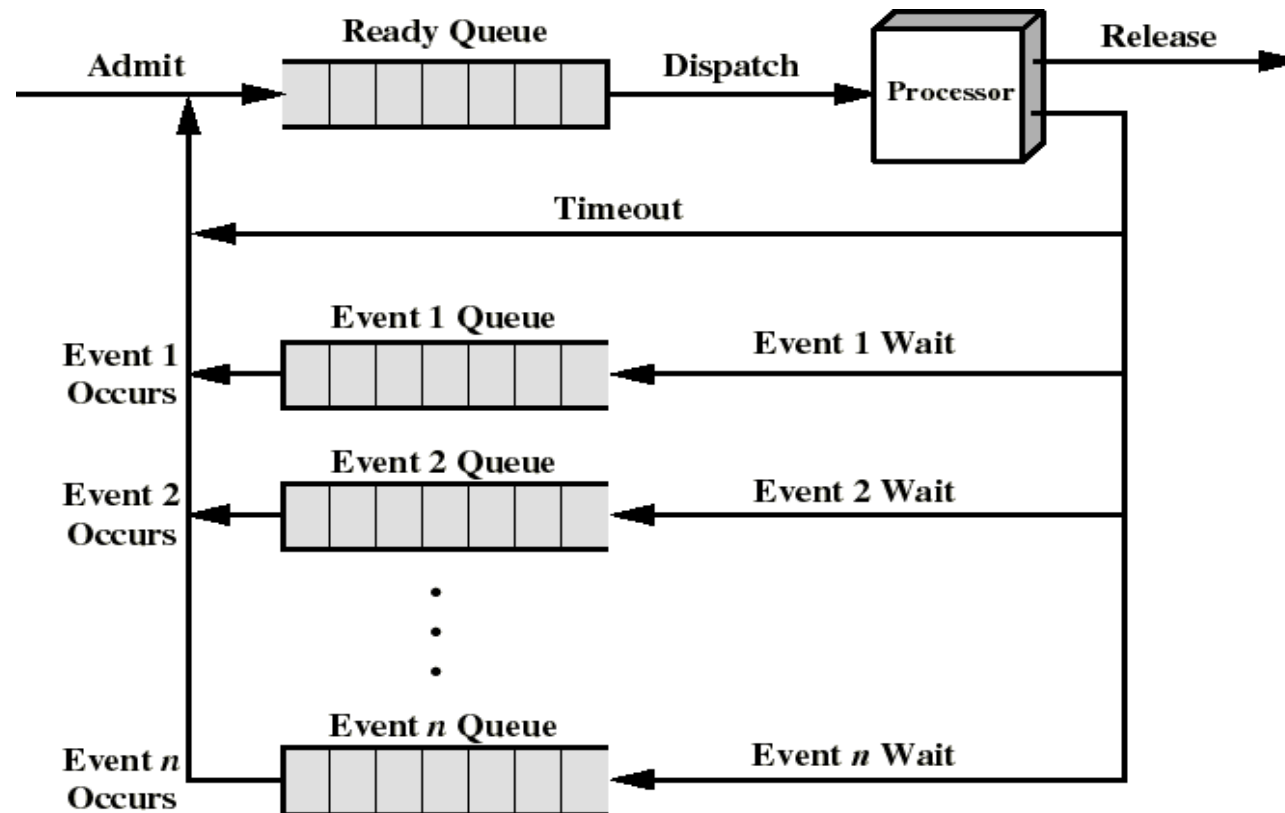
- Queuing diagram for process scheduling.



■ Process Scheduling Queues

■ A Queuing Discipline


- When some event occurs, the corresponding process is moved into the ready queue.



■ Process Switching


■ Multiprogramming Scheduling.

One program counter




(a)

Four program counters



(b)



(c)



■ Process Switching

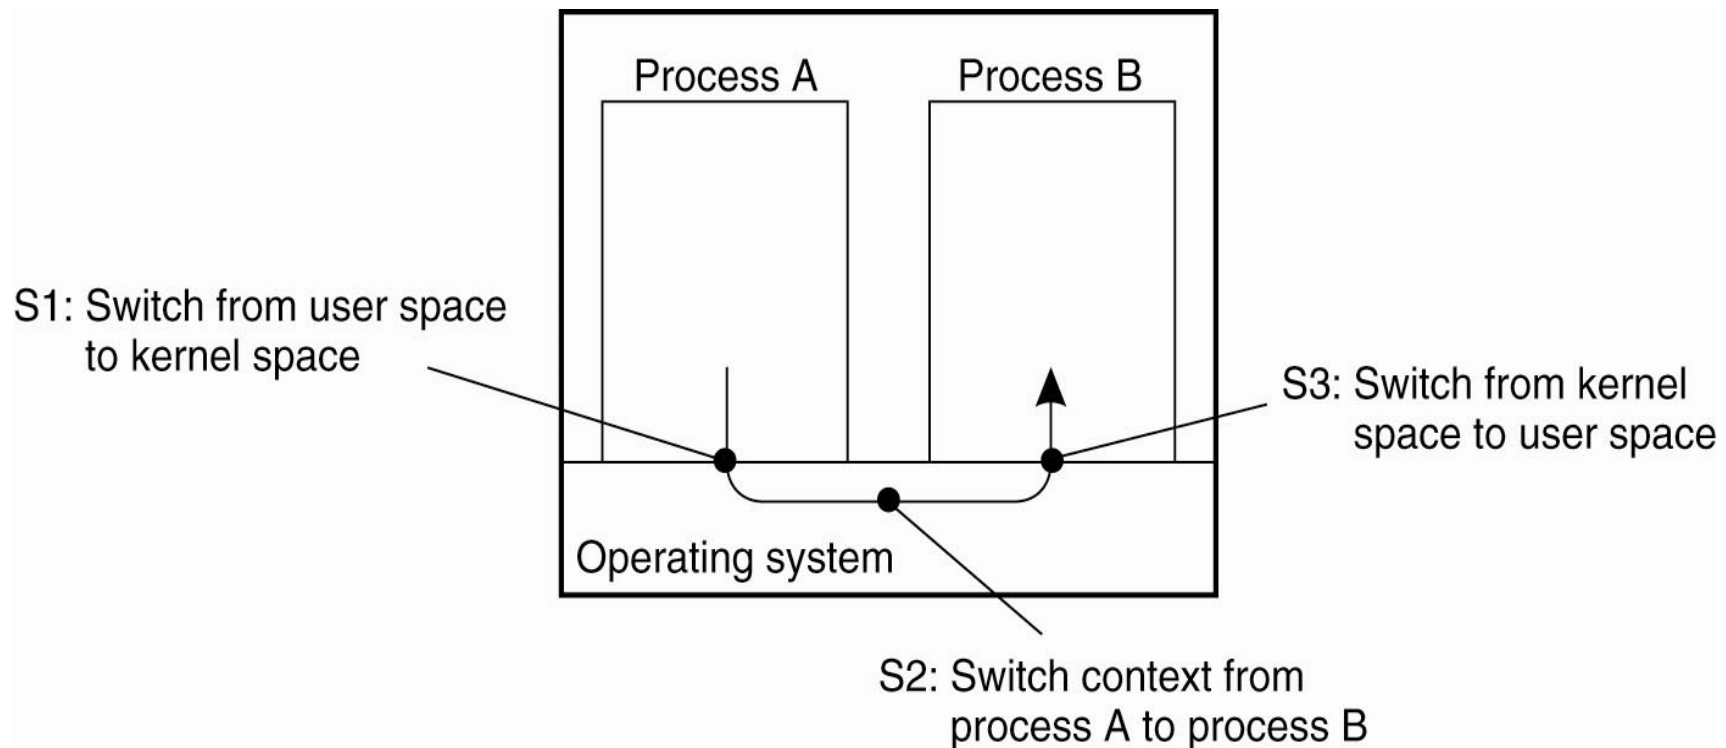
- A process switch (进程切换) may occur whenever the OS has gained control of CPU. i.e., when:
 - Supervisor Call
 - explicit request by the program (example: file open)
 - The process will probably be blocked.
 - Trap
 - an error resulted from the last instruction
 - It may cause the process to be moved to terminated state.
 - Interrupt
 - the cause is external to the execution of the current instruction
 - Control is transferred to Interrupt Handler.

■ Context Switching

- When CPU switches to another process, the system must save the state of the rolling out process and load the saved state of the rolling in process.
 - This is called *Context Switch* (上下文切换).
 - Context of a process is represented in the PCB.
- The time it takes is dependent on hardware support.
- Context-switch time is overhead (开销).
 - The system does no useful work while switching.

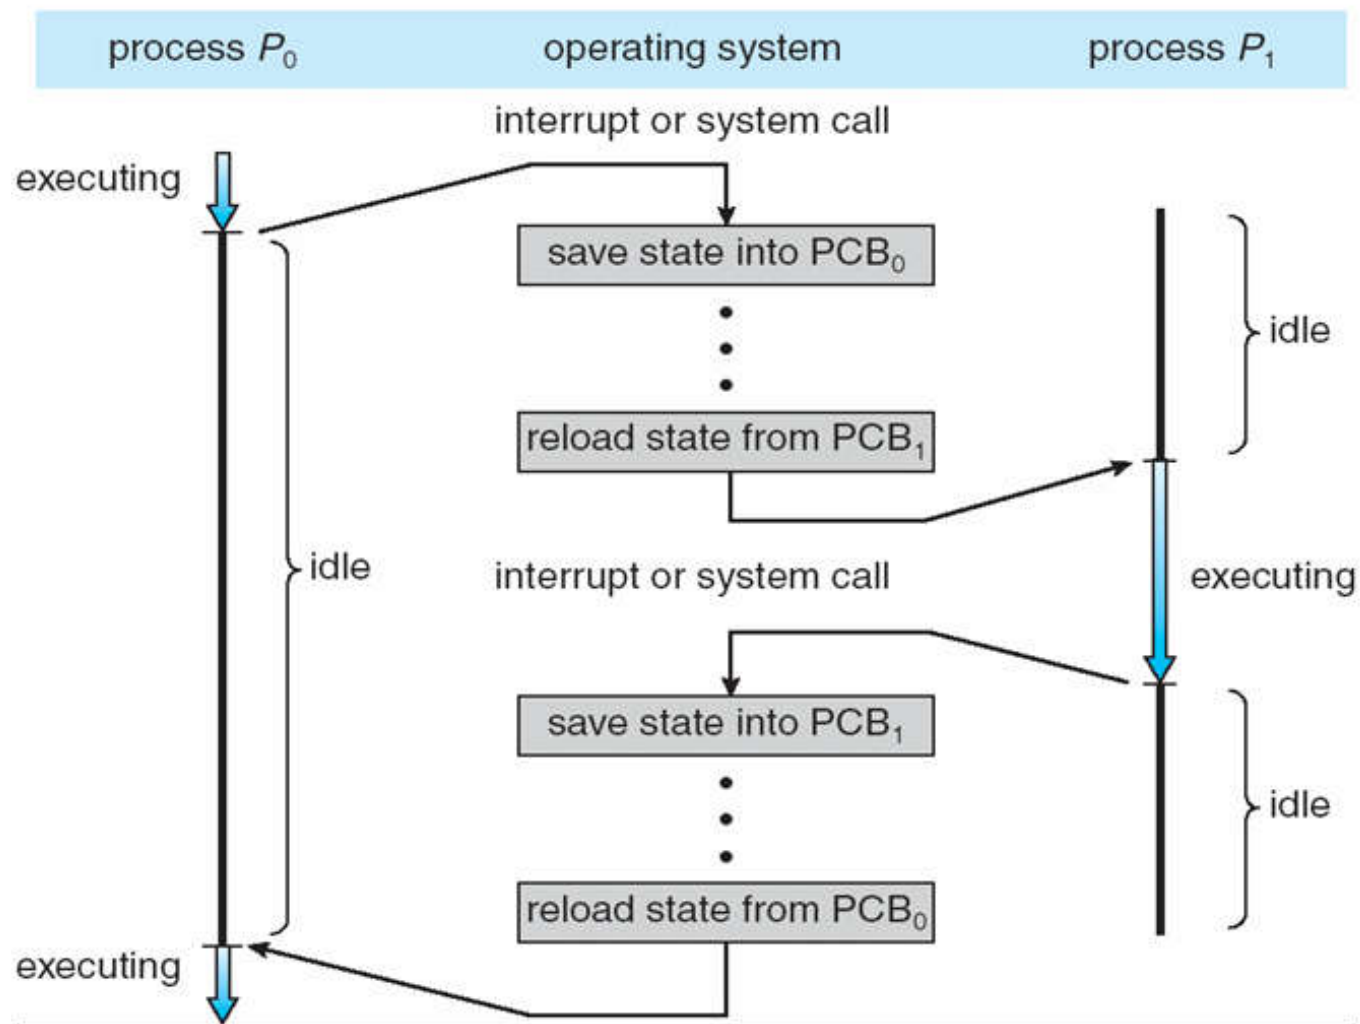
■ Context Switching

- Context Switch between Processes.



Context Switching

- Context Switch between Processes.



■ Context Switching

■ Steps in Context Switch

- Save **CPU context** including program counter and other registers.
- Update the **PCB** of the running process with its new state and other associate information.
- Move PCB to an appropriate **queue** – ready, blocked,
- **Select** another process (for next execution).
- Update **PCB** of the selected process.
- Restore **CPU context** from that of the selected process.



■ Mode Switching

- An interrupt may not produce a context switch.
 - The control can just return to the interrupted program.
- Then only the processor state information needs to be saved on stack.
- This is called *Mode Switch* (模式切换)
 - user mode to kernel mode when going into Interrupt Handler.
- Less overhead
 - no need to update the PCB like for context switch.