

---

# Threads

## Operating Systems

---

School of Data & Computer Science  
Sun Yat-sen University

Lecture Notes: [os\\_sysu@163.com](mailto:os_sysu@163.com)  
Instructor: Guoyang Cai  
email: [isscg@mail.sysu.edu.cn](mailto:isscg@mail.sysu.edu.cn)





## ■ Contents

- Overview
- Multicore Programming
- User and Kernel Level Threads
- Multithreading Models
- Threads Libraries
  - POSIX Pthreads
- Implicit Threading
  - Threads Pools
  - OpenMP
- Threading Issues
  - Semantics of `fork()` and `exec()` system calls
  - Signal handling
  - Thread cancellation
  - Thread-local storage
  - Scheduler activations.
- Linux clone()
- Example: Windows Threads

## ■ Thread Libraries

- A *thread library* provides the programmer with an **API** for creating and managing threads.
- There are two primary ways of implementing a thread library.
  - Provide a library entirely in user space with no kernel support.
    - All code and data structures for the library exist in user space.
    - Invoking a function in the library results in a **local function call in user space** and not a system call.
  - Implement a kernel-level library supported directly by OS.
    - Code and data structures for the library exist in kernel space.
    - Invoking a function in the API for the library typically results in a **system call to the kernel**.

## ■ Thread Libraries

- Three main thread libraries are in use today:
  - POSIX Pthreads
    - Pthreads, the threads extension of the POSIX standard, may be provided as either a **user-level** or a **kernel-level** library.
    - Used on UNIX and Linux systems.
  - The Windows thread library
    - A kernel-level library available on Windows systems.
  - The Java thread API
    - It allows threads to be created and managed directly in Java programs.
    - In most instances the JVM is running on top of a host operating system, the Java thread API is generally implemented using a thread library available on the host system.
    - This means that on Windows systems, Java threads are typically implemented using the Windows API.

## ■ Thread Libraries

- Asynchronous threading and Synchronous threading
  - Asynchronous threading and synchronous threading are two general strategies for creating multiple threads.
  - *Asynchronous threading*
    - Once the parent creates a child thread, the parent resumes its execution, so that the parent and child execute concurrently.
    - Each thread runs independently of every other thread, and the parent thread need not know when its child terminates.
    - There is typically **little data sharing** between threads.
  - *Synchronous threading*
    - The parent thread creates one or more children and then waits for all of its children to terminate before it resumes — the so-called *fork-join* strategy (分支聚合策略).
    - The threads created by the parent perform work concurrently. Once each thread has finished its work, it terminates and joins with its parent. Only after *all* of the children have joined can the parent resume execution.
    - It involves **significant data sharing** among threads.

## ■ Thread Libraries

### ■ POSIX Pthreads

- *POSIX Pthreads* refers to the POSIX standard (IEEE 1003.1c) defining an API for thread creation and synchronization. It provide support either for ULT or KLT.
- Pthreads is a *specification* for thread behavior. Operating-system designers may implement the specification in any way they wish.
- Numerous systems implement the Pthreads specification.
  - UNIX-type systems, including Linux, Mac OS X, and Solaris.

Thread call	Description
pthread_create	Create a new thread
pthread_exit	Terminate the calling thread
pthread_join	Wait for a specific thread to exit
pthread_yield	Release the CPU to let another thread run
pthread_attr_init	Create and initialize a thread's attribute structure
pthread_attr_destroy	Remove a thread's attribute structure

Some of the Pthreads function calls

## ■ Thread Libraries

### ■ POSIX Pthreads

- Windows doesn't support Pthreads natively.
  - Some thirdparty implementations for Windows are available.
- The C program shown in next slide demonstrates the basic Pthreads API for constructing a multithreaded program that calculates the summation of a nonnegative integer in a separate thread:

$$sum = \sum_{i=0}^N i$$



## Thread Libraries

### Pthreads – create & join

#### alg.13-1-pthread-create.c (1)

```
/* gcc -lpthread | -pthread */
int sum; /* shared by threads */
static void *runner(void *); /* thread function */
int main(int argc, char *argv[])
{
    if(argc < 2) {
        printf("usage: ./a.out <positive integer value>\n");
        return -1;
    }
    pthread_t ptid; /* thread identifier */
    pthread_attr_t attr; /* thread attributes structure */
    pthread_attr_init(&attr); /* set the default attributes */
    /* create the thread - runner with argv[1] */
    ret = pthread_create(&ptid, &attr, &runner, argv[1]);
    if (ret != 0) {
        perror("pthread_create()");
        return 1;
    }
    ret = pthread_join(ptid, NULL); /* join() waiting until thread tid returns */
    if (ret != 0) {
        perror("pthread_join()");
        return 1;
    }
    printf("sum = %d\n", sum);
    return 0;
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
```



## ■ Thread Libraries

### ■ Pthreads – create & join

#### ■ [alg.13-1-pthread-create.c \(2\)](#)

```
/* The thread will begin control in this function */
static void *runner(void *param)
{
    int i, upper;

    upper = atoi(param);
    sum = 0;
    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```



## ■ Thread Libraries

- Pthreads – create & join

- alg.13-1-pthread-create.c (2)

/\* The thread will begin control in this function \*/

```
iisscgy@ubuntu:/mnt/os-2020$ gcc alg.13-1-pthread-create.c -pthread
iisscgy@ubuntu:/mnt/os-2020$ ./a.out
usage: a.out <positive integer value>
iisscgy@ubuntu:/mnt/os-2020$ ./a.out 10
sum = 55
iisscgy@ubuntu:/mnt/os-2020$ ./a.out 100
sum = 5050
iisscgy@ubuntu:/mnt/os-2020$ ./a.out -10
sum = 0
iisscgy@ubuntu:/mnt/os-2020$ ./a.out asd
sum = 0
iisscgy@ubuntu:/mnt/os-2020$
```

## ■ Thread Libraries

### ■ Pthreads – create & join

#### ■ [alg.13-1-pthread-create.c](#)

- [pthread.h](#) must be included by any Pthreads programs.
- When the program begins, a single thread of control begins in [main\(\)](#). After some initialization, [main\(\)](#) creates a second thread that begins control in the [runner\(\)](#) function. Both threads *share the global data sum*.
- [pthread\\_t tid](#) declares the identifier [tid](#) for the thread we will create.
- [pthread\\_attr\\_t attr](#) declares the attributes for the thread, set by [pthread\\_attr\\_init\(&attr\)](#). The default attributes are provided without explicitly setting.
- [pthread\\_create\(\)](#) creates a separate thread with the attributes [attr](#). The thread identifier is passing to [tid](#). The name of the function where the new thread will begin execution, the [runner\(\)](#), is also passed. Last, the integer parameter that was provided on the command line, [argv\[1\]](#), is passed.

## ■ Thread Libraries

### ■ Pthreads – create & join

#### ■ [alg.13-1-pthread-create.c](#)

- Now there are two threads: the *parent thread* in `main()` and the *child thread* performing the summation in the `runner()`.
- The program follows the *fork-join synchronous strategy*: after creating the summation thread, the parent thread will *wait* for it to terminate by calling the `pthread_join()` function.
- The summation thread will terminate when it calls the function `pthread_exit()`. Once the summation thread has returned, the parent thread will resume, output the value of the shared data `sum`.
- This example program creates only a single thread. A simple method for waiting on several threads using the `pthread_join()` function is to enclose the operation within a simple `for` loop.

```
#define NUM_THREADS 10
pthread_t workers[NUM_THREADS];

... ..
for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```



## ■ Thread Libraries

### ■ Pthreads – create & join

#### ■ [alg.13-1-pthread-create.c](#)

- `void pthread_exit(void *retval); int pthread_join(tid1, &tret);`
  - If the second para of `pthread_join()` is not `NULL`, `pthread_exit()` transfers `*retval` back.

```
/* alg.13-1-pthread-create-1-1.c */
... ..
int main(int argc, char *argv[])
{
    ... ..
    pthread_create(&ptid, &attr, &runner, argv[1]);
    ... ..
    int *retptr;
    pthread_join(ptid, (void **)&retptr);
    printf("runner returned val = %d\n", *retptr);
    ... ..
}

static void *runner(void *param)
{
    ... ..
    int *retptr = (int *)malloc(sizeof(int)); /* allocated in process space */
    *retptr = 16;
    pthread_exit((void *)retptr);
}
```



## ■ Thread Libraries

### ■ Pthreads – create & join

#### ■ [alg.13-1-pthread-create.c](#)

- `void pthread_exit(void *retval); int pthread_join(tid1, &tret);`
  - If the second para of `pthread_join()` is not `NULL`, `pthread_exit()` transfers `*retval` back.

```
/* alg.13-1-pthread-create-1-2.c */
int sum;
... ..
int main(int argc, char *argv[])
{
    ... ..
    pthread_create(&ptid, &attr, &runner, argv[1]);
    ... ..
    int *retptr;
    pthread_join(ptid, (void **)&retptr);
    printf("runner returned val = %d\n", *retptr);
    ... ..
}

static void *runner(void *param)
{
    ... ..

    pthread_exit((void *)&sum); /* also: return (void *)&sum; */
}:wq
```



## ■ Thread Libraries

### ■ Pthreads – create & join

#### ■ [alg.13-1-pthread-create.c](#)

- `void pthread_exit(void *retval); int pthread_join(tid1, &tret);`
  - If the second para of `pthread_join()` is not `NULL`, `pthread_exit()` transfers `*retval` back.

```
/* alg.13-1-pthread-create-1-3.c */
/* without global variable sum defined */
int main(void)
{
    int sum;
    int *retptr = &sum;    ... ..

    pthread_create(&ptid, &attr, &runner, &sum);
    ... ..
    pthread_join(ptid, (void **)&sum);
    printf("runner returned val = %d\n", *retptr);
    ... ..
}

static void *runner(void *param)
{
    int *sum = (int *)param;
    ... ..
    pthread_exit((void *)sum); /* return the address in process stack segment */
}
```



## ■ Thread Libraries

### ■ Pthreads – create & join

#### ■ [alg.13-1-pthread-create.c](#)

- `void pthread_exit(void *retval); int pthread_join(tid1, &tret);`
  - If the second para of `pthread_join()` is not NULL, `pthread_exit()` transfers `*retval` back.

```
/* alg.13-1-pthread-create-2.c */
... ..
#include <string.h>
int main(int argc, char *argv[])
{
    ... ..
    pthread_create(&tid, &attr, &runner, argv[1]);
    ... ..
    char *retptr;
    pthread_join(tid, (void **)&retptr);
    printf("runner returned val = %s\n", retptr);
    ... ..
}

static void *runner(void *param)
{
    ... ..
    char msg[] = "Hello world!";
    char *retptr = (char *)malloc((strlen(msg) + 1)*sizeof(char));
    strcpy(retptr, msg);
    pthread_exit((void *)retptr);
}
```





## ■ Thread Libraries

### ■ Pthreads – create & join

#### ■ [alg.13-1-pthread-create-3.c \(1\)](#)

```
#include <stdlib.h>
#include <string.h>
#include <unistd.h>
#include <pthread.h>

static void *ftn(void *arg)
{
    int *numptr = (int *)arg;
    int num = *numptr;
    char *retval = (char *)malloc(80*sizeof(char)); /* allocated in process heap */

    sprintf(retval, "This is thread-%d, ptid = %lu", num, pthread_self( ));
    printf("%s\n", retval);

    pthread_exit((void *)retval); /* or return (void *)retval; */
}

int main(int argc, char *argv[])
{
    int max_num = 5;
    int i, ret;

    printf("Usage: ./a.out total_thread_num\n");
    if(argc > 1)
        max_num = atoi(argv[1]);
    printf("main(): pid = %d, ptid = %lu.\n", getpid( ), pthread_self( ));
```



## Thread Libraries

### Pthreads – create & join

#### alg.13-1-pthread-create-3.c (2)

```
pthread_t ptid[max_num];
for(i = 0; i < max_num; i++) {
    ret = pthread_create(&ptid[i], NULL, ftn, (void *)&i);
    if(ret != 0) {
        fprintf(stderr, "pthread_create error: %s\n", strerror(ret));
        exit(1);
    }
}

for(i = 0; i < max_num; i++) {
    char *retptr; /* retptr pointing to address allocated by ftn() */
    ret = pthread_join(ptid[i], (void **)&retptr);
    if(ret!=0) {
        fprintf(stderr, "pthread_join error: %s\n", strerror(ret));
        exit(1);
    }
    printf("thread-%d: retval = %s\n", i, retptr);
    free(retptr);
    retptr = NULL; /* preventing ghost pointer */
}

return 1;
}
```



### ■ Thread Libraries

- Pthreads – create & join
  - [alg.13-1-pthread-create-3.c](#)

```
iisscgy@ubuntu:/mnt/os-2020$ gcc alg.13-1-pthread-create-3.c -pthread
iisscgy@ubuntu:/mnt/os-2020$ ./a.out 5
Usage: ./a.out total_thread_num
main(): pid = 18960, ptid = 140253726336832.
This is thread-2, ptid = 140253717825280
This is thread-3, ptid = 140253709432576
This is thread-4, ptid = 140253701039872
This is thread-0, ptid = 140253684254464
This is thread-0, ptid = 140253692647168
thread-0: retval = This is thread-2, ptid = 140253717825280
thread-1: retval = This is thread-3, ptid = 140253709432576
thread-2: retval = This is thread-4, ptid = 140253701039872
thread-3: retval = This is thread-0, ptid = 140253692647168
thread-4: retval = This is thread-0, ptid = 140253684254464
iisscgy@ubuntu:/mnt/os-2020$
```



## Thread Libraries

### Pthreads – create & join

#### alg.13-1-pthread-create-3-1.c

```
pthread_t ptid[max_num];
for(i = 0; i < max_num; i++) {
    ret = pthread_create(&ptid[i], NULL, ftn, (void *)&i);
    if(ret != 0) {
        fprintf(stderr, "pthread_create error: %s\n", strerror(ret));
        exit(1);
    }
    sleep(1);
}

for(i = 0; i < max_num; i++) {
    char *retptr; /* retptr pointing to address allocated by ftn() */
    ret = pthread_join(ptid[i], (void **)&retptr);
    if(ret!=0) {
        fprintf(stderr, "pthread_join error: %s\n", strerror(ret));
        exit(1);
    }
    printf("thread-%d: retval = %s\n", i, retptr);
    free(retptr);
    retptr = NULL; /* preventing ghost pointer */
}

return 1;
}
```

*i* is not consistent  
to ptid[*i*]

## ■ Thread Libraries

- Pthreads – create & join

- [alg.13-1-pthread-create-3-1.c](#)

```
iisscgy@ubuntu:/mnt/os-2020$ gcc alg.13-1-pthread-create-3-1.c -pthread
iisscgy@ubuntu:/mnt/os-2020$ ./a.out 5
Usage: ./a.out total_thread_num
main(): pid = 18977, ptid = 139918462719808.
This is thread-0, ptid = 139918454208256
This is thread-1, ptid = 139918443620096
This is thread-2, ptid = 139918435227392
This is thread-3, ptid = 139918357559040
This is thread-4, ptid = 139918349166336
thread-0: retval = This is thread-0, ptid = 139918454208256
thread-1: retval = This is thread-1, ptid = 139918443620096
thread-2: retval = This is thread-2, ptid = 139918435227392
thread-3: retval = This is thread-3, ptid = 139918357559040
thread-4: retval = This is thread-4, ptid = 139918349166336
iisscgy@ubuntu:/mnt/os-2020$
```

```
}
```

```
return 1;
```

```
}
```



## ■ Thread Libraries

### ■ Pthreads – create & join

#### ■ [alg.13-1-pthread-create-4.c](#)

```
int main(int argc, char *argv[])
{
    int max_num = 5;
    int i, ret;

    printf("Usage: ./a.out total_thread_num\n");
    if(argc > 1) {
        max_num = atoi(argv[1]);
    }

    int thread_num[max_num];
    for (i = 0; i < max_num; i++) {
        thread_num[i] = i;
    }

    printf("main(): pid = %d, ptid = %lu.\n", getpid( ), pthread_self( ));

    pthread_t ptid[max_num];
    for(i = 0; i < max_num; i++) {
        ret = pthread_create(&ptid[i], NULL, ftn, (void *)&thread_num[i]);
        if(ret != 0) {
            fprintf(stderr, "pthread_create error: %s\n", strerror(ret));
            exit(1);
        }
    }
}
```



## ■ Thread Libraries

- Pthreads – create & join
  - [alg.13-1-pthread-create-4.c](#)

```
isscgy@ubuntu:/mnt/os-2020$ gcc alg.13-1-pthread-create-4.c -pthread
isscgy@ubuntu:/mnt/os-2020$ ./a.out 5
Usage: ./a.out total_thread_num
main(): pid = 19020, ptid = 140096962860864.
This is thread-0, ptid = 140096954349312
This is thread-1, ptid = 140096945956608
This is thread-3, ptid = 140096858744576
This is thread-4, ptid = 140096850351872
This is thread-2, ptid = 140096867137280
thread-0: retval = This is thread-0, ptid = 140096954349312
thread-1: retval = This is thread-1, ptid = 140096945956608
thread-2: retval = This is thread-2, ptid = 140096867137280
thread-3: retval = This is thread-3, ptid = 140096858744576
thread-4: retval = This is thread-4, ptid = 140096850351872
isscgy@ubuntu:/mnt/os-2020$
```

```
    if(ret != 0) {
        fprintf(stderr, "pthread_create error: %s\n", strerror(ret));
        exit(1);
    }
}
```



## ■ Thread Libraries

### ■ Pthreads – sharing memory

#### ■ [alg.13-2-pthread-shm.c \(1\)](#)

```
struct msg_stru {
    char msg1[MSG_SIZE], msg2[MSG_SIZE], msg3[MSG_SIZE];
}; /* global variable */

static void *runner1(void *);
static void *runner2(void *);

int main(void)
{
    pthread_t tid1, tid2;
    pthread_attr_t attr = {0};
    struct msg_stru msg; /* storage in main-thread stack */

    sprintf(msg.msg1, "message 1 by parent");
    sprintf(msg.msg2, "message 2 by parent");
    sprintf(msg.msg3, "message 3 by parent");
    printf("\nparent say:\n%s\n%s\n%s\n", msg.msg1, msg.msg2, msg.msg3);

    pthread_attr_init(&attr);
    if(pthread_create(&tid1, &attr, &runner1, (void *)&msg) != 0) {
        perror("pthread_create()"); return 1;
    }
    if(pthread_create(&tid2, &attr, &runner2, (void *)&msg) != 0) {
        perror("pthread_create()"); return 1;
    }
}
```

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <pthread.h>
#define MSG_SIZE 1024
```



## ■ Thread Libraries

### ■ Pthreads – sharing memory

#### ■ [alg.13-2-pthread-shm.c \(2\)](#)

```
        if(pthread_join(tid1, NULL) != 0) {
            perror("pthread_join()"); return 1;
        }
        if(pthread_join(tid2, NULL) != 0) {
            perror("pthread_join()"); return 1;
        }

        printf("\nparent say:\n%s\n%s\n%s\n", msg.msg1, msg.msg2, msg.msg3);
        return 0;
    }

    static void *runner1(void *param)
    {
        struct msg_stru *ptr = (struct msg_stru *)param;
        sprintf(ptr->msg1, "message 1 changed by child1");
        pthread_exit(0);
    }

    static void *runner2(void *param)
    {
        struct msg_stru *ptr = (struct msg_stru *)param;
        sprintf(ptr->msg2, "message 2 changed by child2");
        pthread_exit(0);
    }
```



### ■ Thread Libraries

- Pthreads – sharing memory
  - [alg.13-2-pthread-shm.c \(2\)](#)

```
if (pthread_join(tid1, NULL) != 0) {  
iisscgy@ubuntu:/mnt/os-2020$ gcc alg.13-2-pthread-shm.c -pthread  
iisscgy@ubuntu:/mnt/os-2020$ ./a.out  
  
parent say:  
message 1 by parent  
message 2 by parent  
message 3 by parent  
  
parent say:  
message 1 changed by child1  
message 2 changed by child2  
message 3 by parent  
iisscgy@ubuntu:/mnt/os-2020$
```

```
static void *runner2(void *param)  
{  
    struct msg_stru *ptr = (struct msg_stru *)param;  
    sprintf(ptr->msg2, "message 2 changed by child2");  
    pthread_exit(0);  
}
```



## ■ Thread Libraries

## ■ Pthreads - pthread\_attr\_setstack testing

alg.13-3-pthread-stack.c (1)

```
#include <stdio.h>
#include <pthread.h>
#include <unistd.h>
#include <malloc.h>
#define STACK_SIZE (524288-10000)*4096 /* nearly 2G: 2^19 = 524288 */

int i = 0;
static void *test(void *arg)
{
    char buffer[1024]; /* 1KiB saved to the thread stack */
    if(i > 5 && i < 1965030)
        printf("\b\b\b\b\b\b\b\b\b%8d", i);
    else
        printf("\niteration = %8d", i);
    i++;
    test(arg); /* recursive calling until segmentation fault */
}
```



### ■ Thread Libraries

#### ■ Pthreads - pthread\_attr\_setstack testing

##### ■ [alg.13-3-pthread-stack.c \(1\)](#)

```
int main(void)
{
    int ret;
    pthread_t ptid;
    pthread_attr_t tattr = {0};
    char *stackptr = malloc(STACK_SIZE);
    if(!stackptr) {
        perror("malloc()");
        return 1;
    }
    pthread_attr_init(&tattr);
    pthread_attr_setstack(&tattr, stackptr, STACK_SIZE);
    ret = pthread_create(&ptid, &tattr, &test, NULL);
    if(ret) {
        perror("pthread_create()");
        return 1;
    }
    sleep(1);
    pthread_join(ptid, NULL);

    return 0;
}
```



### ■ Thread Libraries

- Pthreads - pthread\_attr\_setstack testing

```
gcc alg.13-3-pthread-stack.c -pthread
iisscgy@ubuntu:/mnt/os-2020$ gcc alg.13-3-pthread-stack.c -pthread
iisscgy@ubuntu:/mnt/os-2020$ ./a.out
```

```
iteration =      0
iteration =      1
iteration =      2
iteration =      3
iteration =      4
iteration = 1965029
iteration = 1965030
iteration = 1965031
iteration = 1965032
```

```
Segmentation fault (core dumped)
```

```
iisscgy@ubuntu:/mnt/os-2020$
```

```
pthread_join(tid, NULL);
```

```
return 0;
```

```
}
```

$514288 * 4096 - 1965032 * 1024 = 94330880$   
 $94330880 / 1965032 = 48$  (bytes), for system overhead of each iteration

## ■ Implicit Threading

- *Implicit Threading* is a strategy to give better supports for the design of multithreaded applications. The creation and management of threading are transferred from application developers to compilers and run-time libraries.
- There are three alternative approaches for designing multithreaded programs that can take advantage of multicore processors through implicit threading.
  - Threads Pools
  - OpenMP
  - Grand Central Dispatch (Apple iOS)

## ■ Implicit Threading

### ■ Thread Pools

#### ■ Question:

- Unlimited threads could exhaust system resources, such as CPU time or memory.

#### ■ Solution.

- Create a number of threads at process startup and place them into a *thread pool*, where they sit and wait for work.
- When a server receives a request, it awakens a thread from this pool—if one is available—and passes it the request for service.
- Once the thread completes its service, it returns to the pool and awaits more work. If the pool contains no available thread, the server waits until one becomes free.

#### ■ Example.

- The max number of threads per process in IA-32.
  - The number is about 300 with *10M* default stack-size for each thread in 3G address space.
- A pool of less than 255 threads may be created.

## ■ Implicit Threading

### ■ Thread Pools

#### ■ Benefits of thread pools

- Usually *slightly faster* to service a request with an existing thread than wait to create a new thread.
- A thread pool *limits the number of threads* that exist at any one point. This is particularly important on systems that cannot support a large number of concurrent threads.
- Separating the task to be performed from the mechanics of creating the task allows us to use different *strategies* for running the task.

#### ■ Size of a thread pool

- The number of threads in the pool can be set heuristically based on factors such as the number of CPUs in the system, the amount of physical memory, and the expected number of concurrent client requests.
- More sophisticated thread-pool architectures such as Apple's Grand Central Dispatch can dynamically adjust the number of threads in the pool according to usage patterns.



## ■ Implicit Threading

### ■ OpenMP

- OpenMP is a set of *compiler directives* as well as an *API* for programs written in C, C++, or FORTRAN that provides support for parallel programming in *shared-memory* environments.
- OpenMP identifies *parallel regions* as blocks of code that may run in parallel.
  - Application developers insert compiler directives into their code at parallel regions, and these directives instruct the OpenMP run-time library to execute the region in parallel.
- The following C program illustrates a compiler directive above the parallel region containing the `printf()` statement.
  - When OpenMP encounters the directive

```
#pragma omp parallel
```

it creates as many threads as there are processing cores in the system (e.g., two threads per core for Intel CPU). All the threads simultaneously execute the parallel region. As each thread exits the parallel region, it is terminated.



## ■ Implicit Threading

### ■ OpenMP

#### ■ [alg.13-4-openmp-demo.c](#)

```
#include <stdio.h>
#include <omp.h>
#include <unistd.h>
```

```
/* compiling: gcc -fopenmp */
#define __NR_gettid 186 /* 186 for x86-64; 224 for i386-32 */
#define gettid() syscall(__NR_gettid)

int main()
{
    #pragma omp parallel
    {
        printf("region 1. pid = %d, tid = %ld\n", getpid(), gettid());
    }
    #pragma omp parallel num_threads(2)
    {
        printf("region 2, num_thread(2). pid = %d, tid = %ld\n", getpid(), gettid());
    }
    #pragma omp parallel num_threads(4)
    {
        printf("region 3, num_thread(4). pid = %d, tid = %ld\n", getpid(), gettid());
    }
    #pragma omp parallel num_threads(6)
    {
        printf("region 4, num_thread(6). pid = %d, tid = %ld\n", getpid(), gettid());
    }

    return 0;
}
```

## ■ Implicit Threading

### ■ OpenMP

#### ■ [alg.13-4-openmp-demo.c](#)

```
#include <stdio.h>
#include <omp.h>
#include <unistd.h>
```

```
i SSCgy@ubuntu:/mnt/os-2020$ gcc alg.13-4-openmp-demo.c -fopenmp
i SSCgy@ubuntu:/mnt/os-2020$ ./a.out
parallel region 1. pid = 19375, tid = 19375
parallel region 1. pid = 19375, tid = 19378
parallel region 1. pid = 19375, tid = 19377
parallel region 1. pid = 19375, tid = 19376
parallel region 2 with num_thread(2). pid = 19375, tid = 19376
parallel region 2 with num_thread(2). pid = 19375, tid = 19375
parallel region 3 with num_thread(4). pid = 19375, tid = 19375
parallel region 3 with num_thread(4). pid = 19375, tid = 19379
parallel region 3 with num_thread(4). pid = 19375, tid = 19376
parallel region 3 with num_thread(4). pid = 19375, tid = 19380
parallel region 4 with num_thread(6). pid = 19375, tid = 19379
parallel region 4 with num_thread(6). pid = 19375, tid = 19380
parallel region 4 with num_thread(6). pid = 19375, tid = 19376
parallel region 4 with num_thread(6). pid = 19375, tid = 19375
parallel region 4 with num_thread(6). pid = 19375, tid = 19382
parallel region 4 with num_thread(6). pid = 19375, tid = 19381
i SSCgy@ubuntu:/mnt/os-2020$
```

tid 19376 is reused.

## ■ Implicit Threading

### ■ OpenMP

- OpenMP provides several additional directives for running code regions in parallel, including parallelizing loops.
  - Assume we have two arrays **a** and **b** of size **N**. We wish to sum their contents and place the results in array **c**. We can have this task run in parallel by using the following code segment, which contains the compiler directive for parallelizing for loops:

```
#pragma omp parallel for
/* divides the work contained in the for loop among
the threads it has created */
for (i = 0; i < N; i++) {
    c[i] = a[i] + b[i];
}
```



## ■ Implicit Threading

### ■ OpenMP

#### ■ alg.13-5-openmp-matrixadd.c (1)

```
/* compiling: gcc -fopenmp */
/* omp works for matrix addition: when od = 12200 with 4 threads, system utility may
have 30% improvment */
#include <stdio.h>
#include <stdlib.h>
#include <sys/time.h>
#include <ctype.h>
#include <omp.h>
#define MAX_N 12288

int a[MAX_N][MAX_N], b[MAX_N][MAX_N];
int ans[MAX_N][MAX_N];
int od = 10;

void matrixadd(void)
{
    for(int i = 0; i < od; i++){
        for(int j = 0; j < od; j++){
            ans[i][j] = a[i][j] + b[j][j];
        }
    }

    return;
}
```

## ■ Implicit Threading

### ■ OpenMP

#### ■ [alg.13-5-openmp-matrixadd.c \(2\)](#)

```
int main(int argc, void *argv[])
{
    int i, iteration;

    if(argc > 1)
        od = atoi(argv[1]);
    if (od > MAX_N || od < 0)
        return 1;
    i = MAX_N;
    printf("MAX_N = %d, od = %d\n", i, od);
    for(int i = 0; i < od; i++)
        for(int j = 0; j < od; j++)
            a[i][j] = 20;
    for(int i = 0; i < od; i++)
        for(int j = 0; j < od; j++)
            b[i][j] = 30;

    /* with no omp */
    long start_us, end_us;
    struct timeval t;
    gettimeofday(&t, 0);
    start_us = (long)(t.tv_sec * 1000 * 1000) + t.tv_usec;
    matrixadd();
    gettimeofday(&t, 0);
    end_us = (long)(t.tv_sec * 1000 * 1000) + t.tv_usec;
    printf("Overhead time usec = %ld, with no omp\n", end_us-start_us);
}
```

## ■ Implicit Threading

### ■ OpenMP

#### ■ alg.13-5-openmp-matrixadd.c (3)

```
    /* with 2 threads */
    gettimeofday(&t, 0);
    start_us = (long)(t.tv_sec * 1000 * 1000) + t.tv_usec;
    #pragma omp parallel num_threads(2)
    {
        matrixadd();
    }
    gettimeofday(&t, 0);
    end_us = (long)(t.tv_sec * 1000 * 1000) + t.tv_usec;
    printf("Overhead time usec = %ld, omp thread = 2\n", end_us-start_us);

    /* with 4 threads */
    gettimeofday(&t, 0);
    start_us = (long)(t.tv_sec * 1000 * 1000) + t.tv_usec;
    #pragma omp parallel num_threads(4)
    {
        matrixadd();
    }
    gettimeofday(&t, 0);
    end_us = (long)(t.tv_sec * 1000 * 1000) + t.tv_usec;
    printf("Overhead time usec = %ld, omp thread = 4\n", end_us-start_us);

    return 0;
}
```



## ■ Implicit Threading

### ■ OpenMP

#### ■ alg.13-5-openmp-matrixadd.c (3)

/\* with 2 threads \*/

```
isscgy@ubuntu:/mnt/os-2020$ gcc alg.13-5-openmp-matrixadd.c -fopenmp
isscgy@ubuntu:/mnt/os-2020$ ./a.out 10
MAX_N = 12288, od = 10
Overhead time usec = 17, with no omp
Overhead time usec = 96, omp thread = 2
Overhead time usec = 2779, omp thread = 4
isscgy@ubuntu:/mnt/os-2020$ ./a.out 1000
MAX_N = 12288, od = 1000
Overhead time usec = 15797, with no omp
Overhead time usec = 14968, omp thread = 2
Overhead time usec = 13119, omp thread = 4
isscgy@ubuntu:/mnt/os-2020$ ./a.out 12200
MAX_N = 12288, od = 12200
Overhead time usec = 6487102, with no omp
Overhead time usec = 4929145, omp thread = 2
Overhead time usec = 3538193, omp thread = 4
isscgy@ubuntu:/mnt/os-2020$
```



## ■ Threading Issues

- Some of the issues to consider in designing multithreaded programs.
  - Semantics of `fork()` and `exec()` system calls
  - Signal handling
  - Thread cancellation
    - Asynchronous or deferred.
  - Thread-local storage
  - Scheduler activations.

## ■ `fork()` and `exec()`

- The `fork()` system call is used to create a separate, duplicate **process**. But the semantics of the `fork()` and `exec()` system calls change in a multithreaded program.
  - If one thread in a program calls `fork()`, the new process may
    - duplicate all threads, or
    - only the thread that invoked the `fork()` system call (in Ubuntu).
      - It may be in high risk.
  - Some UNIX systems have two versions of `fork()`.
- The `exec()` system call typically works in the way that if a thread invokes the `exec()` system call, the program specified in the parameter to `exec()` will replace the calling process including all of its threads.
  - If `exec()` is called immediately after forking, the forked process needs only duplicating the calling thread.
    - duplicating all threads is unnecessary as the program specified in the parameters to `exec()` will replace the calling process.
  - Otherwise, the forked process does not call `exec()` after forking, it should duplicate all threads of the calling process.
- **Suggestion: avoid using `fork()` in multithreading.**



## ■ fork() and exec()

### ■ fork() and Multithreading

#### ■ [alg.13-6-fork-pthread-demo1.c \(1\)](#)

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <unistd.h>
#include <sys/wait.h>

int i = 0;

static void *thread_worker(void *args)
{
    while (1) {
        printf("%d\n", i);
        /* will print '0' only, by thread_worker of parent main() */
        sleep(1);
    }

    pthread_exit(0);
}
```

## ■ fork() and exec()

### ■ fork() and Multithreading

#### ■ [alg.13-6-fork-pthread-demo1.c \(2\)](#)

```
int main(void)
{
    pthread_t ptid;
    pthread_create(&ptid, NULL, &thread_worker, NULL);

    pid_t pid = fork(); /* child duplicates parent's main thread only,
without thread_worker() */
    if(pid == 0){ /* child process */
        i = 1;
        printf("In child\n");
        system("ps -l -T");
        while (1) ;
        exit (0);
    }

    wait(&pid);
    printf("In parent\n");
    system("ps -l -T");
    while (1) ;

    return 0;
}
```

## ■ fork() and exec()

```

isscg@ubuntu:/mnt/os-2020$ gcc alg.13-6-fork-pthread-demo1.c -pthread
isscg@ubuntu:/mnt/os-2020$ ./a.out
in child
0
F S  UID  PID  SPID  PPID  C PRI  NI ADDR SZ WCHAN  TTY          TIME CMD
0 S  1000  1856  1856  1846  0  80   0 - 6124 wait  pts/0      00:00:04 bash
0 S  1000  19526  19526  1856  0  80   0 - 20107 wait  pts/0      00:00:00 a.out
1 S  1000  19526  19527  1856  0  80   0 - 20107 hrtim  pts/0      00:00:00 a.out
1 S  1000  19528  19528  19526  0  80   0 - 3723 wait  pts/0      00:00:00 a.out
0 S  1000  19529  19529  19528  0  80   0 - 1158 wait  pts/0      00:00:00 sh
0 R  1000  19530  19530  19529  0  80   0 - 7667 -      pts/0      00:00:00 ps
in parent
F S  UID  PID  SPID  PPID  C PRI  NI ADDR SZ WCHAN  TTY          TIME CMD
0 S  1000  1856  1856  1846  0  80   0 - 6124 wait  pts/0      00:00:04 bash
0 S  1000  19526  19526  1856  0  80   0 - 20107 wait  pts/0      00:00:00 a.out
1 S  1000  19526  19527  1856  0  80   0 - 20107 hrtim  pts/0      00:00:00 a.out
0 S  1000  19531  19531  19526  0  80   0 - 1158 wait  pts/0      00:00:00 sh
0 R  1000  19532  19532  19531  0  80   0 - 7667 -      pts/0      00:00:00 ps
0
0
0
0
^C
isscg@ubuntu:/mnt/os-2020$

```

printing 0s only, by thread\_worker of parent main() pro

spid=19526, pid=19526: parent main() thread  
 spid=19527, pid=19526: thread\_worker  
 spid=19528, pid=19526: forked child pro



## ■ fork() and exec()

### ■ fork() and Multithreading

#### ■ [alg.13-7-fork-pthread-demo2.c \(1\)](#)

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <unistd.h>

int i = 0;
static void *thread_worker(void *args)
{
    pid_t pid = fork(); /* the forked child takes thread_worker as main thread. This
may cause unexpect behaviours in synchronization or signal handling */
    if(pid < 0 )
        return (void *)EXIT_FAILURE;
    if(pid == 0) { /* child pro */
        i = 1;
        printf("In thread_worker's forked child\n");
        system("ps -l -T");
    }

    sleep(2);
    while (1) {
        printf("%d\n", i);
        /* will print '0' by thread_worker of parent main(); '1' by forked child pro */
        sleep(2);
    }
    pthread_exit(0);
}
```

## ■ fork() and exec()

### ■ fork() and Multithreading

#### ■ [alg.13-7-fork-pthread-demo2.c \(2\)](#)

```
int main(void)
{
    pthread_t tid;
    pthread_create(&tid, NULL, &thread_worker, NULL);

    sleep(2) ;
    printf("In start main()\n");
    system("ps -l -T");

    while (1) ;

    pthread_join(tid, NULL);

    return EXIT_SUCCESS;
}
```

## ■ fork() and exec()

### ■ fork() and Multithreading

#### ■ alg.13-7-fork-pthread-demo2.c (2)

```

isscgy@ubuntu:/mnt/os-2020$ gcc alg.13-7-fork-pthread-demo2.c -pthread
isscgy@ubuntu:/mnt/os-2020$ ./a.out
in thread_worker's forked child
0 S 1000 19605 19605 1856 0 80 0 - 3723 hrttime pts/0 00:00:00 a.out
1 S 1000 19605 19606 1856 0 80 0 - 3723 hrttime pts/0 00:00:00 a.out
1 S 1000 19607 19607 19605 0 80 0 - 3723 wait pts/0 00:00:00 a.out
in start main()
0
1
0 S 1000 19605 19605 1856 0 80 0 - 3723 wait pts/0 00:00:00 a.out
1 S 1000 19605 19606 1856 0 80 0 - 3723 hrttime pts/0 00:00:00 a.out
1 S 1000 19607 19607 19605 0 80 0 - 3723 hrttime pts/0 00:00:00 a.out
0
1
0
1
0
1
^C
isscgy@ubuntu:/mnt/os-2020$

```

printing 0s by thread\_worker of start main(); 1s by forked child

spid=19605, pid=19605: start main() pro  
 spid=19606, pid=19605: thread\_worker  
 spid=19607, pid=19607: forked child pro



## ■ fork() and exec()

### ■ fork() and Multithreading

#### ■ [alg.13-7-fork-pthread-demo2.c \(2\)](#)

```
int main(void)
{
    pthread_t ptid;
    pthread_create(&ptid, NULL, &thread_worker, NULL);

    sleep(2) ;
    printf("in start main()\n");
    system("ps -l -T");

    return 1; /* what will happen? you may have to pkill the forked child process */

    while (1) ;

    pthread_join(tid, NULL);

    return EXIT_SUCCESS;
}
```



## ■ fork() and exec()

### ■ fork() and Multithreading

#### ■ alg.13-7-fork-pthread-demo2.c (2)

```
isscgy@ubuntu:/mnt/os-2020$ ./a.out
in thread_worker's forked child
0 S 1000 19647 19647 1856 0 80 0 - 3723 hrttime pts/0 00:00:00 a.out
1 S 1000 19647 19648 1856 0 80 0 - 3723 hrttime pts/0 00:00:00 a.out
1 S 1000 19649 19649 19647 0 80 0 - 3723 wait pts/0 00:00:00 a.out
0
in start main()
1
0 S 1000 19647 19647 1856 0 80 0 - 20107 wait pts/0 00:00:00 a.out
1 S 1000 19647 19648 1856 0 80 0 - 20107 hrttime pts/0 00:00:00 a.out
1 S 1000 19649 19649 19647 0 80 0 - 3723 hrttime pts/0 00:00:00 a.out
isscgy@ubuntu:/mnt/os-2020$ 1
1
1
1
ps
  PID TTY          TIME CMD
  1856 pts/0        00:00:04 bash
  19649 pts/0        00:00:00 a.out
  19656 pts/0        00:00:00 ps
isscgy@ubuntu:/mnt/os-2020$ 1
1
^C
isscgy@ubuntu:/mnt/os-2020$ 1
pk1
ill -f a1
.out
isscgy@ubuntu:/mnt/os-2020$
```

spid=19647, pid=19647: start main() pro  
spid=19648, pid=19647: thread\_worker  
spid=19649, pid=19649: forked child pro

The forked child is still working, printing  
1s. It cannot response to ^C and we have  
to pkill it.  
`$pkill -f a.out`

## ■ Signal Handling

- A signal is used in UNIX systems to notify a process that a particular event has occurred.
  - A signal may be received either synchronously or asynchronously
- All signals should follow the pattern:
  - A signal is generated by the occurrence of a particular event.
  - The signal is delivered to a process.
  - Once delivered, the signal **must** be handled.
- Signals are handled by one of these two signal handlers
  - Default handler run by kernel
  - User-defined handler that can override default handler.
- For single-threaded, one signal is delivered to a process.
- Where should a signal be delivered for multi-threaded?
  - Deliver the signal to the thread to which the signal applies.
  - Deliver the signal to every thread in the process.
  - Deliver the signal to certain threads in the process.
  - Assign a specific thread to receive all signals for the process.
  - Discuss later with `CLONE_THREAD` flag.

## ■ Signal Handling

### ■ Example

- Standard UNIX function for delivered a signal to a specified process

```
int kill(pid_t pid, int signal)
```

- POSIX Pthreads function pthread\_kill() allows a signal to be delivered to a specified thread

```
int pthread_kill(pthread_t ptid, int signal)
```

## ■ Signal Handling

### ■ Linux

```
iisscgy@ubuntu:/mnt/hgfs/os-2020$ kill -l
 1) SIGHUP      2) SIGINT      3) SIGQUIT     4) SIGILL      5) SIGTRAP
 6) SIGABRT     7) SIGBUS     8) SIGFPE     9) SIGKILL     10) SIGUSR1
11) SIGSEGV    12) SIGUSR2   13) SIGPIPE   14) SIGALRM    15) SIGTERM
16) SIGSTKFLT  17) SIGCHLD  18) SIGCONT   19) SIGSTOP    20) SIGTSTP
21) SIGTTIN    22) SIGTTOU  23) SIGURG    24) SIGXCPU    25) SIGXFSZ
26) SIGVTALRM  27) SIGPROF  28) SIGWINCH  29) SIGIO      30) SIGPWR
31) SIGSYS     34) SIGRTMIN  35) SIGRTMIN+1 36) SIGRTMIN+2 37) SIGRTMIN+3
38) SIGRTMIN+4 39) SIGRTMIN+5 40) SIGRTMIN+6 41) SIGRTMIN+7 42) SIGRTMIN+8
43) SIGRTMIN+9 44) SIGRTMIN+10 45) SIGRTMIN+11 46) SIGRTMIN+12 47) SIGRTMIN+13
48) SIGRTMIN+14 49) SIGRTMIN+15 50) SIGRTMAX-14 51) SIGRTMAX-13 52) SIGRTMAX-12
53) SIGRTMAX-11 54) SIGRTMAX-10 55) SIGRTMAX-9 56) SIGRTMAX-8 57) SIGRTMAX-7
58) SIGRTMAX-6 59) SIGRTMAX-5 60) SIGRTMAX-4 61) SIGRTMAX-3 62) SIGRTMAX-2
63) SIGRTMAX-1 64) SIGRTMAX
iisscgy@ubuntu:/mnt/hgfs/os-2020$
```

- 1-31号是常规信号（非实时信号），没有阻塞队列，多次产生时只有一次记录。
- 32-64号是实时信号，支持排队策略。



### ■ Signal Handling

#### ■ Linux – Data Structures

```
struct sigaction {
    union {
        __sig_handler_t _sa_handler;
        void (*_sa_sigaction)(int, struct siginfo *, void *);
    } _u;
    sigset_t sa_mask;
    unsigned long sa_flags;
}

struct sigaction {
    void (*sa_handler)(int); /* func name | SIG_IGN | SIG_DFL */
    void (*sa_sigaction)(int, siginfo_t *, void *); /* seldom
used */
    sigset_t sa_mask; /* 64bit mask, valid only for the action */
    int sa_flags; /* 0: mask itself */
    void (*sa_restorer)(void); /* abandoned */
};
```

- SIG\_IGN and SIG\_DFL are handles defined in the kernel which means ignoring signals or dealing with by default.



## ■ Signal Handling

### ■ Linux – Signal Registering

```
int sigaction(int signum, const struct sigaction *act, struct
sigaction *oldact);
/* signum: the registering signal, other than 9 SIGKILL and 19
SIGSTOP; *act: the new sigaction; *oldact: the old sigaction
saved, NULL if careless */
```

### ■ Linux – Setting `sigset_t sa_mask`

```
int sigemptyset(sigset_t *set) /* clear all */
int sigfillset(sigset_t *set) /* set all */
int sigaddset(sigset_t *set, int signum) /* add */
int sigdelset(sigset_t *set, int signum) /* delete */
int sigismember(sigset_t *set, int signum) /* query */
```

### ■ Linux – Sending Signals

```
int sigqueue(pid_t pid, int signo, const union sigval value);
int kill(pid_t pid, int signal)
int tgkill(int tgid, int tid, int sigal);
int tkill(int tid, int sigal); /* obsolete predecessor to tgkill()
*/
```





### ■ Signal Handling

#### ■ alg.13-8-sigaction-demo.c (1)

```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <signal.h>

void my_handler(int signo) /* user signal handler */
{
    printf("\nhere is my_handler");
    printf("\nsignal caught: signo = %d", signo);
    printf("\nCtrl+\ is masked");
    printf("\nsleeping for 10 seconds ... \n");
    sleep(10);
    printf("my_handler finished\n");
    printf("after returned to the main(), Ctrl+\ is unmasked\n");
    return;
}
```





## ■ Signal Handling

### ■ alg.13-8-sigation-demo.c (2)

```
int main(void)
{
    int ret;
    struct sigaction newact;

    newact.sa_handler = my_handler; /* set the user-defined handler */
    sigemptyset(&newact.sa_mask); /* clear the mask */
    sigaddset(&newact.sa_mask, SIGQUIT); /* sa_mask, set signo=3
(SIGQUIT:Ctrl+\) */
    newact.sa_flags = 0; /* default */

    printf("now start catching Ctrl+c\n");

    ret = sigaction(SIGINT, &newact, NULL); /* register signo=2
(SIGINT:Ctrl+C) */
    if (ret == -1) {
        perror("sigaction()");
        exit(1);
    }

    while (1);
    return 0;
}
```



## ■ Signal Handling

### ■ alg.13-8-sigaction-demo.c (2)

```
int main(void)
```

```
iisscgy@ubuntu:/mnt/os-2020$ gcc alg.13-8-sigaction-demo.c
iisscgy@ubuntu:/mnt/os-2020$ ./a.out
now start catching Ctrl+c
^C
here is my_handler
signal caught: signo = 2
Ctrl+\ is masked
sleeping for 10 seconds ...
^\\my_handler finished
after returned to the main(), Ctrl+\ is unmasked
Quit (core dumped)
iisscgy@ubuntu:/mnt/os-2020$
```

```
    perror("sigaction()");
    exit(1);
}
```

```
while (1);
return 0;
```

```
}
```

signo 2 (Ctrl+\) is masked and pending until my\_handler finished, causing core dumped.

## ■ Thread Cancellation

- *Thread Cancellation* involves terminating a thread (here called the *target thread*) *before* it has completed.
- Cancellation of a target in two approaches:
  - Asynchronous cancellation
    - terminates the target thread immediately.
  - Deferred cancellation
    - allows the target thread to periodically check if it should be cancelled.
- The difficulty with cancellation occurs in situations where resources have been allocated to a canceled thread or where a thread is canceled while in the midst of updating data it is sharing with other threads.



### ■ Thread Cancellation

#### ■ Example

- In POSIX Pthreads, the `pthread_cancel()` function cancels a specified thread. The identifier of the target thread is passed as a parameter to the function.
- The following code illustrates creating—and then canceling—a thread.

```
pthread_t ptid;  
    /* create the thread */  
pthread_create(&ptid, 0, &thread_worker, NULL);  
    . . .  
    /* cancel the thread */  
pthread_cancel(ptid);
```

## ■ Thread Cancellation

### ■ Pthreads Cancellation Modes

- Pthreads supports three cancellation modes. Each mode is defined as a state and a type.
  - Off Mode
    - State = PTHREAD\_CANCEL\_DISABLE;
    - no Type defined
  - **Deferred Mode**
    - State = PTHREAD\_CANCEL\_ENABLE;
    - Type = PTHREAD\_CANCEL\_DEFERRED
  - Asynchronous Mode
    - State = PTHREAD\_CANCEL\_ENABLE;
    - Type = PTHREAD\_CANCEL\_ASYNCHRONOUS
- The default cancellation type is deferred cancellation. Cancellation request is pending until the thread reaches a ***cancellation-point***.

## ■ Thread Cancellation

### ■ Pthreads Cancellation Points

#### ■ *Cancellation-point*

- By POSIX, the functions of `pthread_join()`, `pthread_testcancel()`, `pthread_cond_wait()`, `pthread_cond_timedwait()`, `sem_wait()`, `sigwait()` and the system calls `read()`, `write()` that may cause system blocking are Cancellation-point.

#### ■ The function

```
void pthread_testcancel(void)
```

can be used for setting a cancellation point, or cancelling the thread whose cancellation request is deferred.

#### ■ Example:

```
pthread_testcancel();  
retcode = read(fd, buffer, length);  
    /* be set as a cancellation point */  
pthread_testcancel();
```



### ■ Thread Cancellation

#### ■ Pthreads functions with respect to thread cancellation

- `int pthread_create(pthread_t *restrict ptid, const pthread_attr_t *restrict_attr, void*(*start_rtn)(void*), void *restrict arg);`
- `int pthread_cancel(pthread_t ptid);`
- `void pthread_testcancel(void);`
- `int pthread_setcancelstate(int state, int *oldstate);`
- `int pthread_setcanceltype(int type, int *oldtype);`
- `void pthread_cleanup_push(void (*routine) (void *), void *arg);`
- `void pthread_cleanup_pop(int execute);`