

---

# Memory Segmentation & Paging

## Operating Systems

---

School of Data & Computer Science  
Sun Yat-sen University

Lecture Notes: [os\\_sysu@163.com](mailto:os_sysu@163.com)  
Instructor: Guoyang Cai  
email: [isscgymail@mail.sysu.edu.cn](mailto:isscgymail@mail.sysu.edu.cn)





## ■ Contents


- Simple Segmentation
- Simple Paging
- Address Translation
- Page Tables
- Examples

## ■ Segmentation Concepts

- Segmentation is memory-management scheme that supports **user view** of memory.
- A program is a collection of segments.
  - A segment is a logical unit such as:
    - main program
    - procedure
    - function
    - method
    - object
    - local variables, global variables
    - common block
    - stack
    - symbol table
    - arrays.

## ■ Segmentation Concepts

- User View of a Program.



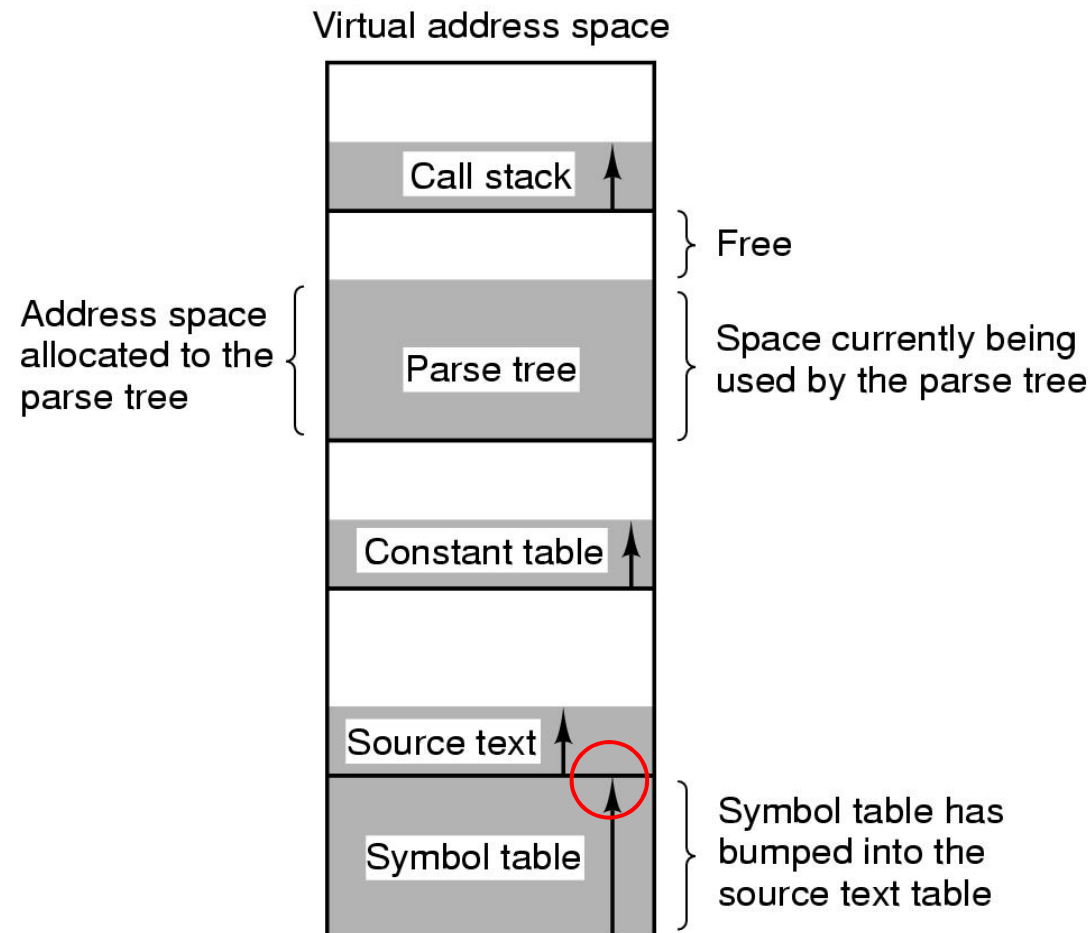
Logical address space.

## ■ Segmentation Concepts

- User View of a Program – Example.
  - A compiler has many tables that are built up as compilation proceeds, possibly including:
    - The source text being saved for the printed listing (on batch systems).
    - The symbol table – the names and attributes of variables.
    - The table containing integer, floating-point, constants used.
    - The parse tree (语法树), the syntactic analysis of the program.
    - The stack used for procedure calls within the compiler.

## Segmentation Concepts

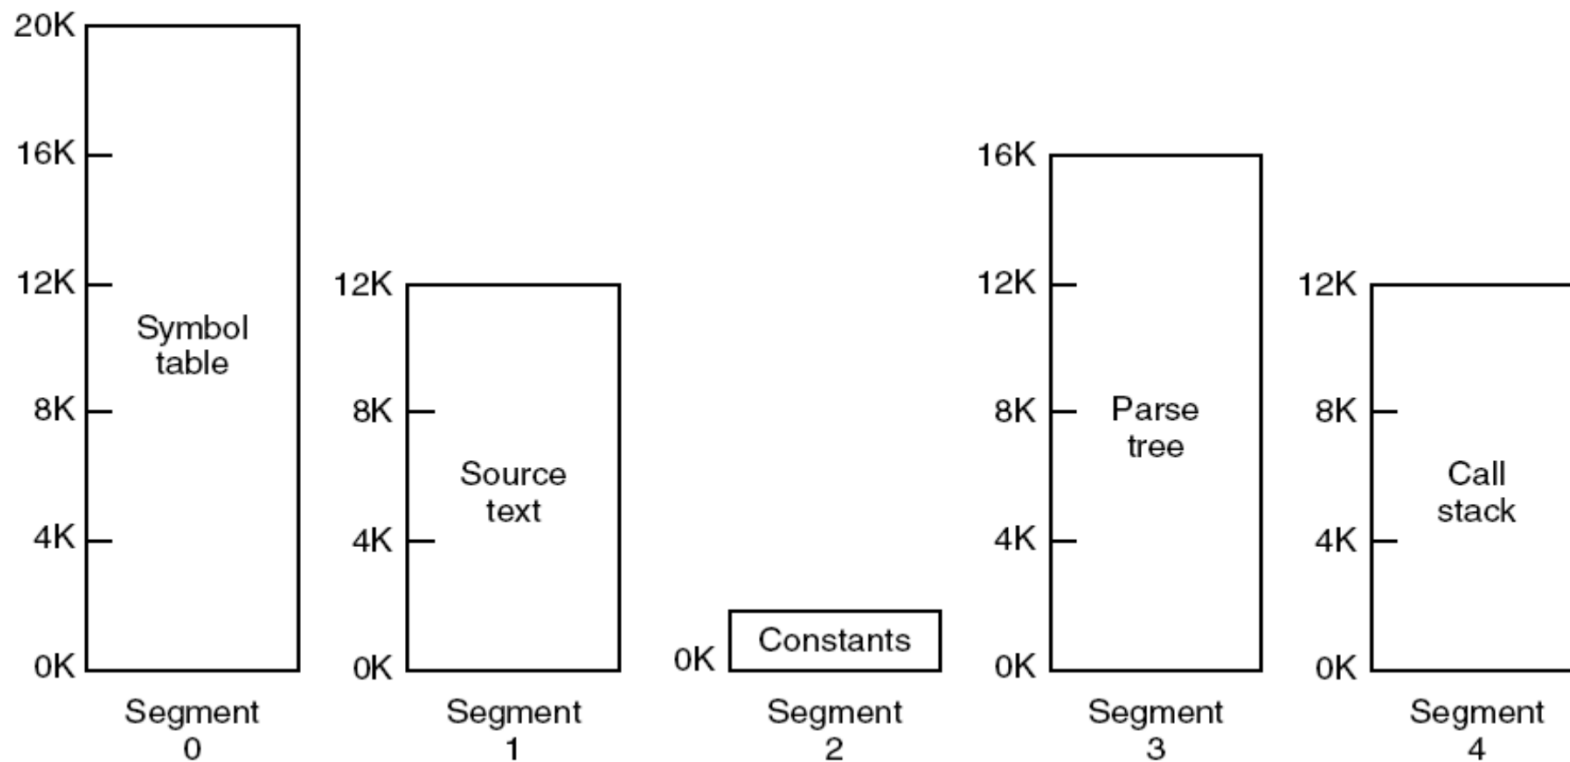
- One-dimensional address space
  - In a **one-dimensional address space** with growing tables, one table may bump into another.



## ■ Segmentation Concepts

### ■ Segmentation Solution

- Segmentation allows each table to grow or shrink independently of the other tables.




## ■ Dynamics of Simple Segmentation


- Each program is subdivided into blocks of non-equal size called segments.
- When a process gets loaded into main memory, its different segments can be located anywhere.
- Each segment is fully packed with instructions/data; **no internal fragmentation.**
- There is external fragmentation; it is reduced when using small segments.
- In contrast with paging, segmentation is **visible to the programmer.**
  - provided as a convenience to organize logically programs (e.g., data in one segment, code in another segment).
  - must be aware of segment size limit.
- OS maintains a **segment table** (段表) for each process. Each entry contains:
  - the starting *physical addresses* of that segment
  - the length of that segment for protection.



## ■ Logical view of simple segmentation



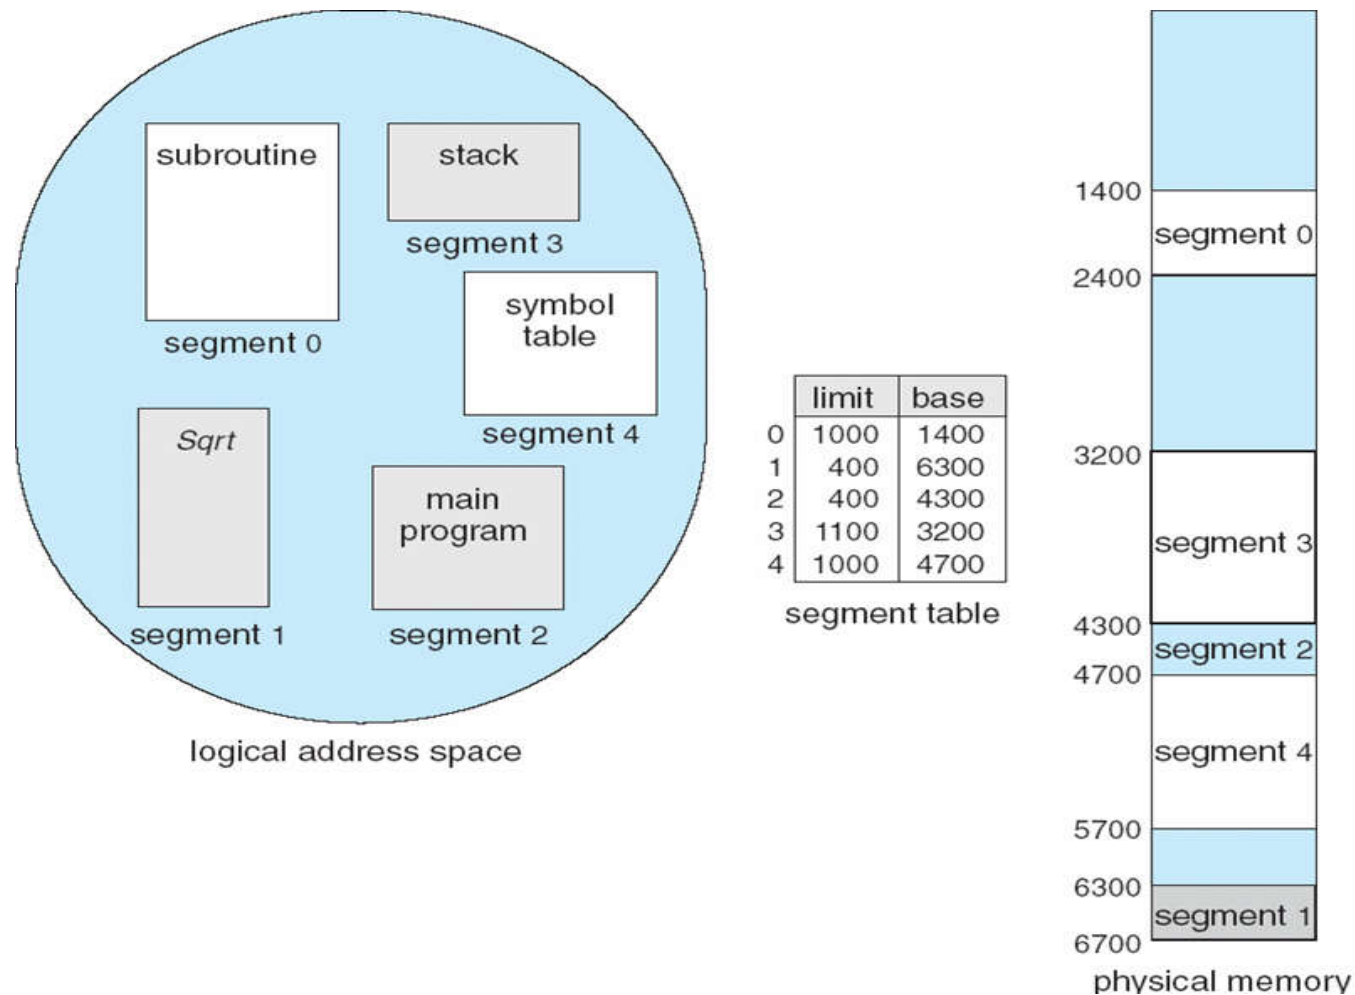
user space



physical memory space

## Address Translation in segmentation

### Example.



## ■ Segmentation Architecture

- A logical address consists of a two-tuple:

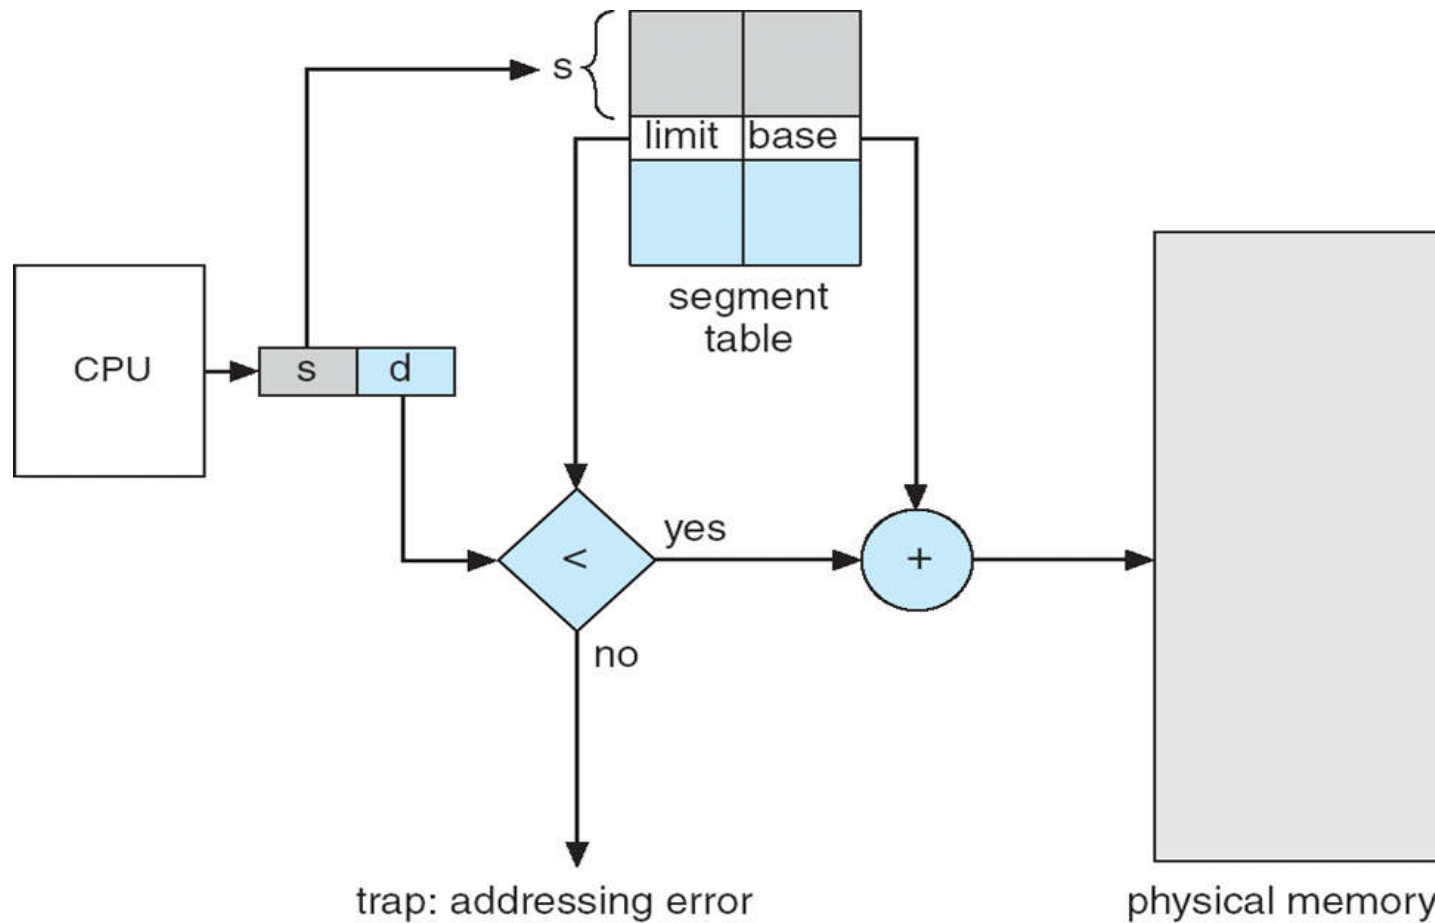
$\langle \text{segment\_number}, \text{offset} \rangle$ ,

段号      偏移

- *Segment table* – maps two-dimensional physical addresses; each table entry has:
  - *base* – contains the starting physical address where the segments reside in memory.
  - *limit* – specifies the length of the segment.
- *Segment-table base register (STBR)* points to the location of the segment table in memory.
- *Segment-table length register (STLR)* indicates the *number of segments* used by a program
  - segment-number  $s$  is legal if  $s < \text{STLR}$ .

## Segmentation Architecture

- Address Translation Architecture - Segmentation Hardware.



## ■ Segmentation Architecture

### ■ Address Translation in Segmentation

- When a process enters the Running state, a dedicated register gets loaded with the starting address of the process's segment table.
- Presented with a logical address

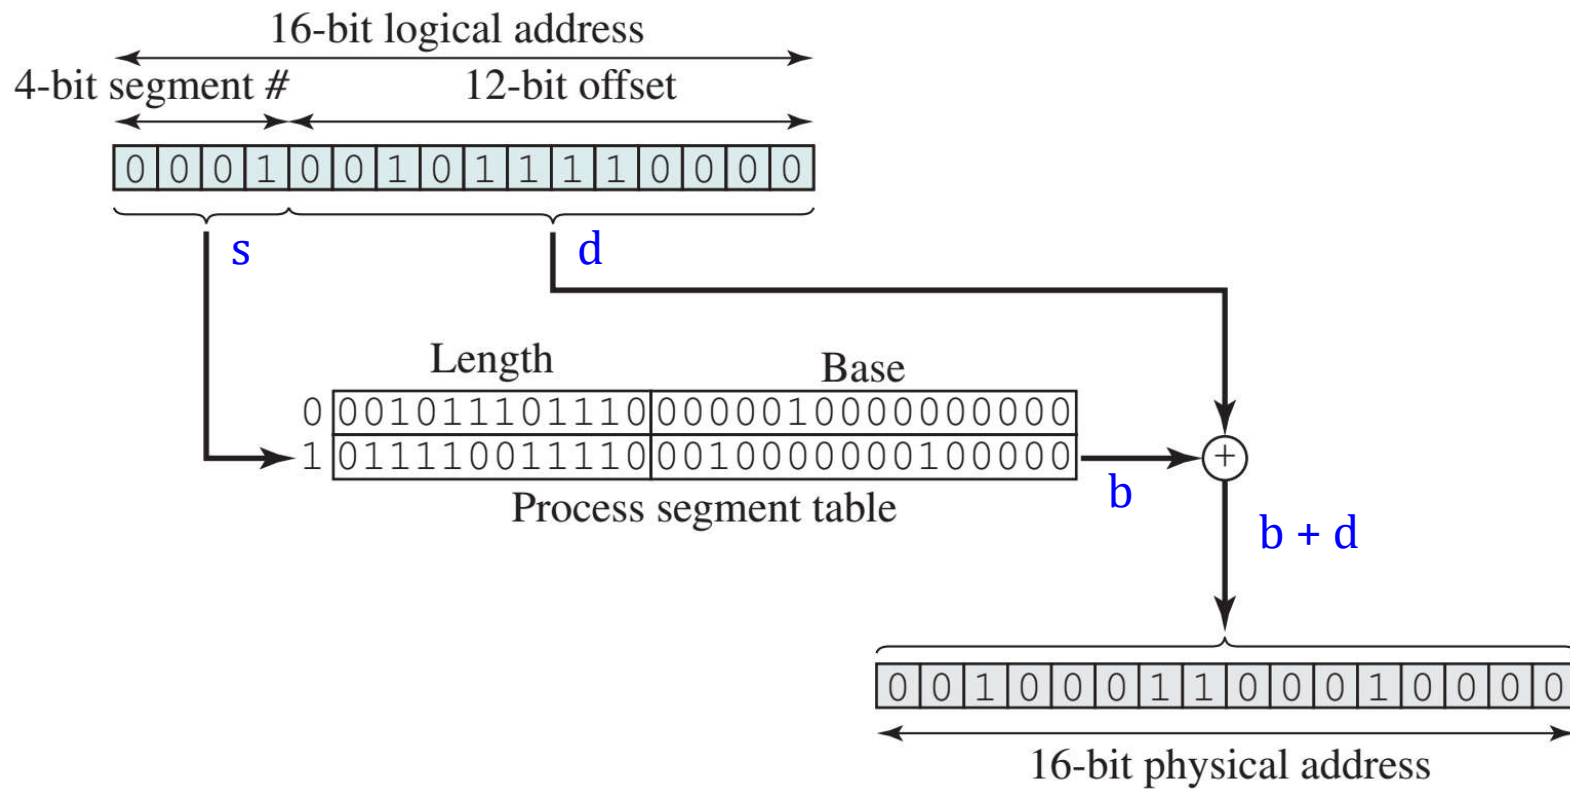
$$\langle \text{segment\_number, offset} \rangle = \langle s, d \rangle,$$

the CPU indexes (with  $s$ ) the segment table to obtain the starting physical address  $b$  and the length limit  $l$  of that segment.

- The physical address is obtained by adding  $d$  to  $b$ .
  - The hardware also compares the offset/displacement  $d$  with the limit  $l$  of that segment to determine if the address is valid.

## Segmentation Architecture

- Address Translation in Segmentation
  - Example.



## ■ Segmentation Architecture

- Protection
  - With each entry in segment table associate:
    - validation bit = 0  $\Rightarrow$  illegal segment
    - read/write/execute privileges.
- Protection bits associated with segments; code sharing occurs at segment level.
- Since segments vary in length, memory allocation is a dynamic storage-allocation problem.

## ■ Shared Segments

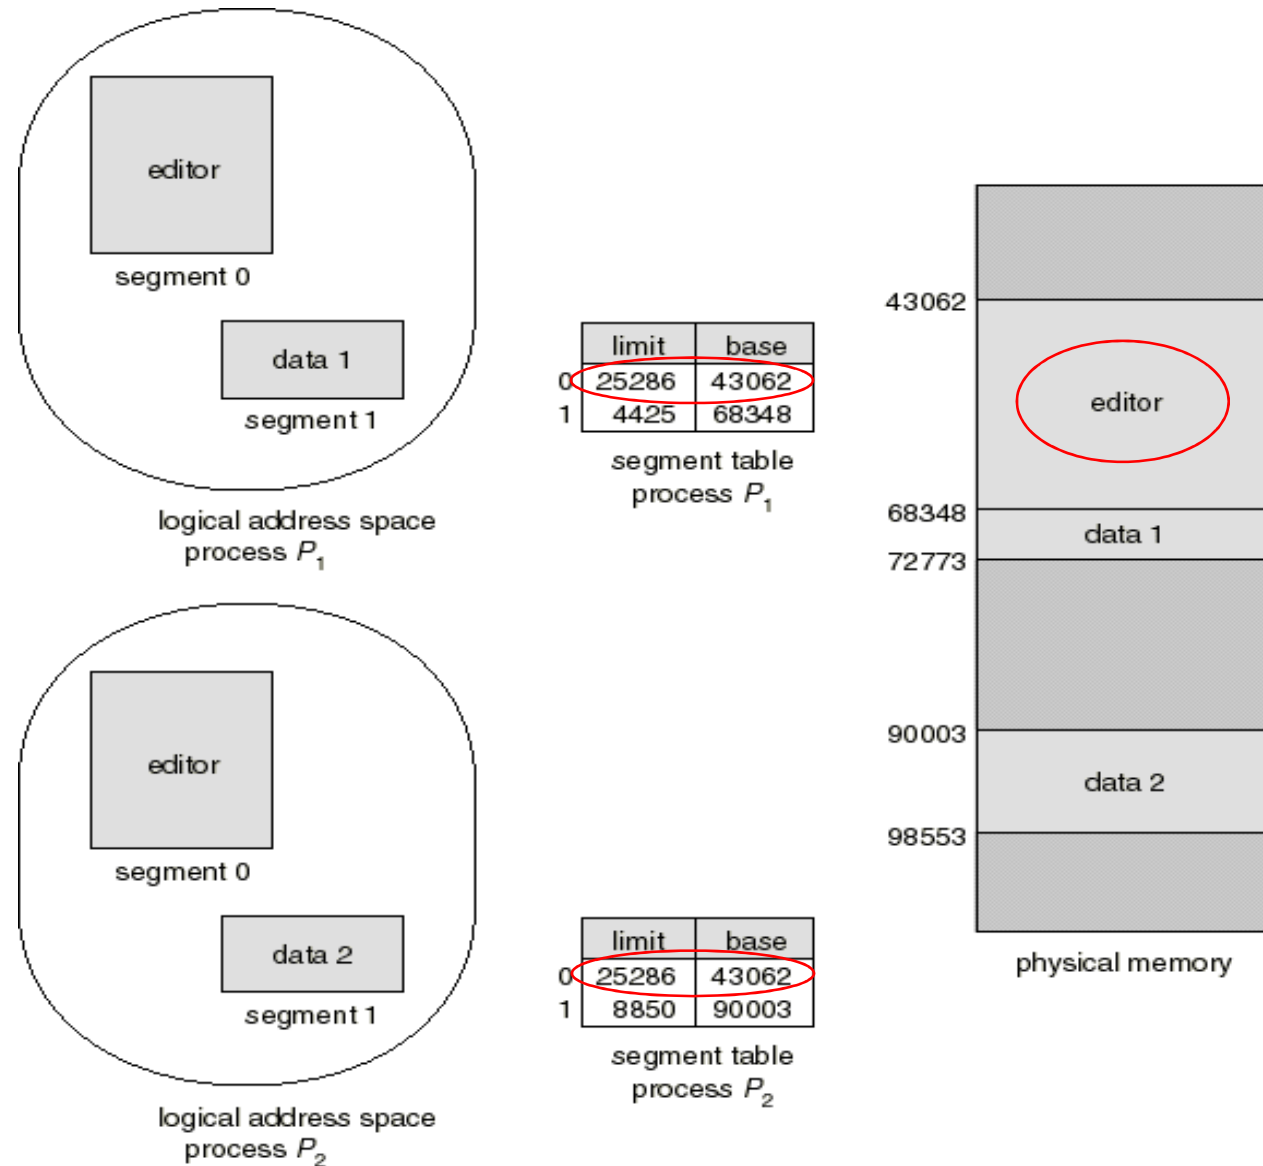
### ■ Sharing in Segmentation Systems

- Segments are shared when entries in the segment tables of 2 different processes point to the same physical locations.
- Example.
  - The same code of a text editor can be shared by many users.
  - Only one copy is kept in main memory.
- But each user would still need to have its own private data segment.



## Shared Segments

### Example.

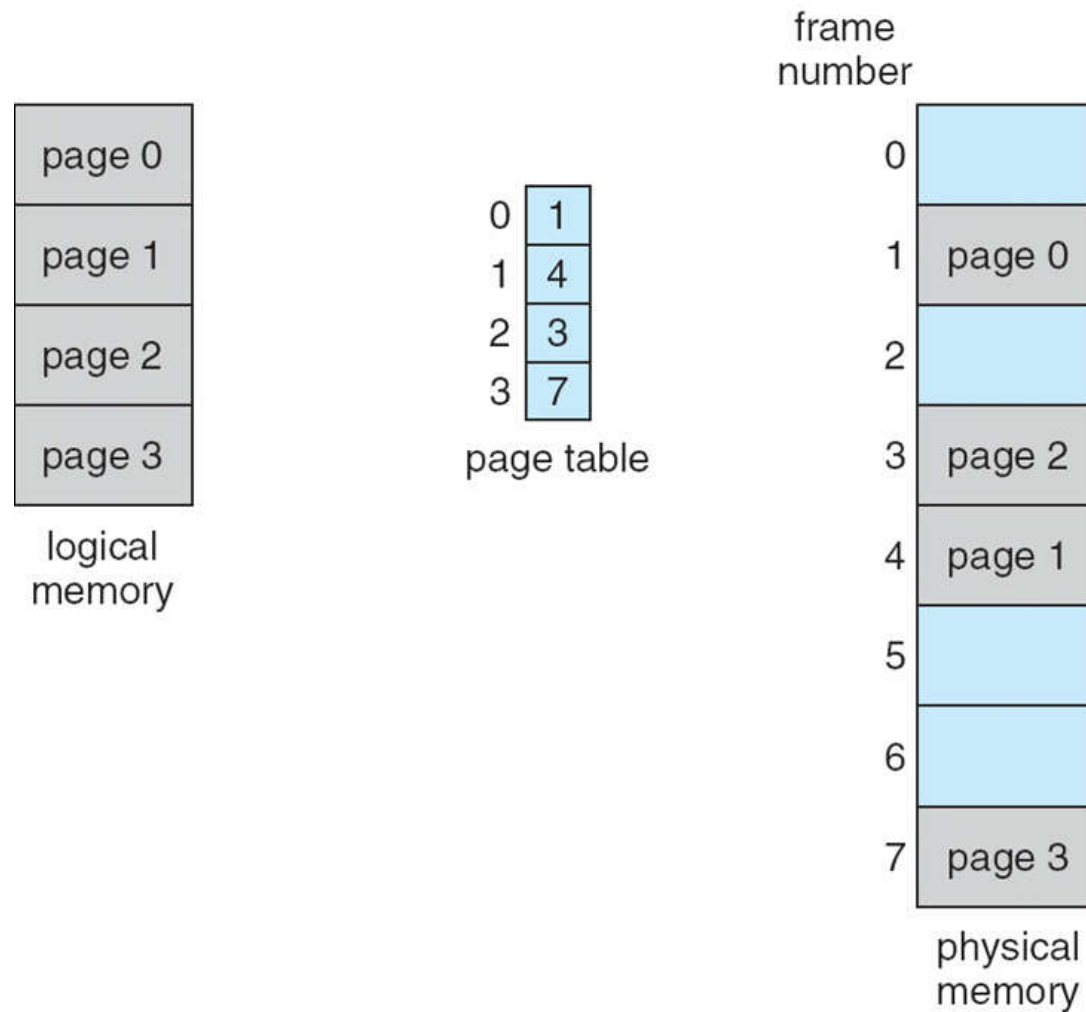


## ■ Simple Paging Concepts

- Idea: Physical address space of a process can be **noncontiguous**; process is allocated physical memory whenever the latter is available in size.
  - Avoids **external fragmentation**.
  - Avoids problem of varying sized memory chunks.
- Divide **physical memory** into fixed-sized chunks/blocks called **frames** (页帧, 页框, 物理页), size of power of 2, usually between 512 bytes and 16 MB.
- Divide **logical memory** into blocks of the same size as physical frames, called **pages** (页, 页面, 逻辑页).
- Process pages can thus be assigned to any free frames in main memory; a process does not need to occupy a contiguous portion of physical memory. To run a program of size  $n$  pages, we need to find  $n$  free frames of physical memory and load the program.
  - So we need to use a **free-frame list** to keep track of all free frames in physical memory.
  - And we need to set up a **page table** (页表) to translate logical pages to physical frames.
- **Internal fragmentation** possible only for the page at the end of program.


## Simple Paging Concepts

### Paging Example.



## Simple Paging Concepts

- Free-frame list example.



## Simple Paging Concepts

- OS now needs to maintain in main memory a *page table* for each process.
  - Each entry of a page table consists of the frame number where the corresponding page is physically located.
  - The corresponding page table is indexed by the page number to obtain the frame number.
- A *free frame table/list*, of available frames, is also maintained.
  - Example.

0	0
1	1
2	2
3	3

Process A  
page table

0	—
1	—
2	—

Process B  
page table

0	7
1	8
2	9
3	10

Process C  
page table

0	4
1	5
2	6
3	11
4	12

Process D  
page table

13
14

Free-frame  
list

## ■ Calculating Internal Fragmentation

### ■ Example.

- Page/frame size: 2,048 bytes
- Suppose some process has size: 72,766 bytes
  - $72766 = 2048 \times 35 + 1,086$ .
  - Need to allocate 36 frames to the process.
  - Internal fragmentation:  $2,048 - 1,086 = 962$  (bytes)
- Worst case fragmentation: 2048 - 1 (byte)
  - For processes having size of  $2048p + 1$  bytes.
- On average internal fragmentation: 1/2 frame size.

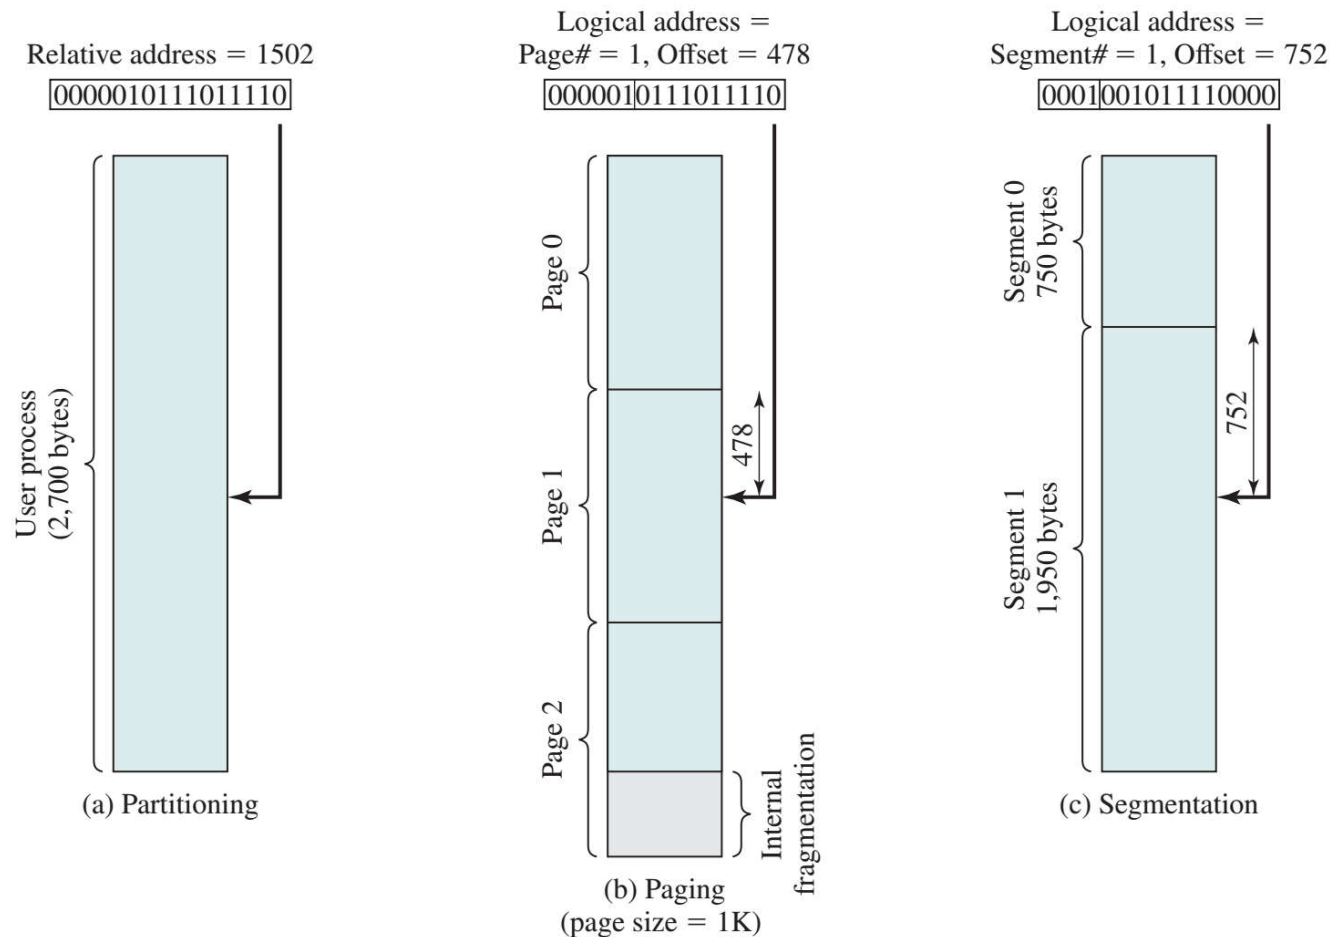
### ■ So small frame sizes desirable?

- But each page table entry takes memory to track.
- Page sizes growing over time:
  - Linux page sizes: 4KB / 2MB
    - `$getconf PAGESIZE`
    - `$cat /proc/meminfo | grep Huge`
  - Solaris page sizes: 8KB / 4MB.



## Logical Address in Paging

- The logical address becomes a relative address when the page size is a power of 2.
- Comparison of Partitioning, Paging and Segmentation.



## ■ Logical Address in Paging

- The logical address becomes a relative address when the page size is a power of 2.
  - Example.
    - If 16-bit addresses are used and page size = 1K, we need 10 bits for offset and have 6 bits available for page number.
    - Then the 16 bit address, obtained with the 10 least significant bits as offset and 6 most significant bits as page number, is a location relative to the beginning of the process.



## ■ Logical Address in Paging

- Within each program, each logical address must consist of a page number and an offset/displacement within the page

$$\langle \text{page\_number}, \text{offset} \rangle = \langle p, d \rangle.$$

- A dedicated register always holds the starting physical address of the page table of the currently running process.
- Presented with the logical address


$$\langle \text{page\_number}, \text{offset} \rangle = \langle p, d \rangle,$$

the processor accesses the page table to obtain the physical address

$$\langle \text{frame\_number}, \text{offset} \rangle = \langle f, d \rangle.$$

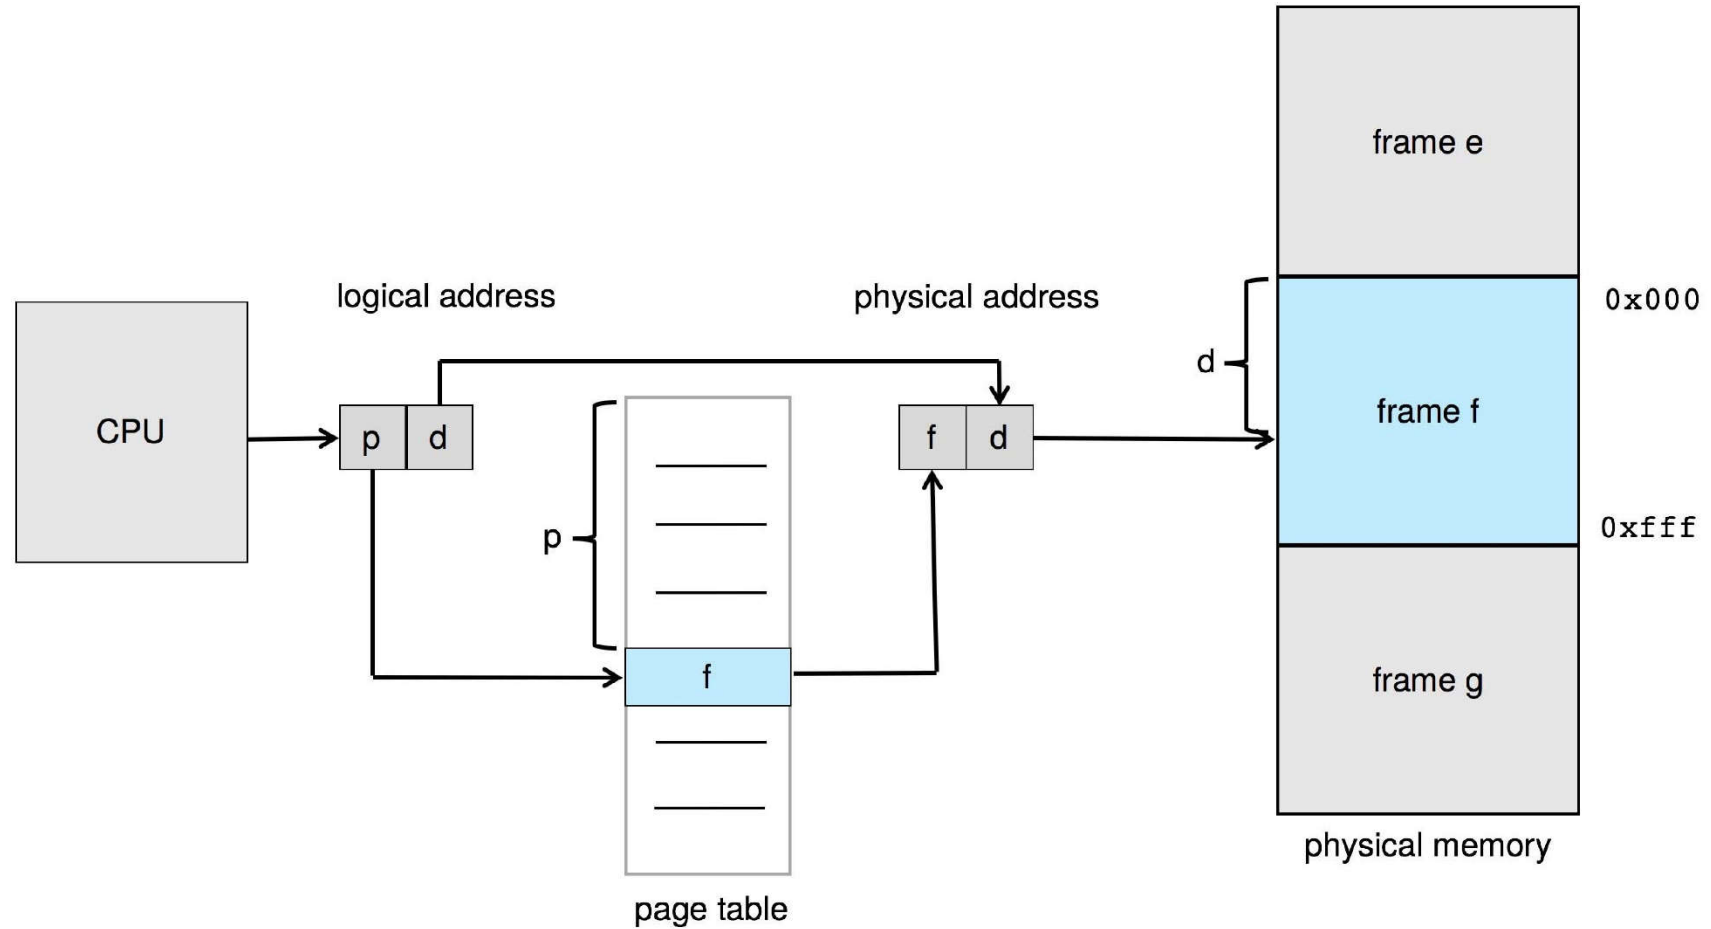
## ■ Address Translation Scheme

- Logical address generated by CPU is divided into two parts:
  - Page number (**p**) – used as an index into a page table which contains the base address of each page in physical memory.
  - Page offset/displacement (**d**) – combined with base address to define the physical memory address that is sent to the memory unit.
- For given logical address space  $2^m$  and page size  $2^n$  ( $n < m$ ),



- By using a page size of a power of 2, the pages are *invisible* to the programmer, compiler/assembler, and the linker.
- Address translation at run-time is then easy to implement in hardware:
  - Logical address **<p, d>** gets translated to physical address **<f, d>** by indexing the page table with **p** and *appending* the same offset **d** to the frame of number **f**.


## ■ Address Translation Architecture



## Address Translation Architecture

### Example.

$m = 4, n = 2$ , page size =  $2^2 = 4$  (bytes), memory size = 32 (bytes).

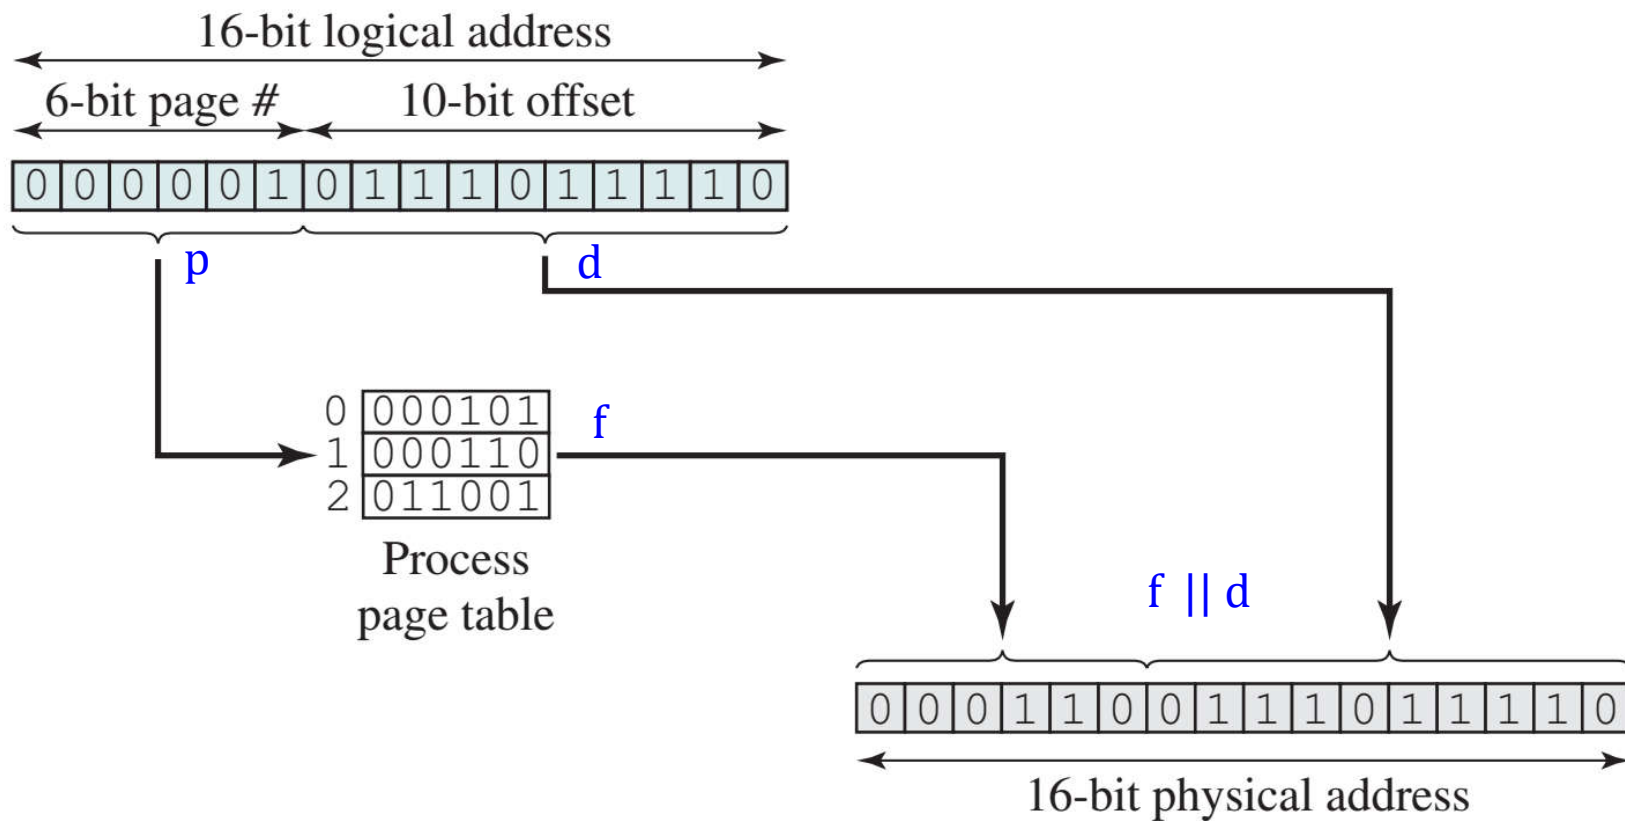




## Address Translation Architecture

### Example.

$m = 16, n = 10$ , page size =  $2^{10} = 1024$  (bytes).






### ■ Implementing Page Table

- If we keep Page Table in main memory
  - *Page-table base register (PTBR)* points to the location of the page table in memory.
  - *Page-table length register (PTLR)* indicates the size of the page table.
  - Every data/instruction access requires *two memory accesses*.
    - one for the page table, and one for the data/instruction
- If we keep Page Table in hardware (in MMU?)
  - However, page table can be large – too expensive.
- The two memory accesses problem can be solved by combining these two mechanisms.
  - Use a special fast-lookup hardware cache called *Associative Memory* (Registers) or *Translation Look-aside Buffer (TLB, 快表)* – enables fast parallel search.
  - Address translation <p, d>
    - If p is in associative register, get frame number out.
    - Otherwise get frame number from page table in memory.

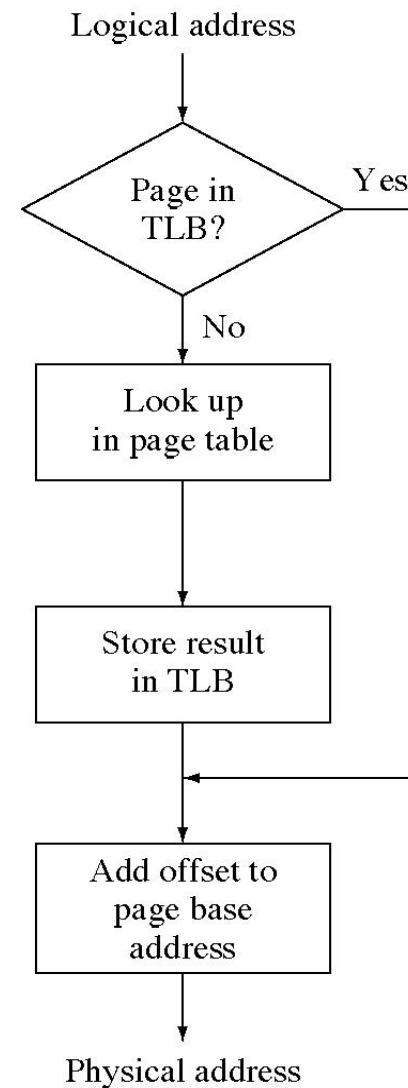
## ■ Implementing Page Table

### ■ Paging Hardware with TLB.



## ■ Implementing Page Table

- Paging Hardware with TLB.
- TLB Flow Chart.





## ■ Implementing Page Table

- Paging Hardware with TLB
  - TLB takes advantage of the *Locality Principle* (局部性原理).
  - TLB uses associative mapping hardware to simultaneously interrogate all TLB entries to find a match/hit on page number.
  - TLB hit rates are 90%+.
  - TLB must be flushed (erased) each time a new process enters the running state.
    - There is only one global TLB serving for all processes.
  - TLB information maybe kept/loaded in/from process context.



### ■ Implementing Page Table

#### ■ Paging Hardware with TLB

- Each associative (memory) lookup time =  $\epsilon$  time unit.
  - time unit: memory access time
  - can be < 10% of (far faster than) memory access time
- Hit ratio =  $\alpha$ 
  - percentage of times that a page number is found in the associative memory;
  - related to number of associative registers.
- Effective Access Time (EAT)
$$\text{EAT} = \alpha(\epsilon + 1) + (1 - \alpha)(\epsilon + 2) = 2 + \epsilon - \alpha.$$
Unit: memory access time.
- EAT is between  $1 + \epsilon$  and  $2 + \epsilon$  access times.
  - should be closer to 1.

## ■ Implementing Page Table

### ■ Paging Hardware with TLB

#### ■ Effective Access Time (EAT)

$$\text{EAT} = \alpha(\epsilon + 1) + (1 - \alpha)(\epsilon + 2) = 2 + \epsilon - \alpha.$$

Unit: memory access time.

#### ■ Examples: Assume that memory cycle time is 100ns and $\epsilon = 20\text{ns}$ for TLB search.

- Consider  $\alpha = 80\%$ . Then

$$\text{EAT} = 0.80 \times (20 + 100) + 0.20 \times (20 + 200) = 140(\text{ns}), \text{ or}$$

$$\text{EAT} = 2 \times 100 + 20 - 1 \times 100 \times 0.8 = 140(\text{ns})$$

So it is  $(140 - 100) / 100 = 40\%$  slowdown in memory access time.

- Consider more realistic hit ratio  $\alpha = 98\%$ . Then

$$\text{EAT} = 0.98 \times 120 + 0.02 \times 220 = 122(\text{ns}), \text{ or}$$

$$\text{EAT} = 2 \times 100 + 20 - 1 \times 100 \times 0.98 = 122(\text{ns})$$

So it is only  $(122 - 100) / 100 = 22\%$  slowdown in memory access time.

## ■ Implementing Page Table

### ■ Paging Hardware with TLB

#### ■ Advanced TLB Aspects

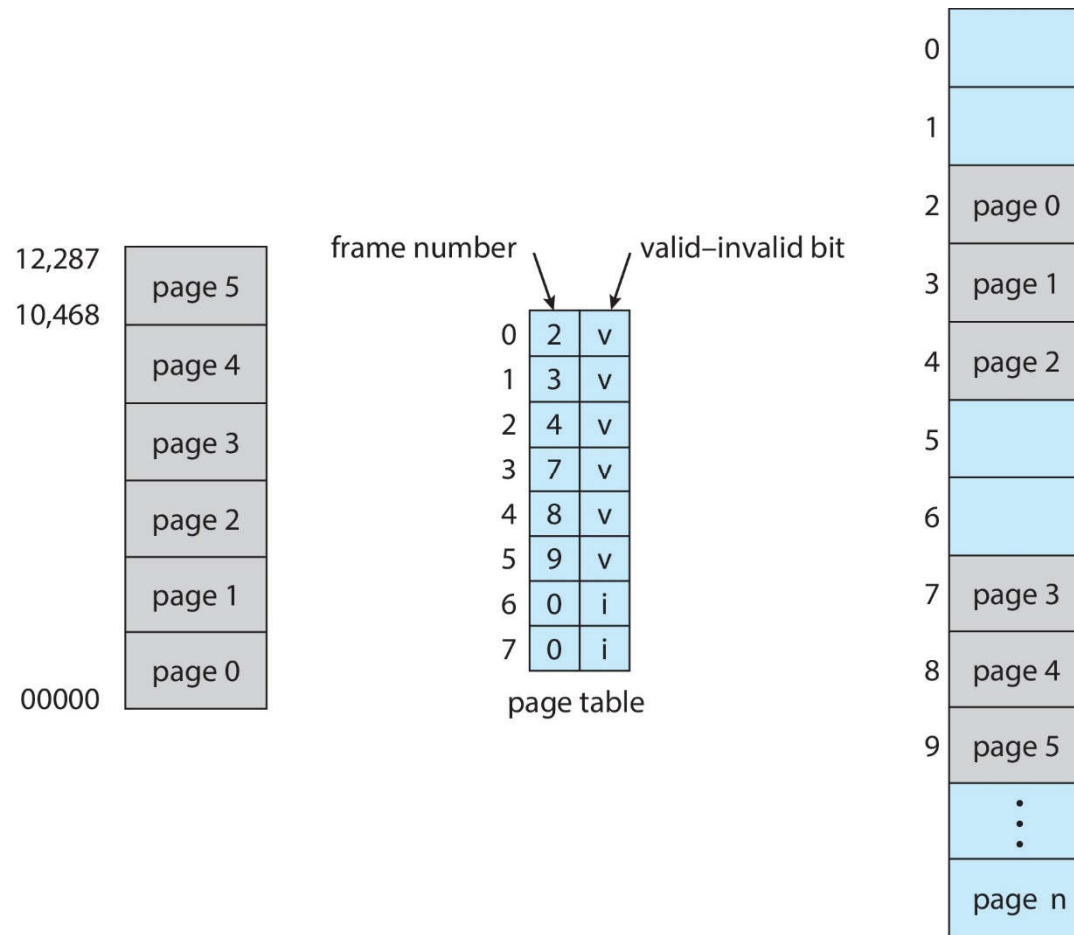
- Some TLBs store address-space identifiers (ASIDs) in each TLB entry – uniquely identifies each process to provide address-space protection for that process:
  - Otherwise need to flush at every context switch.
- TLBs typically small (64 to 1,024 entries).
- On a TLB miss, value is loaded into the TLB for faster access next time:
  - Replacement policies must be considered.
  - Some entries can be *wired down* (not removed from the TLB) for permanent fast access.

## ■ Memory Protection


- Memory protection is implemented by associating protection bit with *each frame* to indicate if read-only or read-write access is allowed.
  - can also add more bits to indicate page execute-only, and so on.
- Valid-invalid bit is attached to each entry in the page table.
  - “valid” indicates that the associated page is in the process’ logical address space, and is thus a legal page.
  - “invalid” indicates that the page is not in the process’ logical address space.
  - Or use *page-table length register (PTLR)*
- Any violations result in a trap to the kernel.

## Memory Protection

- Valid (v) or Invalid (i) Bit in Page Table.



## ■ Transfer of a Paged Memory to Contiguous Disk Space



## ■ Shared Pages

### ■ Shared code

- One copy of read-only (*reentrant*) code shared among processes (e.g., text editors, compilers, window systems).
  - Similar to multiple threads sharing the same process space.
- Shared code must appear in the same physical location in the logical address spaces of all processes.

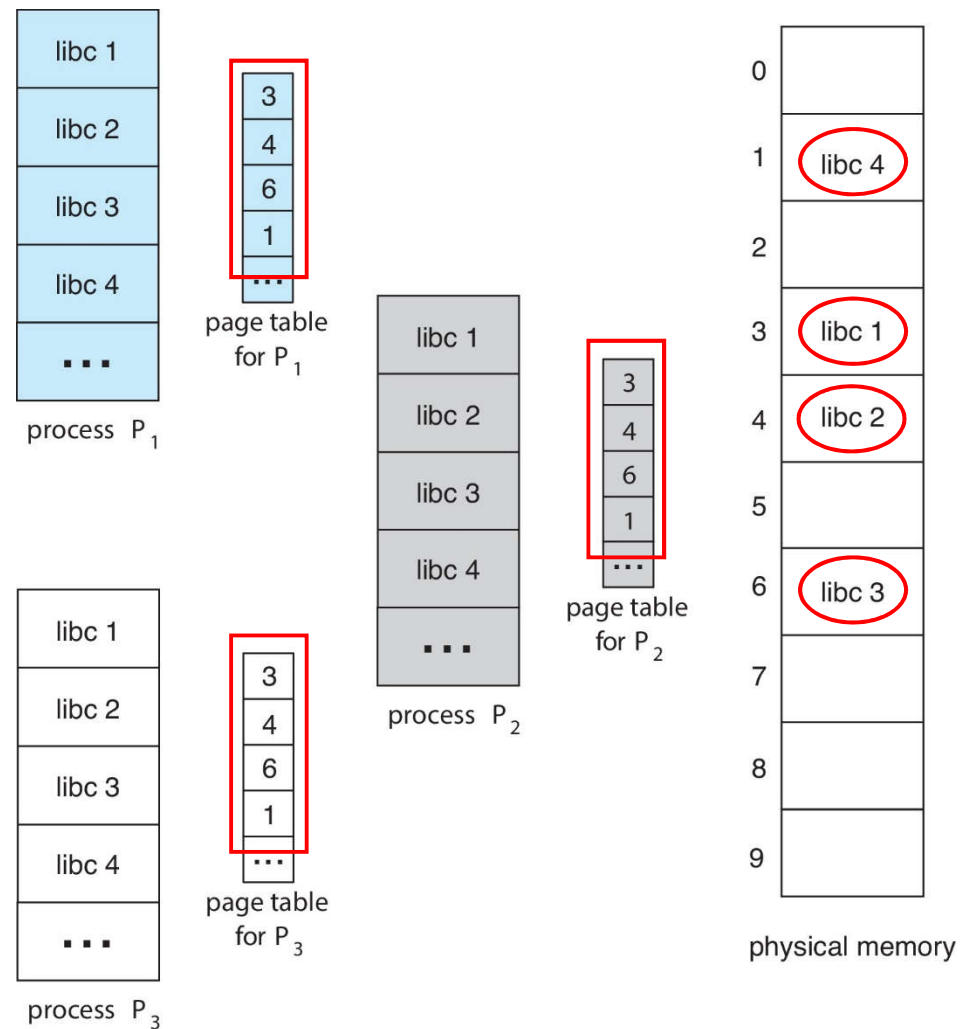
### ■ Private code and data

- Each process keeps separate copy of code and data.
- The pages for the private code and data can appear anywhere in the logical address space.



## Shared Pages

### Shared Pages Example.






### ■ Structures of Page Tables

- Memory structures for paging can get huge using straight-forward methods.
  - Consider a 32-bit logical address space as on modern computers with page size of 4KB ( $2^{12}$ ).
  - Page table would have 1 million entries ( $2^{32 - 12} = 2^{20}$ ).
  - If each entry is 4 bytes, then 4MB of physical address space is needed for page table alone.
    - That amount of memory used to cost a lot.
    - It is hard to allocate that contiguously in main memory.
- There are three structures for page tables.
  - Hierarchical Page Tables
  - Hashed Page Tables
  - Inverted Page Tables.

## Structures of Page Tables

### Hierarchical Page Tables

- Break up the logical address space into multiple page tables.
  - A simple technique is a two-level page table.



A two-level page-table scheme




## Structures of Page Tables

### Hierarchical Page Tables

Consider a 32-bit logical address space.

- A logical address (on 32-bit machine with  $2^{12} = 4K$  page size) is divided into:
  - a page number consisting of 20 bits
  - a page offset consisting of 12 bits.
- Since the page table is paged, the page number is further divided into:
  - a 10-bit outer page number
  - a 10-bit inner page offset.
- Thus, a logical address is as follows:

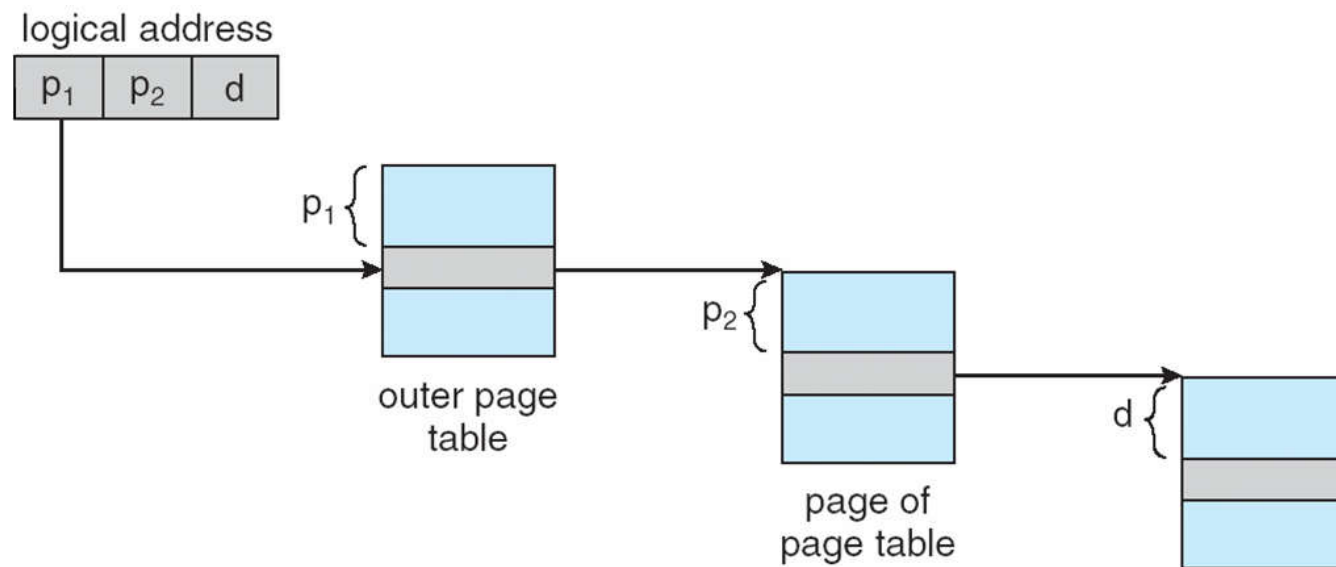


where  $p_1$  is an index into the **outer** page table, and  $p_2$  is the **displacement** within the page of the **inner** page table.

- Known as **forward-mapped page table**

## Structures of Page Tables

### Hierarchical Page Tables.




Two-level Paging Scheme.

## ■ Structures of Page Tables

### ■ Hierarchical Page Tables

#### ■ Consider a 64-bit Logical Address Space.


- Even two-level paging scheme is not sufficient.
- If page size is 4KB ( $= 2^{12}$ ), then page table has  $2^{52}$  entries.
  - For two level scheme, inner page tables could have  $2^{10}$  entries of 4 bytes length. Address would look like



- Outer page table has  $2^{42}$  entries ( $2^{44}$  bytes, 4 bytes for each entry).
- One solution is to add a 2<sup>nd</sup> outer page table.
- But in the following example the 2<sup>nd</sup> outer page table is still  $2^{34}$  bytes in size.
  - And possibly 4 memory access to get to one physical memory location.

## ■ Structures of Page Tables

- Hierarchical Page Tables
  - Three-level Paging Scheme.



## ■ Structures of Page Tables

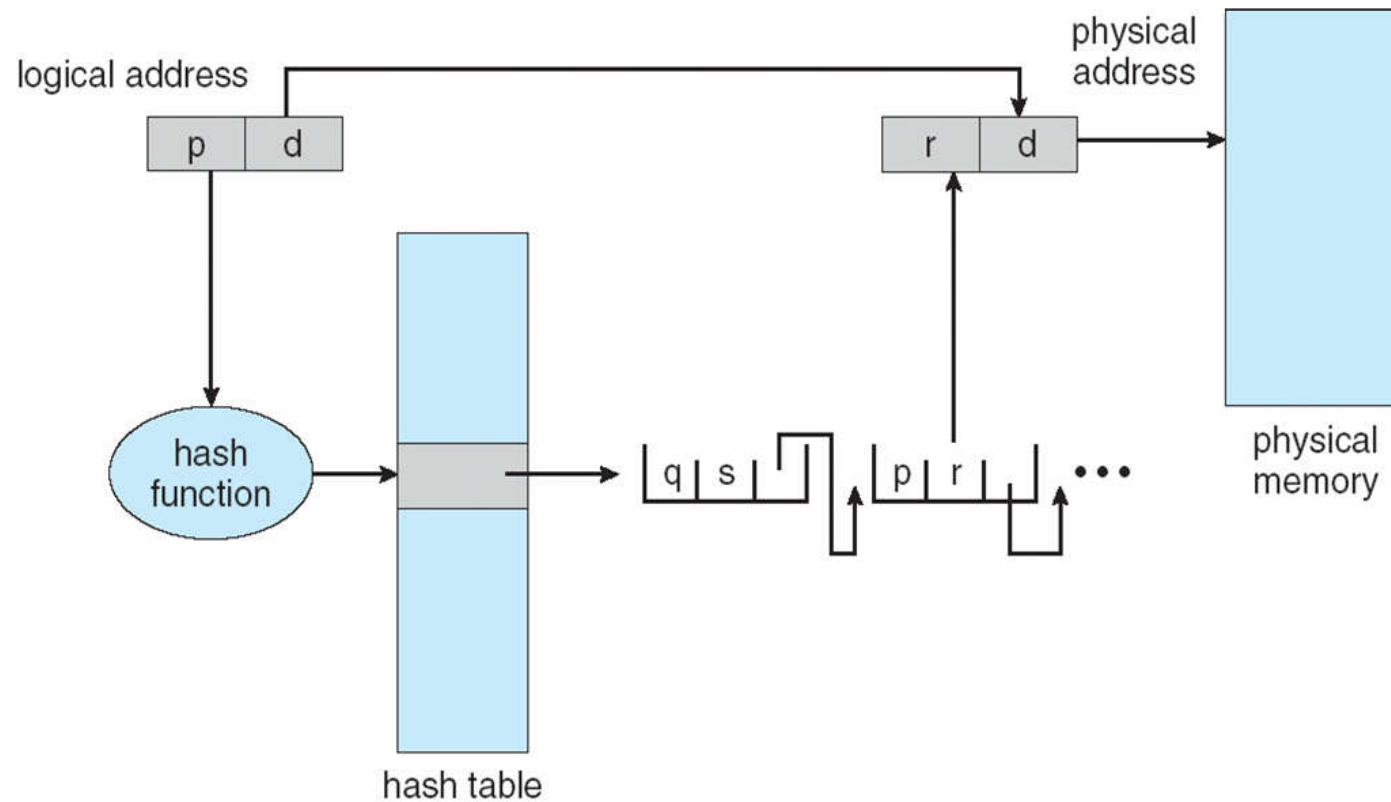
### ■ Hashed Page Tables

- Common in address spaces  $> 32$  bits.
- The virtual page number is hashed into a page table.
  - This page table contains a chain of elements hashing to the same location.
- Each element of the chain table contains:
  - (1) The virtual page number
  - (2) The value of the mapped page frame
  - (3) A pointer to the next element.
- Virtual page numbers are compared in this chain searching for a match.
  - If a match is found, the corresponding physical frame is extracted.
- Variation for 64-bit addresses is *clustered page tables*.
  - Similar to hashed but each entry refers to several pages (such as 16) rather than 1.
  - Especially useful for *sparse* address spaces (where memory references are non-contiguous and scattered).



## Structures of Page Tables

### Hashed Page Tables



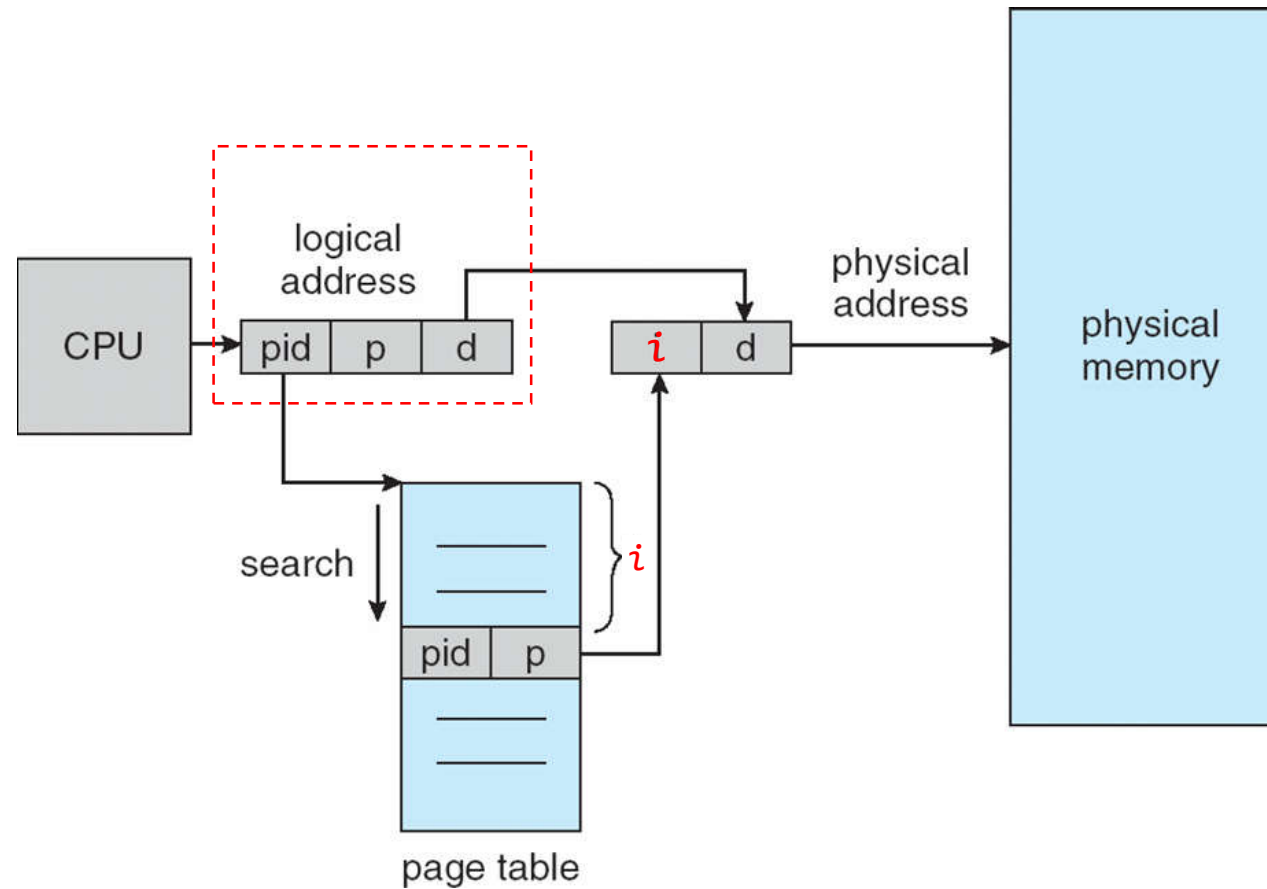
## ■ Structures of Page Tables

### ■ Inverted Page Table

- Rather than each process having a page table and keeping track of all possible logical pages, track all physical pages.
- Entries of an inverted page table consists of the virtual address of the page stored in that real (physical) memory location, with information about the process that owns that page.
  - one entry for each real page of memory
    - *e.g.*, <process-id, page-number>
  - used by 64-bit Ultra SPARC, PowerPC, ...
- This scheme decreases memory needed to store each page table, but increases time needed to search the table when a page reference occurs.
- Use hash table to limit the search to one — or at most a few — page-table entries.
  - TLB can accelerate access.
- But how to implement shared memory?
  - One mapping of a virtual address to the shared physical address.

## Structures of Page Tables

### Inverted Page Table



Inverted page table architecture

## ■ Simple Segmentation/Paging Comparison

- Segmentation is visible to the programmer whereas Paging is transparent.
- Naturally supports protection/sharing.
- Segmentation can be viewed as commodity offered to the programmer to logically organize a program into segments while using different kinds of protection (example: execute-only for code but read-write for data).
- Segments are variable-size; Pages are fixed-size.
- Segmentation requires more complicated hardware for address translation than Paging.
- Segmentation suffers from external fragmentation. Paging only yields a small internal fragmentation.
- Maybe combine Segmentation and Paging?





### ■ The Intel 32 Architecture

- Dominant industry chips
- Pentium CPUs are 32-bit and called IA-32 architecture
  - Current Intel CPUs are 64-bit and called IA-64 architecture
- Many variations in the chips, cover the main ideas here
- Supports both segmentation and segmentation with paging
  - Each segment can be 4GB
  - Up to 16K segments per process
  - Divided into two partitions
    - First partition of up to 8K segments are private to process (kept in *local descriptor table* (LDT))
    - Second partition of up to 8K segments shared among all processes (kept in *global descriptor table* (GDT))
- CPU generates logical address
  - Given to segmentation unit which produces linear addresses.
  - Linear address given to paging unit:
    - Which generates physical address in main memory.
    - Paging units form equivalent of MMU.
    - Pages sizes can be 4KB or 4MB.



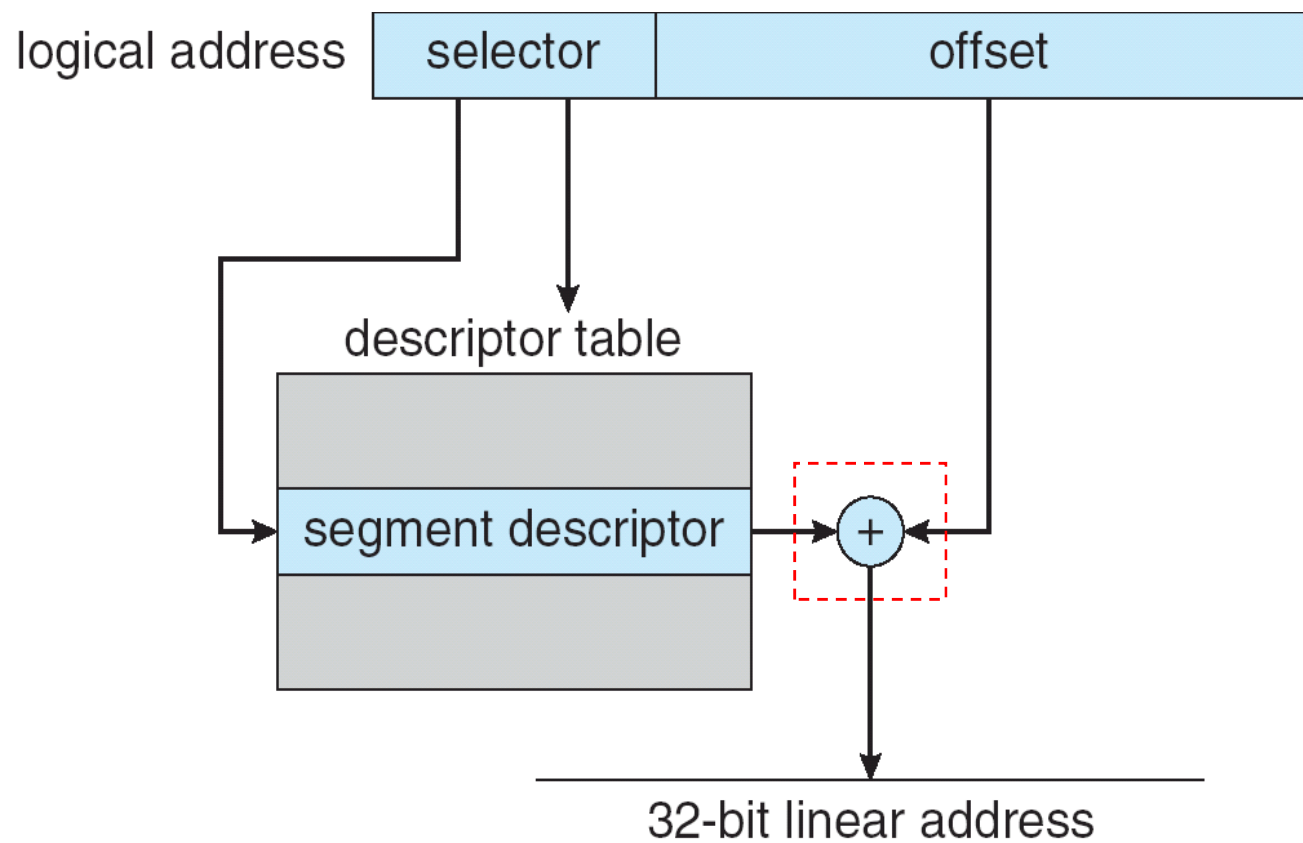
## ■ The Intel 32 Architecture

- Logical to Physical Address Translation.




## ■ The Intel 32 Architecture

- Intel Pentium Segmentation.



## ■ The Intel 32 Architecture

### ■ Intel Pentium Paging Architecture.



Two-level paging scheme  
for 4 KB page.


Direct paging scheme  
for 4 MB page.



## ■ The Intel 32 Architecture

### ■ Intel IA-32 Page Address Extensions

- 32-bit address limits led Intel to create *page address extension (PAE)*, allowing 32-bit apps access to more than 4GB of memory space.
  - Paging went to a three-level scheme
  - Top two bits refer to a *page directory pointer table*
  - Page-directory and page-table entries moved to 64-bits in size
  - Net effect is increasing address space to 36 bits – 64GB of physical memory






### ■ The Intel 32 Architecture

#### ■ Intel X86-64

- current generation Intel x86 architecture
- 64 bits is ginormous (> 16 exabytes)
- in practice only implement 48 bit addressing
  - Page sizes of 4KB / 2MB / 1GB
  - Four levels of paging hierarchy
- can also use PAE so virtual addresses are 48 bits and physical addresses are 52 bits



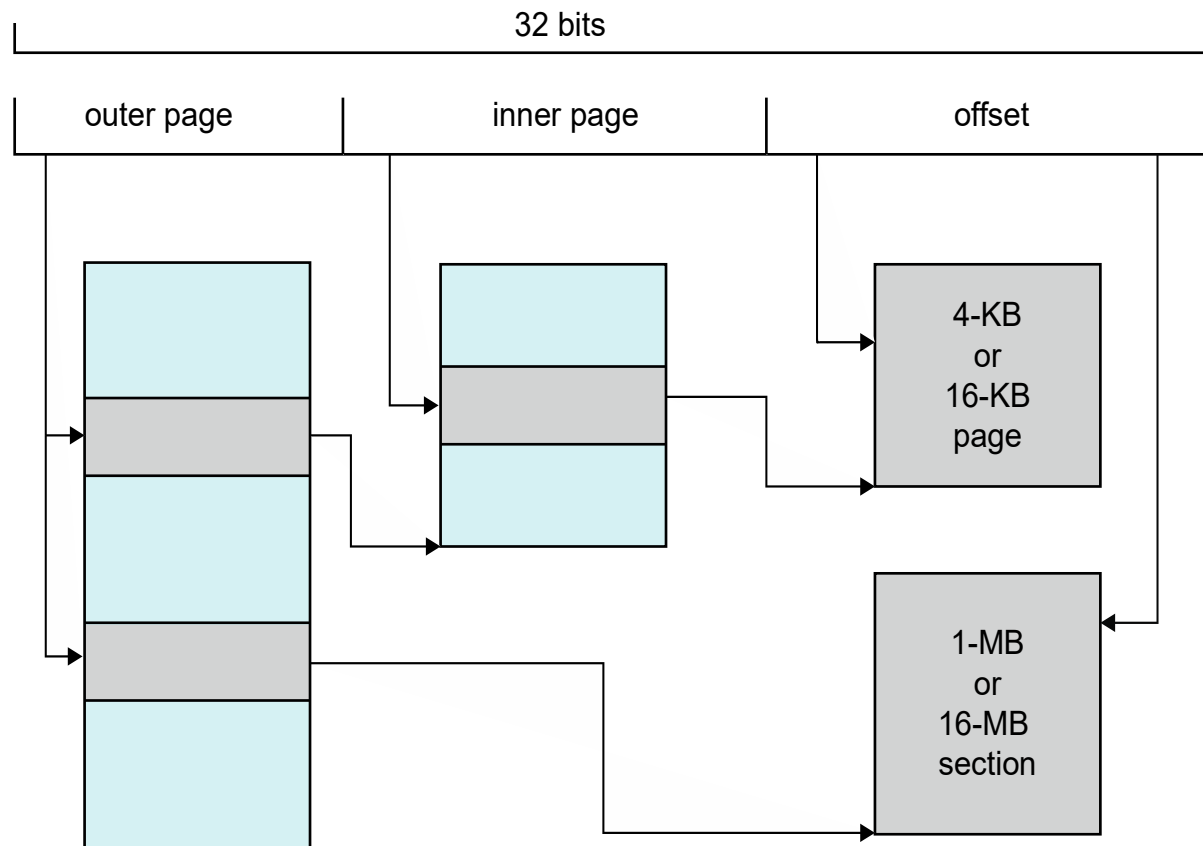


### ■ ARM Architecture

- Dominant mobile platform chip (Apple iOS and Google Android devices for example)
- Modern, energy efficient, 32-bit CPU
- 4KB / 16KB pages
- 1MB / 16MB pages (termed *sections*)
- One-level paging for sections, two-level for smaller pages
- Two levels of TLBs
  - Outer level has two micro TLBs (one for data, one for instruction)
  - Inner is single main TLB
  - First inner is checked, on miss outers are checked, and on miss page table walk performed by CPU



## ■ ARM Architecture



## Linux

### Three-level Paging in Linux.

