

---

# Virtual Memory Policies

## Operating Systems

---

School of Data & Computer Science  
Sun Yat-sen University

Lecture Notes: [os\\_sysu@163.com](mailto:os_sysu@163.com)  
Instructor: Guoyang Cai  
email: [isscg@mail.sysu.edu.cn](mailto:isscg@mail.sysu.edu.cn)





## ■ Contents

- Virtual Memory Policies Need to decide on:
  - Fetch policy
  - Placement policy
  - Replacement policy
  - Resident Set Management
  - Load Control



### ■ **Memory Management Software**

- Memory management software depends on whether the hardware supports paging or segmentation or both.
- Pure segmentation systems are rare. Segments are usually paged -- memory management issues are then those of paging.
- We shall thus concentrate on issues associated with paging.
- To achieve good performance we need a low page fault rate.



### ■ Fetch policy

- Determines when a page should be brought into main memory. Two common policies:
  - Demand Paging only brings pages into main memory when a reference is made to a location on the page (i.e., paging on demand only).
    - many page faults when process first starts but should decrease as more pages are brought in.
  - Prepaging brings in pages whose use is anticipated:
    - locality of references suggest that it is more efficient to bring in pages that reside contiguously on the disk.
    - efficiency not definitely established: the extra pages brought in are “often” not referenced.



### ■ Placement Policy

- Determines where in real memory a process piece resides.
- For paging (and paged segmentation):
  - the hardware decides where to place the page: the chosen frame location is irrelevant since all memory frames are equivalent (not an issue).
- For pure segmentation systems:
  - first-fit, next fit... are possible choices (a real issue).



### ■ Replacement Policy

- Deals with the selection of a page in main memory to be replaced when a new page is brought in.
- This occurs whenever main memory is full (no free frame available).
- Occurs often since the OS tries to bring into main memory as many processes (pages) as it can to increase the multiprogramming level.
- Not all pages in main memory can be selected for replacement.
- Some frames are locked (cannot be paged out):
  - much of the kernel is held on locked frames as well as key control structures and I/O buffers.
- The OS might decide that the set of pages considered for replacement should be:
  - limited to those of the process that has suffered the page fault.
  - the set of all pages in unlocked frames.

## ■ Replacement Policy

- The decision for the set of pages to be considered for replacement is related to the resident set management strategy:
  - how many page frames are to be allocated to each process?
- No matter what is the set of pages considered for replacement, the replacement policy deals with algorithms that will choose the page within that set.



### ■ Allocation of Frames

#### ■ Resident Set Management

- The OS must decide how many page frames to allocate to a process:
  - large page fault rate if too few frames are allocated.
  - low multiprogramming level if too many frames are allocated.

#### ■ Resident Set Size

##### ■ Fixed-allocation policy:

- allocates a fixed number of frames that remains constant over time:
  - the number is determined at load time and depends on the type of the application.

##### ■ Variable-allocation policy:

- the number of frames allocated to a process may vary over time:
  - may increase if page fault rate is high.
  - may decrease if page fault rate is very low.
- requires more OS overhead to assess behavior of active processes.



## ■ Allocation of Frames

### ■ Fixed Allocation

#### ■ Equal Allocation

- for example, if there are 100 frames and 5 processes, give each process 20 frames.
- keep some as free frame buffer pool

#### ■ Proportional Allocation

- allocate according to the size of process.
- dynamic as degree of multiprogramming, process sizes change

Let  $s_i$  = size of process  $P_i$ ,  $m$  = total number of frames

Then  $S = \sum s_i$

$a_i$  = allocation for  $p_i = (s_i / S) \times m$ .

- Example

$$s_1 = 10, s_2 = 127, m = 64$$

$$S = \sum s_i = 10 + 127 = 137$$

$$a_1 = (10 / 137) \times 64 \approx 5$$

$$a_2 = (127 / 137) \times 64 \approx 59$$

## ■ Allocation of Frames

### ■ Fixed Allocation

#### ■ Priority Allocation

- Use a proportional allocation scheme using priorities rather than size.
- If process  $P_i$  generates a page fault:
  - select for replacement one of its frames.
  - select for replacement a frame from a process with lower priority number.

## ■ Allocation of Frames

### ■ Replacement Scope

- Replacement scope is the set of frames to be considered for replacement when a page fault occurs.
- *Global replacement* – process selects a replacement frame from the set of all frames; one process can take a frame from another
  - But then process execution time can vary greatly
  - But greater throughput so more common
- *Local replacement* – each process selects from only its own set of allocated frames
  - More consistent per-process performance
  - But possibly underutilized memory
- We will consider the possible combinations of replacement scope and resident set size policy.



### ■ Allocation of Frames

- Fixed Allocation + Local Scope
  - Each process is allocated a fixed number of pages:
    - determined at load time; depends on application type.
  - When a page fault occurs, page frames considered for replacement are local to the page-fault process:
    - the number of frames allocated is thus constant.
    - previous replacement algorithms can be used.
  - Problem: difficult to determine ahead of time a good number for the allocated frames:
    - if too low: page fault rate will be high.
    - if too large: multiprogramming level will be too low.
- Fixed Allocation + Global Scope
  - Impossible to achieve:
    - if all unlocked frames are candidate for replacement, the number of frames allocated to a process will necessary vary over time.

## ■ Allocation of Frames

- Variable Allocation + Global Scope
  - Simple to implement -- adopted by many OSs (like Unix SVR4).
  - A list of free frames is maintained:
    - when a process issues a page fault, a free frame (from this list) is allocated to it.
    - Hence the number of frames allocated to a page fault process increases.
    - The choice for the process that will loose a frame is arbitrary; It's far from optimal.
  - Page buffering can alleviate this problem since a page may be reclaimed if it is referenced again soon.



### ■ Allocation of Frames

- Variable Allocation + Local Scope
  - May be the best combination (used by Windows NT).
  - Allocate at load time a certain number of frames to a new process based on application type:
    - use either pre-paging or demand paging to fill up the allocation.
  - When a page fault occurs, select the page to replace from the resident set of the process that suffered the page fault.
  - Reevaluate periodically the allocation provided and increase or decrease it to improve the overall performance.

## ■ Allocation of Frames

### ■ The Working Set Strategy

- The working set strategy is a variable-allocation method with local scope based on the assumption of locality of references.
- The working set for a process at time  $t$ ,  $W(D, t)$ , is the set of pages that have been referenced in the last  $D$  virtual time units:
  - virtual time = time elapsed while the process was in execution.
  - $D$  is a window of time.
  - $W(D, t)$  is an approximation of the program's locality.
- The working set of a process first grows when it starts executing.
- Then stabilizes by the principle of locality.
- It grows again when the process enters a new locality (transition period):
  - up to a point where the working set contains pages from two localities.
- Then decreases after a sufficient long time spent in the new locality.

## ■ Allocation of Frames

### ■ The Working Set Strategy

- The working set concept suggests the following strategy to determine the resident set size:
  - Monitor the working set for each process.
  - Periodically remove from the resident set of a process those pages that are not in the working set.
  - When the resident set of a process is smaller than its working set, allocate more frames to it:
    - If not enough free frames are available, suspend the process (until more frames are available), i.e., a process may execute only if its working set is in main memory.



## ■ Allocation of Frames

### ■ The Working Set Strategy

#### ■ Working-Set Model

- $\Delta$  = working-set window = a fixed number of page references
  - Example: 10,000 instructions
- $WSS_i$  (working set of Process  $P_i$ ) = total number of pages referenced in the most recent  $\Delta$  (varies in time):
  - if  $\Delta$  too small it will not encompass entire locality.
  - if  $\Delta$  too large it will encompass several localities.
  - if  $\Delta = \infty \Rightarrow$  will encompass entire program.
- $D = \sum WSS_i \equiv$  total demand frames
- if  $D > m \Rightarrow$  Thrashing
- Policy if  $D > m$ , then suspend one of the processes




## Allocation of Frames


- The Working Set Strategy
  - Working-Set Model.

page reference table

... 2 6 1 5 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...



$WS(t_1) = \{1, 2, 5, 6, 7\}$



$WS(t_2) = \{3, 4\}$

## ■ Allocation of Frames

### ■ The Working Set Strategy

#### ■ Keeping Track of the Working Set

- Approximate with interval timer + a reference bit.
- Example:  $\Delta = 10,000$ 
  - Timer interrupts after every 5000 time units.
  - Keep in memory 2 bits for each page.
  - Whenever a timer interrupts, copy and set the values of all reference bits to 0.
  - If one of the bits in memory = 1  $\Rightarrow$  page in working set.
- Why is this not completely accurate?
- Improvement = 10 bits and interrupt every 1000 time units.



### ■ Allocation of Frames

#### ■ The Working Set Strategy

##### ■ Practical problems with this working set strategy:


- measurement of the working set for each process is impractical:
  - necessary to time stamp the referenced page at every memory reference.
  - necessary to maintain a time-ordered queue of referenced pages for each process.
- the optimal value for  $D$  is unknown and time varying.

##### ■ Solution: rather than monitor the working set, monitor the page fault rate!

## Allocation of Frames

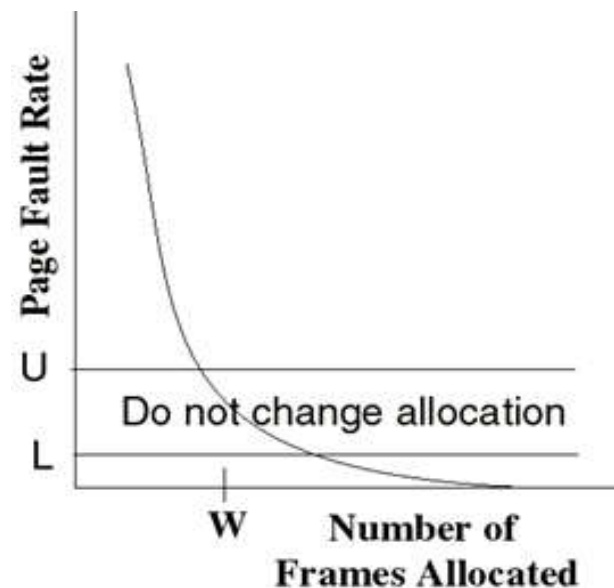
### Page-Fault Frequency (PFF) Scheme

- Establish “acceptable” page-fault rate:
  - If actual rate too low, process loses frame.
  - If actual rate too high, process gains frame.



## Allocation of Frames


- Page-Fault Frequency (PFF) Strategy
  - Define an upper bound  $U$  and lower bound  $L$  for page fault rates.
  - Allocate more frames to a process if fault rate is higher than  $U$ .
  - Allocate less frames if fault rate is less than  $L$ .
  - The resident set size should be close to the working set size  $W$ .
  - We suspend the process if the  $PFF > U$  and no more free frames are available.






### ■ Allocation of Frames

- Working Sets and Page-Fault Rates
  - Direct relationship between working set of a process and its page-fault rate
  - Working set changes over time
  - Peaks and valleys over time



## ■ Load Control

- Determines the number of processes that will be resident in main memory (i.e., the multiprogramming level):
  - Too few processes: often all processes will be blocked and the processor will be idle.
  - Too many processes: the resident size of each process will be too small and flurries of page faults will result: thrashing.





## ■ Load Control

- A working set or page fault frequency algorithm implicitly incorporates load control:
  - only those processes whose resident set is sufficiently large are allowed to execute.
- Another approach is to adjust explicitly the multiprogramming level so that the mean time between page faults equals the time to process a page fault:
  - performance studies indicate that this is the point where processor usage is at maximum.

## ■ Load Control

### ■ Process Suspension

- Explicit load control requires that we sometimes swap out (suspend) processes.
- Possible victim selection criteria:
  - Faulting process
    - this process may not have its working set in main memory so it will be blocked anyway.
  - Last process activated
    - this process is least likely to have its working set resident.
  - Process with smallest resident set
    - this process requires the least future effort to reload.
  - Largest process
    - will yield the most free frames.